Parallelization of the NIM Dynamical Core for GPUs



Mark Govett
Jacques Middlecoff, Tom Henderson,
Jim Rosinski, Craig Tierney



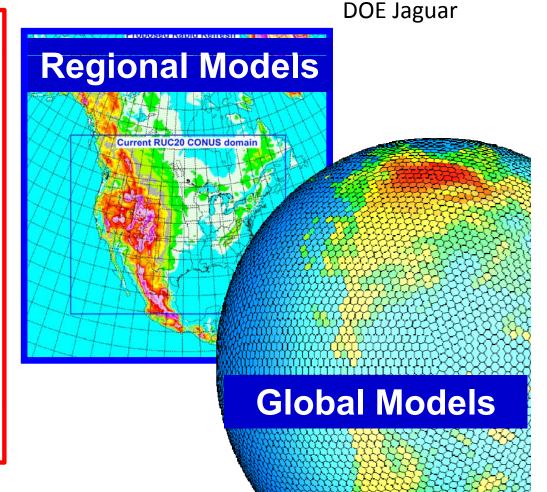
CPU

- Bigger Systems
- More Expensive Facilities
- Bigger Power Bills
- Lower System Reliability

GPU

- Faster
- Less power

Lower cost



GPU & MIC Hardware

NVIDIA: Fermi chip first to support HPC

Tesla (2008)
 240 cores, 933 SP GFlops, CUDA

Fermi (2010)512 cores, 662 DP Gflops, ECC memory

Kepler (2012)
 28nm - - > 3-4x performance / watt

- Maxwell (2014?) Integrated CPU & GPU

AMD/ATI: Graphics chip primarily

- Fusion (2011): integrates CPU & GPU

Intel MIC: Many Integrated Core (2012), 32-64 cores

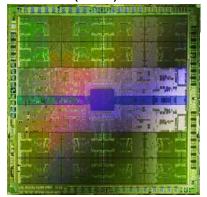
Large memory, cache

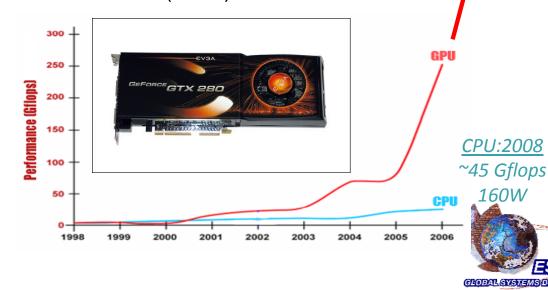
– x86 => existing tools work

TACC plans a large cluster

NVIDIA: Fermi (2010) 3 billion transistors

448 (512) cores





Tesla (2008)

GPU: 2008

933Gflops

150W



Application Performance

GPU

Block (0, 0)

Registers

Thread (0, 0)

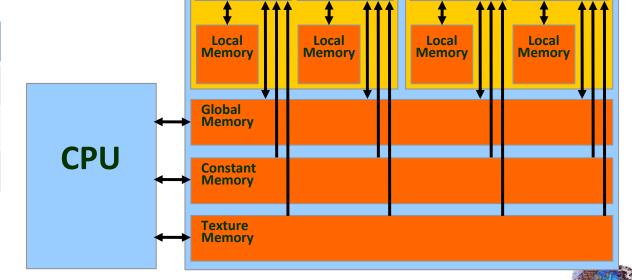
Shared Memory

Registers

Thread (1, 0)

- NVIDIA GPU: a co-processor to CPU
 - Data moved across PCI bus
- Efficient use of memory is critical to good performance
 - 1-2 cycles to access shared memory
 - Hundreds of cycles to access global memory

Memory	Tesla	Fermi
Shared	16K	16/48K
Cache		16/48K
Global	1-2GB	4-6GB





GPU Multi-layer Memory (Tesla)

Block (1, 0)

Registers

Thread (0, 0)

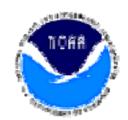
Shared Memory

Registers

Thread (1, 0)

GPU Programming Approaches

- Language Approach
 - CUDA, OpenCL, CUDA Fortran, etc.
 - Less portable to different architectures
 - Requires that separate versions be maintained
 - In practice this rarely works too costly, difficult
- Directive-based Approach
 - Appear as comments in the source
 - !ACC\$DO VECTOR (1)
 - Compilers can analyze and (hopefully) generate efficient code
 - Dependent on maturity





Nonhydrostatic Icosahedral Model (NIM)

- Team of scientists, parallel programmers and computer scientists
 - Designed for GPU, fast parallel execution

NIM Development Team

- •Jin Lee
- A.E. MacDonald
- •Jung-Eun (Esther) Kim
- •Ka Yee Wong
- •Jian-Wen Bao
- Ning Wang
- Jacques Middlecoeff
- Mark Govett
- Tom Henderson
- •Jim Rosinski

Collaborators/Contributors:

•YSU: Prof. Songyou Hong, Sueng-On Hwang(KMA)

•CSU: Prof. David Randall, Todd Jones

•PSD: Drs. George Kiladis, Stefan N. Tulich.

•Purdue U: Prof. Wen-Yih Sun

•GFDL: Dr. S.-J. Lin



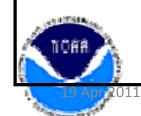


ESRL finite-volume Icos- models (FIM/NIM)

ESMF

- FIM (flow-following finitevolume Icosahedral model): A hydrostatic model for NCEP global model ensemble.
- A hydrostatic model consists of 2-D finite-volume SWM coupled with hybrid σ-θ vertical solver.
- Produce accurate medium-range weather forecasts with scores comparable to GFS.
- Development began in ~2000

- NIM (Nonhydrostatic Icosahedral model): A multi-scale global model for weather and intraseasonal climate predictions.
- Extension of 2-D finite-volume integration into 3-D integration on control volume defined on the height coordinate.
- Use the latest GPU technology to speed up high-resolution model calculations
- Targeting 2KM global scale
- Development began in ~2005





Novel features of FIM/NIM

- Finite-volume Integrations on Local Coordinate
- Efficient Indirect Addressing Scheme on Irregular Grid
 - Adopted by MPAS icos model at NCAR
- FIM: Hybrid σ - θ Coordinate w/ GFS Physics
- Conservative and Monotonic Adams-Bashforth 3rd-order FCT Scheme
- Grid Optimization for Efficiency and Accuracy
- Novel Features of NIM
 - Three dimensional finite-volume integration
 - Runge-Kutta (RK) 4th solver for vertically propagating acoustic waves
 - Conservative and positive definite transport scheme





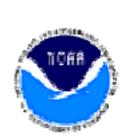
NIM Code Design

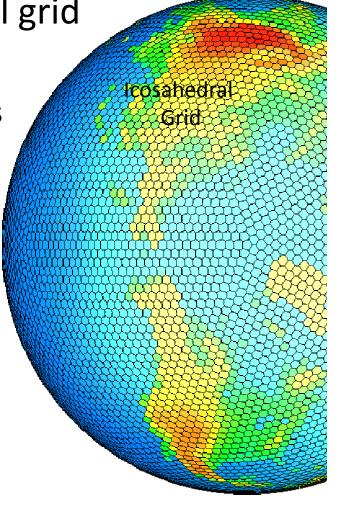
Uniform, hexagonal-based, icosahedral grid

• 1 horizontal dimension

 Novel indirect addressing scheme permits concise, efficient code

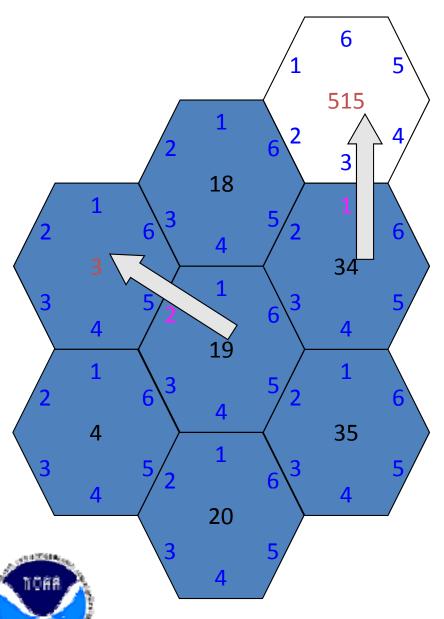
- Separated coarse and fine grain parallelization
 - CPU controls high level flow
 - Distributed memory parallelism (MPI)
 - Initialization, message passing, I/O
 - GPU executes dynamics routines
 - Data is resident in GPU memory
 - Data is passed to CPU only for I/O and inter GPU communications







NIM/FIM Indirect Addressing (MacDonald, Middlecoff)



- Single horizontal index
- Store number of sides (5 or 6) in "nprox" array
 - \blacksquare nprox(34) = 6
- Store neighbor indices in "prox" array
 - prox(1,34) = 515
 - prox(2,19) = 3
- >1% performance impact
 - Indirect reference is not the innermost dimension
- Very compact code



Simple Loop With Indirect Addressing

- Compute sum of all horizontal neighbors
 - nip = number of columns
 - nvl = number of vertical levels





F2C-ACC GPU Compiler

- Developed to speed parallelization of NIM
 - Commercial compilers were not available in 2008
- Translates Fortran to C or CUDA
 - Many (but not all) language features supported
 - Generates readable, debuggable code with original comments retained
- Ten directives for code parallelization, eg.

– !ACC\$REGION! Define GPU regions

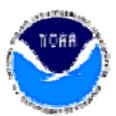
– !ACC\$DO! Identify loop level parallelism

– !ACC\$DATA! Move data between CPU and GPU

!ACC\$INSERT, ACC\$REMOVE
 ! Hand insertions / deletions where

translation is not available

- Continues to be developed
 - Until commercial compilers are better





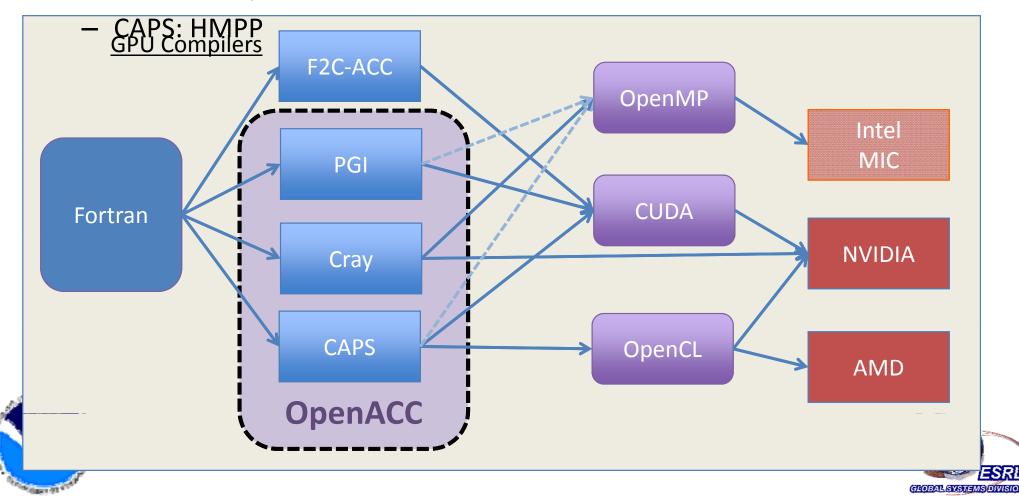
Directive-Based Fortran GPU

Compilers and Portability

OpenACC directives adopted by vendors

– PGI: Accel– Cray: OpenMP

F2C-ACC: OpenSource



Fortran GPU Compiler Results (2011)

Using NIM G5 - 10242 horizontal points, 96 vertical levels Fermi GPU vs. Intel Westmere CPU Socket

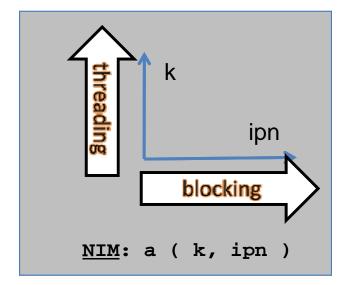
NIM routine	CPU 1- core Time (sec)	CPU 6- core Time (sec)	F2C-ACC GPU Time (sec)	HMPP GPU Time (sec)	PGI GPU Time (sec)	F2C-ACC Speedup vs. 6-core CPU
Total	8654	2068	449			4.6
vdmint s	4559	1062	196	192	197	5.4
vdmint v	2119	446	91	101	88	4.9
flux	964	175	26	24	26	6.7
vdn	131	86	18	17	18	4.8
diag	389	74	42	33		1.8
force	■ Us@d	PAPI peßfo	rmance 🕏	unters of (CPU (GP31	_) 4.7





Parallelization Factors for NIM

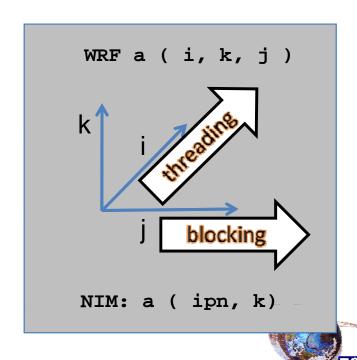
- Code design a dominant factor in performance
 - Weather and climate codes typically have a high memory access to compute ratio
 - Memory accesses limit performance
 - Data alignment led to a 10x improvement
- Data dependencies guide parallelization
 - Dynamics are in the horizontal
 - a [vert, horiz]
 - Physics are in the vertical column
 - a [horiz, vert]
 - Transpose needed to optimize memory accesses





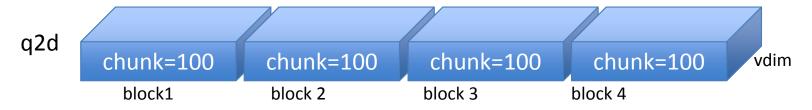
WRF Physics

- Legacy code
 - Community Model used worldwide for more than a decade
 - WRF-ARW, WRF-NMM, WRF-RR, WRF-CHEM, HWRF
- Traditional cartesian grid
 - 3D arrays (horizontal, vertical, horizontal) = = > array3D(i,k,j)
- Designed for CPU architectures
 - Primary calculations on 2D arrays (i,k)
 - Improves cache utilitzation
- Limited ability to change the code
- A challenge for GPU parallelization
 - Dependencies in vertical
 - GPU: threading in horizontal dimensions
- Further challenge for NIM
 - Only 1 horizontal dimension



Application Requirements NIM + WRF Physics

- Chunking
 - Assigning threads and blocks to the same dimension



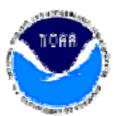
- Approach has been successful
 - ~2-3x performance improvement anticipated
 - No changes to WRF code base
- Supported in F2C-ACC
 - Working with vendors to provide this



Application Requirements

weather and climate codes

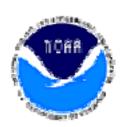
- Promotion of variables
 - Needed for correctness & performance
 - Added dimension needed to store block or thread parallel calculations
 - 2D array + + + > 3D array
- Demotion of variables
 - 1D arrays - -> GPU register variables
 - 2D arrays - -> shared memory arrays
 - Results in huge performance increases





Future Work

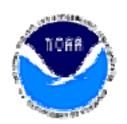
- NIM continues to be developed
 - physics integration
 - Aqua-planet simulations
 - Testing with real data this summer
- Committed to a single source
 - Performance portable between CPU, GPU, serial, parallel
 - NVIDIA, Intel MIC, AMD
 - We anticipate some challenges for legacy codes
 - FIM, WRF, HYCOM
 - Unclear performance of Intel MIC, tools





Conclusion

- F2C-ACC has proven useful
 - For NIM parallelization
 - To evaluate commercial compilers
 - Good interactions with vendors on improvements
- NIM GPU performance results encouraging
 - 5x faster than Intel-Westmere (socket-to-socket)
 - Plan to run NIM on Intel MIC in 2012
- Challenges Remain
 - Codes take too long to port to GPUs
 - MIC may be better, but price / performance unclear
 - Performance portability a concern
 - Standards for GPU directives





Final Thoughts

- Unclear how well industry will support us
 - Climate & weather only a small piece of their business
- Do we need to develop our own tools?
 - Preserve & adapt our models
 - Enhance code portability
 - Improve performance optimizations
 - Architecture specific changes
 - Array re-declarations
 - Loop re-ordering
 - Loop unrolling
 - Promotion & demotion of variables
 - Use of special memories, vector instructions
- These are all possible with existing compiler technology