



# DEEP

## Dynamical Exascale Entry Platform

2<sup>nd</sup> IS-ENES Workshop  
on “High performance computing for climate models”

30.01.2013, Toulouse, France

Estela Suarez


The research leading to these results has received funding from the European Community's Seventh Framework Programme (FP7/2007-2013) under *Grant Agreement* n° 287530

- DEEP: “Dynamical Exascale Entry Platform”
- EU-project funded by the FP7 program:
  - 2011 call for Exascale
  - 3 Projects selected:
    - DEEP, MontBlanc, CRESTA
  - DEEP EU-funding: 8.03 M€
- Duration: 3 years
  - Dec 2011 – Nov 2014
- 16 Partners
  - Coordinator: JSC


[www.deep-project.eu](http://www.deep-project.eu)




- **Build a prototype of an Exascale architecture:**
  - With accelerators that can react autonomously (→ “Booster”)
- **Hardware Development:**
  - Build a Booster with Intel® Xeon Phi™ and EXTOLL network
  - Energy efficient with “hot water” cooling
- **Software Development**
  - Ressource-Management System
  - Programming environment
  - Libraries and Performance analysis tools
- **Porting scientific applications**




**DEEP hardware architecture**




**DEEP software architecture**








128 CNs - Sandy Bridge




512 BNs – Intel® Xeon Phi™




Highly scalable code parts (HSCP) with regular communication patterns are off-loaded to the Booster



**DEEP cold plate**




- Combine task-based and MPI programming
  - Introduce highly parallel tasks implemented using MPI
  - Extend OmpSs to handle these tasks and data transfer between Cluster and Booster
  - Rely on MPI for scalability



- Booster Nodes do not need a host CPU
- Direct communication between BNs through EXTOLL
- Flexible assignment of groups of Cluster Nodes and Booster Nodes possible (also dynamically)
- Large blocks of code can be sent to the Booster
- Minimization of communication between CPU and accelerator
- Booster Nodes run standard Linux
- I/O is done through the Cluster

- **Scientific applications:**

- Brain simulation (EPFL)
- Space weather simulation (KULeuven)
- Climate simulation (CYI)
- Computational fluid engineering (CERFACS)
- High temperature superconductivity (CINECA)
- Seismic imaging (CGGVS)





- **Goals:**


- Evaluate the DEEP concept and its programmability
- Compare its performance with standard architectures
- Create best practice guidelines
- Propose improvements to the DEEP System

- Two coupled models:
    - ECHAM (Atmospheric)
    - **MESSy (Physicochemical interactions)**
- EMAC = ECHAM/MESSy Atmospheric Chemistry


## Structure



## Analysis



## Division



- **Project is now in Month 14**
- **Hardware status:**
  - DEEP Cluster installed at JSC (Juelich)
  - 3x Intel Xeon Phi workstations installed at JSC
  - Booster backplane design finished
  - BNC prototype ready and under test

## DEEP Cluster



## BNC (Booster Node Card) prototype




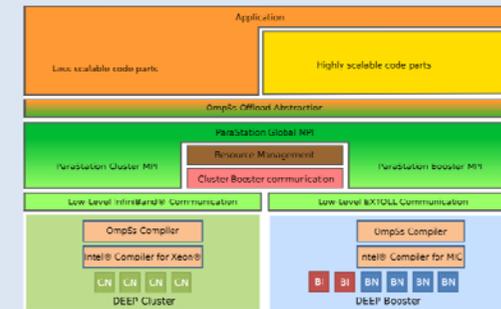
- **Software status:**

- Programming model definition completed:
  - Global MPI + OmpSs
- OmpSs runtime (Nanos++) ported to MIC
- ParaStationMPI port to MIC and EXTOLL ongoing
- Low-level Cluster-Booster protocol implemented
- Mathematical libraries under evaluation

- **Scientific Applications:**

- Structure of applications analysed
- First ansatz on application division (between Cluster and Booster) already defined
- First experiences with Intel Xeon Phi (KNC)

## DEEP software architecture



**Space Weather application**



Thank you for your attention



*Dynamical Exascale  
Entry Platform*




[www.deep-project.eu](http://www.deep-project.eu)

The research leading to these results has received funding from the European Community's Seventh Framework Programme (FP7/2007-2013) under Grant Agreement n° 287530



**BACK**



scalar + parallel  
strong scaling

$$S_N = \frac{1}{s + \frac{p}{N}}$$

weak scaling  
 $W = s + pN$




2 concurrency levels, K and N

$$S_{K,N} = \frac{1}{\frac{1-p}{Kf} + \frac{p}{N}}$$

$p_r = .50, N = 500.000, K = 10.000, f = 1:$   
 $\rightarrow S = 20.000$


$p_r = .95, N = 500.000, K = 10.000, f = 4:$   
 $\rightarrow S = 320.000$




# EXTOLL performance comparison

Technology	Implementation	Internal Clock Frequency	Link Bandwidth [Gb/s]	Latency [ $\mu$ s]	Effective maximum Bandwidth [GB/s]	Message Rate [msg/s]
EXTOLL VENTOUX	Xilinx Virtex6 FPGA	200 MHz	16 Gb/s	1.2 $\mu$ s	Up to 1.4 GB/s	~ 25 millions
EXTOLL TOURMALET	65 nm ASIC	750 MHz	120 Gb/s	0.6 $\mu$ s	Up to 10 GB/s	~ 100 millions
IB FDR*	45 nm ASIC	Unknown	56 Gb/s	0.8 $\mu$ s	Not measured	Not measured
Typical 1GE	ASIC based	e.g. 125 MHz	1 Gb/s	e.g. 40 $\mu$ s	0.11 GB/s	~0.5 millions
10GE	ASIC based	~125 to 312 MHz	12,5 Gb/s	12.5 $\mu$ s	1.2 GB/s	< 2.5 millions
Cray Gemini	90 nm ASIC	650/800 MHz	75 Gb/s	1.5 $\mu$ s	Up to 5.9 GB/s	~ 2 millions
Tianhe-1a	ASIC based	Unknown	80 Gb/s	2.5 $\mu$ s	Up to 6.34 GB/s	~ 1-3 millions
TOFU (K-Computer)	65 nm ASIC	312.5 MHz	50 Gb/s	1.5 $\mu$ s	Up to 4.76 GB/s	> 8 millions


\* Mellanox literature data



- **Low latency**
  - 1-2  $\mu$ sec
- **Large Bandwidth**
  - 32 Gbit/s
- **Clos-Topologie (fat-tree)**
  - 36-port Switches
  - Very flexible
  - But not so scalable



- **Very low latency**
  - $<1$   $\mu$ sec
- **Large Bandwidth**
  - 32 Gbit/s
- **3D-torus Topology**
  - 10-port Switches
  - Highly scalable
  - But low flexibility



- Application's main() part runs on Cluster Nodes (CN) only
- Resources managed statically or dynamically
- OmpSs acts as an abstraction layer
- Actual spawn done via Global MPI
- Spawn is a collective operation of Cluster processes
- Highly scalable code-parts (HSCP) utilise multiple Booster Nodes (BN)