

Using Fine-Grain Architectures for Weather and Climate Models

Mark Govett

Tom Henderson, Jacques Middlecoff, Jim Rosinski, Paul Madden

NOAA Earth System Research Laboratory
Boulder, Colorado

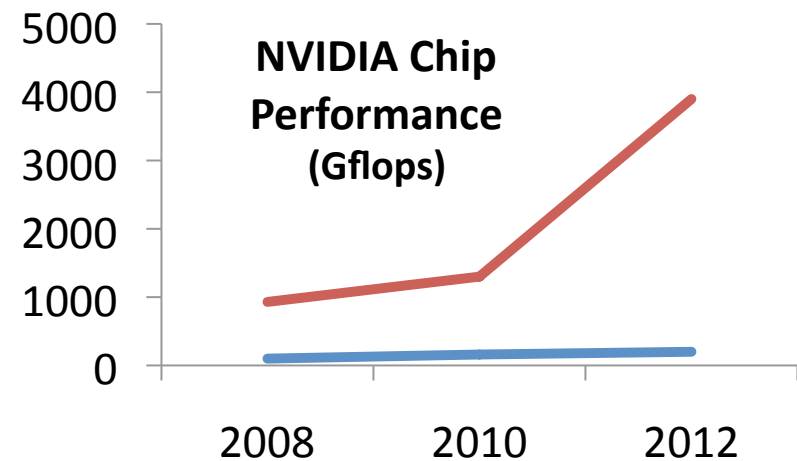
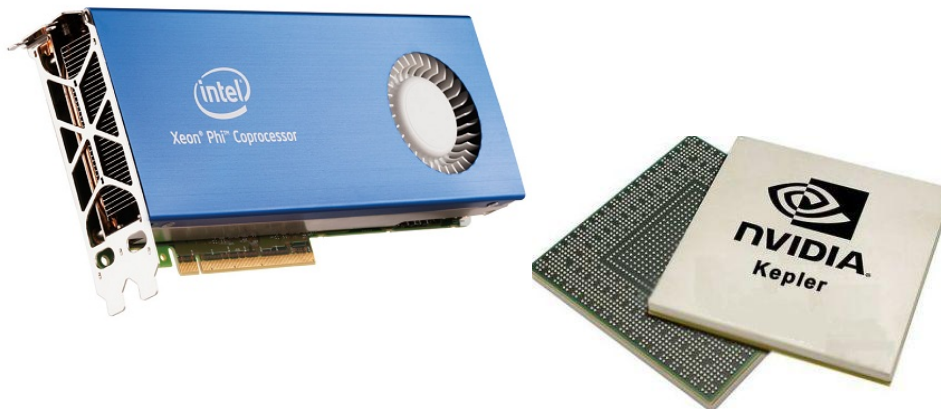
Fine-Grain Computing Hardware

- GPU

- NVIDIA Kepler 2880 cores 250w 3.9 TFlops
- AMD FireStream ~3000 cores 225w 3.23 TFlop

- Many Integrated Core

- Intel Xeon Phi 61 cores 300w 2.1 Tflops



- CPU

- Intel SandyBridge 16 cores 125w 150 GFlops

Fine-Grain Architectures

- Architectures are diverse and continue to evolve
 - Identify more parallelism in the models
- Challenge to maintain single source & performance portability

GPU Chip	Fermi (2010) C2050/70	Kepler (2012) K20x	Intel MIC (2012) Xeon Phi	AMD FireStream 9370
Cores	448	2688	61	1600
- Clock Speed	1.15 GHz	0.73 GHz	1.91 GHz	1.2 GHz
- Flops SP	1.0 TF	3.9 TF	2.1 TF	2.6 TF
Memory	4-6 GB	6 GB	8 GB	4 GB
- Bandwidth	144 GB/sec	250 GB/sec	~200 GB/sec	147 GB/sec
- Shared/L1	64 KB	64 KB		
Power	238 W	235 W	300 W	225 W
Programing Features	Cache Mem	Dynamic Parallelism	Vector units	openCL

Breakthrough Titan Performance

Jaguar Specs (2011)	
Compute Nodes	18,688
Login & I/O Nodes	256
Memory per node	16 GB
# of Opteron cores	224,256
# of NVIDIA K20 “Kepler” accelerators (2013)	N/A
Total System Memory	300 TB
Total System Peak Performance	2.3 Petaflops

Titan Specs (2012)	
Compute Nodes	18,688
Login & I/O Nodes	512
Memory per node	32 GB + 6 GB
# of Opteron cores	299,008
# of NVIDIA K20 “Kepler” accelerators (2013)	18,688
Total System Memory	710 TB
Total System Peak Performance	20+ Petaflops

- Joint agreement to give NOAA 10M node hours
 - 4 km NIM expected to be tested in 2013/14
 - 4000 GPUs needed

Fine-Grain Computing @ NOAA

- Weather Codes
 - FIM, NIM in progress
 - NMM-B, HRRR, GFS planning stages
- Climate Codes
 - NOAA GFDL Cube-Sphere in progress
- Computing Systems
 - Kepler ORNL Titan, NOAA gaea
 - MIC NFS TACC, NASA

Topics

- Maintaining Portability in Model Code
 - Single source, good design
 - Software tools & compilers
 - Validating results
- Performance & Key Optimizations
- Status and Future Work

NVIDIA GPU Architecture

Compute:

- SIMT: 32 cores, called a warp, execute in lock-step on each Streaming MultiProcessor (SM)

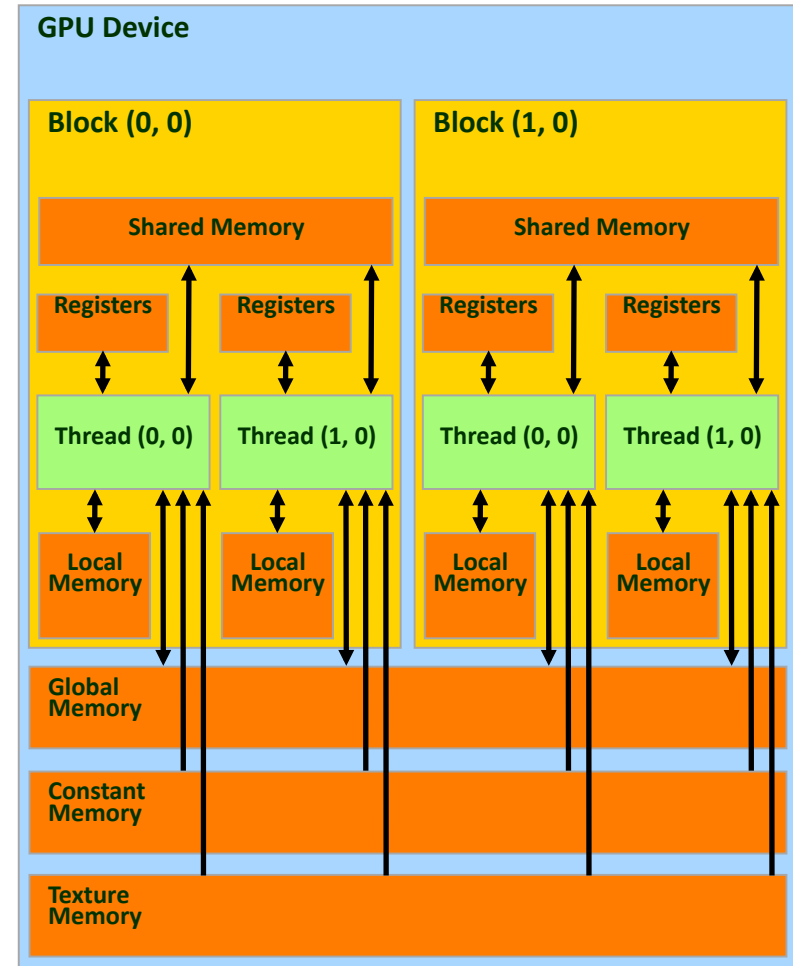
	<u>Fermi</u>	<u>Kepler</u>
SM:	32	14
cores/SM	16	192

Memory:

- 2-4 cycles to access fast memory
 - Registers, cache / shared
- Hundreds of cycles to access slow memory
 - Global, constant, texture

	<u>Fermi</u>	<u>Kepler</u>
Shared:	16/48K	16/48K
Register:	128K	256K

NVIDIA Compute / Memory



Fine-Grain Compilers

- Language-based

- | | | |
|----------------|-------------|----------------------|
| – CUDA | NVIDIA | C, C++ + extensions |
| – OpenCL | AMD, NVIDIA | C, C++ + extensions |
| – CUDA Fortran | PGI | Fortran + extensions |

- Directive-based Compilers

- F2C-ACC



F2C-ACC Compiler

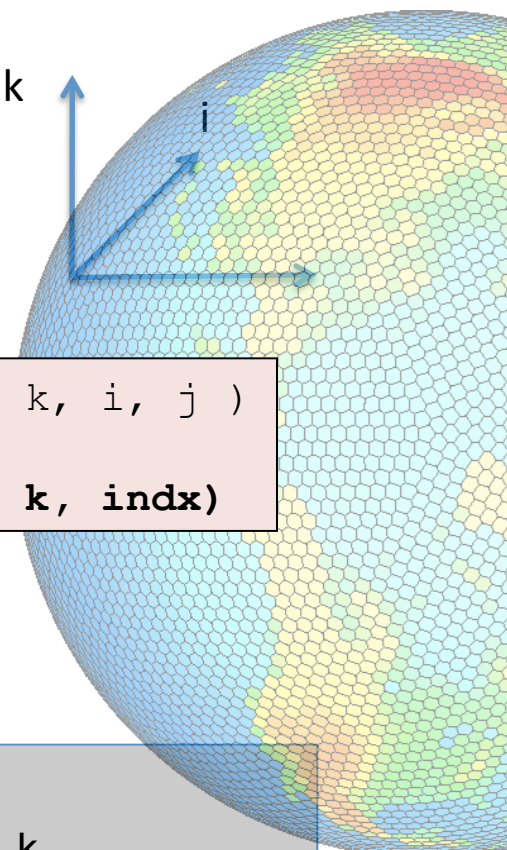
- Developed in 2009, before commercial compilers were available
 - Ten directives for parallelization: `!ACC$ <directive_name>`
- Used to parallelize NIM & FIM dynamics
 - **Focus on minimizing changes to preserve original code**
 - **Single source to preserve performance portability**
 - Run on CPU, GPU, serial, parallel
- Advanced Features
 - Variable promotion, demotion
 - Shared, local, constant memory
 - Increasing parallelism using “chunking” and “blocking”
- Used to evaluate commercial compilers
 - 2011: evaluated CAPS, PGI (Henderson)
 - 2012: shared stand-alone tests with all the vendors
 - performance, correctness, language support
 - Plan another evaluation in 2013
- 0

Programming for MIC (Intel Phi)

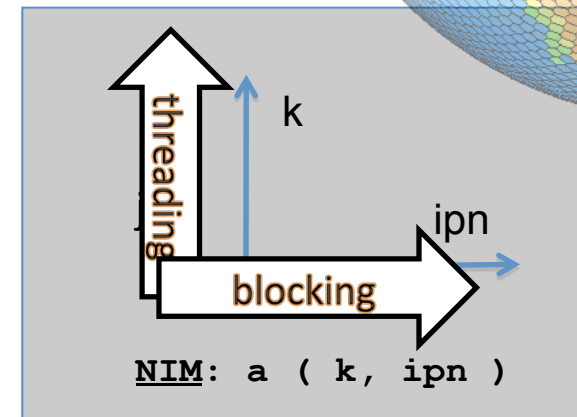
- Relies on OpenMP directives
 - Add an intel compiler flag to invoke MIC
- Two programming modes
 - Native
 - Everything runs on the coprocessor
 - NO code mods required to get it running
 - Can use multiple cores via OpenMP and/or MPI
 - Offload
 - Host offloads part of calculation to coprocessor
 - Compiler directives describe how to move data
 - Similar to GPU in handling of data transfer between 2 address spaces

NIM Parallelization for GPUs

- Uniform, hexagonal-based, icosahedral grid
- Novel indirect addressing scheme permits concise, efficient code
 - MacDonald, Middlecoff, Henderson, et al
 - No performance impact
 - Adopted by NCAR MPAS model
- Designed for fine-grain parallel in 2008
- Dynamics
 - Running on GPU, OpenMP, MIC in progress
- Physics: GFS, YSU
 - GPU, OpenMP, MIC parallelization planned
- Testing at 120, 60, 30km
 - Aqua-Planet to 300 days
- Testing at 120, 60 KM
 - real data runs



```
lat-lon a ( k, i, j )  
NIM:      a [ k, indx)
```



FIM Fine-Grain Parallelization

- Well established code
 - Designed in 2000 for CPUs
 - Near operational status
 - Running daily at 10, 15, 30 KM resolutions
 - Multi-faceted development
 - Ensembles, chemistry, ocean
- Limited ability to change the code
 - Fortran modules
 - Performance improvements
 - No degradation in clarity of code
- No change in results is ideal
 - Otherwise scientists must evaluate



FIM: a [k, indx)

- GPU
 - Blocking in horizontal
 - Threading in vertical
- OpenMP, MIC
 - Threading in horizontal
 - Vector in vertical

Parallelization for GPU & MIC

- Separate CPU routines from accelerator routines
 - Compute on GPU
 - I/O, Comms on CPU
- Parallelize the leaf routines in the call tree
- Use the accelerator style initially
 - Data lives on CPU, copied to GPU to execute each kernel
- Validate output for every step
- Make data resident on device
 - Copy data to GPU
 - Run dynamics on GPU
 - Copy data back to CPU

- FIM Dynamics
 - allocstate
 - dyn_init
 - dyn_run
 - **hybgen**
 - regrid_1d
 - » restp_1d, equilb, inflate
 - remap_1d
 - » inflate, plmadv
 - **hystat**
 - **abstart**
 - **edgvar**
 - **cnuity**
 - cnuity1, cnuity2, cnuity3, cnuity4, cnuity5
 - **trcadv**
 - trcadv1, trcadv2, trcadv3
 - **momtum**
 - del4prep, dissip, momtum1
 - compare_output

Validation of Results

- Prior to CUDA v4.2, the number of digits of accuracy was used to compare FIM / NIM results between the CPU and GPU

Variable	Ndifs	RMS (1)	RMSE	max	DIGITS
rublten	2228	0.1320490309E-03	0.2634E-09	0.3922E-05	5
rvblten	2204	0.2001348128E-03	0.6318E-09	0.2077E-04	4
exch_h	3316	0.1670498588E+02	0.8979E-05	0.8379E-05	5
hpbl	9	0.4522379124E+03	0.2688E-03	0.1532E-04	4
rqiblten	1082	0.2236843110E-09	0.7502E-17	0.6209E-07	7

- Small differences for 1 timestep can become significant when running a model over many timesteps
- CPU and GPU results are identical (Intel versus NVIDIA)
 - When the fused multiply-add instruction is turned off
 - Eliminates truncation of arithmetic operations
 - Significantly speeds parallelization
 - Exceptions:
 - Use of **power function**, and possibly other intrinsics
 - Calculations must be done on the CPU in these cases right now

Simple OpenMP / GPU Kernel

- Placement of directives are generally in the same location for GPU directives & OMP
 - OMP: Need to have sufficient work to overcome startup overhead
- Data movement is implicit
 - F2C accelerator model assumes data resides on CPU

```
!$OMP PARALLEL DO PRIVATE (k,...) SCHEDULE (runtime)
!ACC$REGION(<nvl>,<ime-ims+1>) BEGIN

!  worka(:, :) = tr3d(:, :, 1)           !F2C does not support F90 assignments

!ACC$DO PARALLEL(1)
do ipn=ips,ipe
!ACC$DO VECTOR(1)
  do k=1,nvl
    worka(k,ipn) = tr3d(k,ipn,1)
  end do
end do
!ACC$REGION END
!OMP END PARALLEL DO
```

OpenACC GPU Kernel

- Unclear how OpenACC compilers handles data movement
 - Explicitly listing each variable for each region can be tedious particularly if the region is big and complicated
- Unclear how if OpenACC compilers support Fortran 90 syntax
 - Compiler analysis could determine parallelism implicitly

```
!$ACC DATA [copyin] [copyout]
!$ACC PARALLEL [ gang ] [ worker ] [ vector ]

!  worka(:, :) = tr3d(:, :, 1)          !unclear if openACC can handle this

!$ACC LOOP [gang] [worker] [vector]
do ipn=ips,ipe
!$ACC LOOP [gang] [worker] [vector]
do k=1,nvl
    worka(k,ipn) = tr3d(k,ipn,1)
end do
end do
!$ACC END PARALLEL
```

FIM Performance – CPU, GPU & MIC (2012)

- GPU timings used F2C-ACC compiler
 - Commercial compilers performance between Global and Optimized
- Intel Xeon Phi (MIC) – SE10x Pre-production chip
 - 61 cores, 1.091 GHz, 8GB memory
 - OpenMP + MIC extensions used for parallelization
- Kepler K20x results are from the Fermi F2C-ACC optimized code

FIM Dynamics Routines	Fermi GPU F2C-ACC 1 socket	Fermi GPU Optimized 1 socket	Intel CPU SandyBridge 1 & (2) socket	Intel Xeon Phi – KNC 1 socket	Kepler GPU EARLY RESULTS
trcadv	2.07	1.28	2.10 (1.80)	1.59	0.66
cnuity	5.20	1.04	1.13 (0.76)	0.74	1.15
momtum	0.57	0.41	0.71 (0.60)	2.03	0.41
hybgen		4.13	4.09 (2.06)	3.40	3.43
TOTAL	-	8.01	8.87 (5.97)	9.89	6.60

Variable Promotion for Correctness

Example: NIM vdmintv subroutine (nz=32)

F2C V4: - promote variables using GPU global memory

```
real :: rhsu(nz,nob), rhsv(nz,nob), tgtu(nz,npp), tgtv(nz,npp)

!ACC$REGION (<nz>,<(ipe-ips+1)>,                                     &
!ACC$> <rhsu,rhsv,tgtu,tgtv:none,global,promote(1:block)> ) BEGIN
!ACC$DO PARALLEL(1)
do ipn=ips,ipe
!ACC$DO VECTOR(1,1:nz-1)
  do k=1,nz-1
    rhsu(k,1) = cs(1,ipn)*u(k,ipp1)+sn(1,ipn)*v(k,ipp1) - u(k,ipn)
    rhsu(k,2) = ...
    < Similar calculations on rhsv, tgtu, tgtv >
  enddo
  call solver( ..., rhsu, rhsv, ... )
enddo
!ACC$REGION END
```

Performance: run-time w/ global memory: 12.51 ms

nvcc will use cache by default

Optimization: Shared Memory

Example: NIM vdmintv subroutine (nz=32)

F2C V4: - use GPU shared memory for rhsu,rhsv,tgtu,tgtv

```
real :: rhsu(nz,nob), rhsv(nz,nob), tgtu(nz,npp), tgtv(nz,npp)
!ACC$DATA(<rhsu,rhsv,tgtu,tgtv:none,shared>)      !declaration required
!ACC$REGION (<nz>,<(ipe-ips+1)>,&
!ACC$> <rhsu,rhsv,tgtu,tgtv:none,shared> ) BEGIN
!ACC$DO PARALLEL(1)
do ipn=ips,ipe
!ACC$DO VECTOR(1,1:nz-1)
  do k=1,nz-1
    rhsu(k,1) = cs(1,ipn)*u(k,ipp1)+sn(1,ipn)*v(k,ipp1) - u(k,ipn)
    rhsu(k,2) = ...
    < Similar calculations on rhsv, tgtu, tgtv >
  enddo
  call solver( ..., rhsu, rhsv, ...)
enddo
!ACC$REGION END
```

Performance: run-time w/ shared memory: 7.30 ms

1.7x speedup over global memory w/ cache

Optimization: Variable Demotion

Example: FIM trcadv subroutine

F2C V4: - Demote variables + shared, local or register memory

```
!ACC$REGION(<nvl:block=2>,<ipe-ips+1>,&  
!ACC$> <s_plus,s_mnus:none,local,demote(1)>) BEGIN  
!ACC$DO PARALLEL(1)  
    do ipn=ips,ipe  
!ACC$DO VECTOR(1)  
    do k=1,nvl  
        s_plus(k) = 0.  
        s_mnus(k) = 0.  
    end do  
    do edg=1,nprox(ipn)  
!ACC$DO VECTOR(1)  
    do k=1,nvl  
        s_plus(k) = s_plus(k) - min(0., antiflx(k,edg,ipn))  
        s_mnus(k) = s_mnus(k) + max(0., antiflx(k,edg,ipn))  
    end do  
    end do
```

Performance: 1.8x faster than global memory / cache

NIM Performance (2013)

- 10K horizontal points, 96 vertical levels

NIM	Opteron	Westmere	SandyBridge	Fermi	K20x
1 node	75.4	86.8	31.3	25.0	20.7

- Very early weak scaling results
 - Opteron, K20x from Gaea-GPU + Gemini
 - SandyBridge + Infiniband
 - Communications issue noted on Gaea
 - Contradicts Jaguar test showing Gemini was faster

Number of Nodes		Opteron + Gemini Total (comms)		SB + Infiniband Total (comms)		K20x + Gemini Total (Comms)	
4	(120 km)	79.8	(10.6)	34.0	(5.9)	46	(21.0) = 25.0
16	(60 km)	95.9	(21.5)	43.1	(11.8)	68.7	(37.1) = 31.6
64	(30 km)	83.7	(28.3)	45.9	(14.4)	60.1	(29.6) = 30.5

Conclusion & Next Steps

- Demonstrated single source, performance portability for GPU, MIC, CPU
- Bit reproducibility between GPU & Intel CPU
- Early performance on K20x is encouraging
- Next Steps
 - NIM
 - MIC parallelization of dynamics
 - MIC & GPU parallelization of physics
 - Optimize scaling from 100s to 1000s of GPUs after Titan becomes available in May
 - FIM
 - MIC & GPU parallelization of physics
 - Run 60 member ensembles on Titan