

On the Potential of Concurrency for Scaling W&C Applications in the ExaSCALE Era

Rupert Ford, STFC Ukri

Reinhard Budich, MPI-M

Julia Duras, DKRZ

Why this Presentation?

- Attempt in ESiWACE to tackle the topic „Concurrency“
- People had radiation and OBGC in mind, and may be others
- More activity in that direction as well

- Somehow we got stuck: Why?
- Report* had to be written, the base for this talk

* Available on request from Reinhard

Concurrency: Which One?

- ➔ Rob Pike is a Canadian programmer and author: Go, Bell Labs, now at google.
- ➔ From Wikipedia:
„According to Rob Pike, concurrency is the composition of independently executing computations, [2] and concurrency is not parallelism:
 - ➔ **Concurrency** is about dealing with lots of things at once but **parallelism** is about doing lots of things at once.
 - ➔ **Concurrency** is about structure, **parallelism** is about execution,
 - ➔ Concurrency provides a way to structure a solution to solve a problem that may be (but is not necessarily) parallelizable.[3]“

(2) „Go Concurrency Patterns“. <https://talks.golang.org/2012/concurrency.slide#6> talks.golang.org. Retrieved 2021-04-08.

(3) "Go Concurrency Patterns". <https://talks.golang.org/2012/waza.slide#8> talks.golang.org. Retrieved 2021-04-08.

Corresponding Author: Reinhard Budich, MPI-M, reinhard.budich@mpimet.mpg.de



Concurrency: Which One?

➔ Definitions for C&W computing problems / type of concurrency

➔ a) ensembles / external

➔ b) model components / coupling

➔ c) comm. and comp. / I/O servers

➔ d) data structures / nodes or threads

➔ Here

➔ e) functions / intra-component ➔ **Functional Concurrency**

This Talk

- Intro
- Ansatz for a definition of functional concurrency
- Findings about experiences in the community
- The potential of functional concurrency
- Why it is so hard
- Potential Solutions & Outlook

Functional Concurrency: Ansatz for a Definition

➔ The report describes it as:

„Model components can often also be further split into finer-grain components with some of these components having the potential to run in parallel with each other. This intra-component functional parallelism is currently typically not exploited by model components ...“

➔ Examples include, but are not limited to

➔ atm - rad

➔ oce - obgc

➔ lnd - rivers

➔ ...

Findings about Experiences in the Community

→ Examples

→ MPAS

→ COSMO

→ LFRic

→ ICON


MPAS: Courtesy Rich Loft

(Inter-component task concurrency in ESM, not considered here)

- ➔ Intra-component task concurrency
 - ➔ Look at cost, refactoring ease, expected acceleration, potential for software reuse
 - ➔ Rad and Ind are on CPU, rest on GPU
 - ➔ cpu as „Co-processor“

MPAS: Courtesy Rich Loft

MPAS Call Structure



MPAS: Courtesy Rich Loft, Remarks

➔ Radiation

➔ To GPU

- ➔ Rad RRTMG was not yet ported to GPU, as it is now the case.
- ➔ Rad was 1/3 of the LoC of MPAS-A: Major task!
- ➔ Many lookup-tables problematic for memory access
- ➔ Not clear if spreading the task to many GPUs would scale and perform

➔ On CPU

- ➔ Would rad@cpu be fast enough for the rest? Also depending on ratio between
 - ➔ rad and dycore time step!
 - ➔ Performance of cpus and gpus

MPAS: Courtesy Rich Loft, Conclusions

➡ It depends ...

➡ Please refer to the report: Tremendous work!

➡ ... on implementation of the concurrency in the component(s)

COSMO - Preliminaries: Courtesy Carlos Osuna

- ➔ NWP code of MeteoSuisse parallelized with
 - ➔ OpenACC for Driver, parametrization, assimilation
 - ➔ GridTools for the dycore
- ➔ Observation
 - ➔ Strong scalability saturates at 40 k grid points per GPU device
 - ➔ GPUs can not be fully occupied, adding GPUs does not help
- ➔ Idea
 - ➔ Run dycore and rad (every 60th time-step) in parallel as GPU streams, OpenMP and OpenACC applied (async queues)
- ➔ Hypothesis
 - ➔ This way GPUs can be utilized in a way nearer to saturation

COSMO - Results: Courtesy Carlos Osuna

→ Results

- No bit-identical results, but differences should be small from theoretical considerations - full meteorological evaluation was not performed, though.
- Quote from report: „Even though we observed the desired task parallelism in the radiation scheme, we could not measure a significant improvement in performance corresponding to it.,,

→ Reasons

- (Too) Large kernels in the rad scheme use all resources, and block other kernels
- Controlling launch times for kernels would help here, but was not possible from OpenACC

→ My take

- Smart strategy, but disappointing results
- Strategy dropped for the time being
- **Dependent on implementation of the concurrency in the component(s)**

LFRic: Courtesy Rupert Ford


- ➔ W&C atmosphere code of MetOffice to come
 - ➔ Not yet complete, not yet with the desirable performance, but useable for prototypical case studies
 - ➔ Concurrency for components and data treated as usual, also with DSL (for functional and task low level Conc. on loop level)
- ➔ Approaches: High level concurrency
 - ➔ for slow physics, within a time-step
 - ➔ processes are independent of each other within a timesteps
 - ➔ (Fast physics are not considered and executed sequentially)
 - ➔ for radiation, across time-step
 - ➔ Use physics consideration
 - ➔ Rad status is not (very) dependent on parameters in the same time-step, information can be used from earlier time-steps

LFRic, slow physics: Courtesy Rupert Ford


➔ Time-step in LFRIC as DAG:



➔ Add timing information, TS weighted




LFRic - Slow physics expanded: Courtesy Rupert Ford




LFRic slow physics

Computation at rad time step !

LFRic - Slow physics expanded: Courtesy Rupert Ford



LFRic slow physics




LFRic slow physics TS weighted

Different possibilities for concurrency under consideration, depending on complexity of components

Computation at rad time step !

LFRic, concurrent rad: Courtesy Rupert Ford

➔ Approach:



➔ Results:

➔ It depends


- ➔ on physical assumptions, and cost benefit considerations by scientists
- ➔ on the exact balance of costs for the different components, depending on their complexity and the frequency with which they are computed
- ➔ on implementation of the concurrency in the component(s)

ICON: C. Many Iconists - Marco Giorgetta, Leonidas Linardakis, Julia Duras ...

- ➔ Complex coupled weather and climate model of DKRZ, DWD, KIT and MPI-M
 - ➔ Atm, Lnd, Oce
 - ➔ Provides ample concurrency in all dimensions
 - ➔ Coupling of the atmosphere and the ocean via the coupler library YAC
 - ➔ I/O-Server library CDI-PIO, enabling concurrent execution of I/O
 - ➔ Data parallelism within model components employing OpenMP and MPI
 - ➔ functional concurrency in ICON-A, ICON-O
- ➔ OBGC (HAMOCC)
 - ➔ provides a concurrent ocean case
 - ➔ Described elsewhere
 - ➔ It runs well (T. Ilyina, pers. comm.)
- ➔ Look at rad in ICON-A here

ICON-A: C. Many Iconists - Marco Giorgetta, Julia Duras ...

- ➔ Process splitting for slow physics in ICON-A (M. Giorgetta)
 - ➔ Coarse res / long time-steps can lead to higher tracer concentrations than the reference (neg. concentrations)
-> physically not reasonable, numerically unstable
 - ➔ High res / short TS configurations allow to reduce number of processes (rad, diff, cloud μ -physics)
 - ➔ Sequential coupling currently
 - ➔ But considered for complete parallel coupling (some remarks next slide)
 - ➔ So far, partial parallel coupling was considered



Remarks on Complete Parallel Coupling in ICON-A

- ➔ Parallel coupling and computing of more than one "slow physics" process and dynamics not yet been done
 - ➔ It is a tremendous work, for which sustainability is not guaranteed
 - ➔ Parallel coupling scheme
 - ➔ Two synchronizations necessary:
 - ➔ Combine with some monitoring and regularization scheme to prevent rare exceptions
- ➔ Distribute the common input state to all processes computed in parallel
- ➔ Collect and process their output within the same time step

Remarks on Complete Parallel Coupling in ICON-A

- ➔ Parallel coupling and computing of more than one "slow physics" process and dynamics not yet been done
 - ➔ It is a tremendous work, for which sustainability is not guaranteed
 - ➔ Parallel coupling scheme
 - ➔ Two synchronizations necessary:
 - ➔ Combine with some monitoring and regularization scheme to prevent rare exceptions
- ➔ Scientifically parallel coupling scheme could be developed independently:
 - ➔ Once an accepted coupling scheme for a certain experiment is found, working with acceptable time steps (i.e. computing costs), such a parallel computing scheme could be developed and employed
 - ➔ Sustainability of the implementation questionable
- ➔ Make input data and the tasks available
- ➔ Get work done as fast as possible
- ➔ Collect results from all processes
- ➔ Combining them ready for the next phase of computation

Remarks on Complete Parallel Coupling in ICON-A

- ➔ Parallel coupling and computing of more than one "slow physics" process and dynamics not yet been done
 - ➔ It is a tremendous work, for which sustainability is not guaranteed
 - ➔ Parallel coupling scheme
 - ➔ Two synchronizations necessary:
 - ➔ Combine with some monitoring and regularization scheme to prevent rare exceptions
- ➔ Scientifically parallel coupling scheme could be developed independently:
 - ➔ Once an accepted coupling scheme for a certain experiment is found, working with acceptable time steps (i.e. computing costs), such a parallel computing scheme could be developed and employed
 - ➔ Sustainability of the implementation questionable
- ➔ With a completely parallelized computation of processes, distribution of resources to the tasks can be embedded in the model, invisible to the user
- ➔ Dynamical process scheduling might be an option

Partial parallel coupling in ICON-A

→ Concurrent computation of radiation process

→ Use of longer rad time-step:

$$\Delta t_{\text{rad}} = k \cdot \Delta t, k > 1$$



→ The larger k, the more the radiation will deviate from „true“ solution - affordable

→ Concurrent PSrad was developed

→ Own processes

→ different coupling scheme, different causality, different results

Sequential, „traditional“




Concurrent, PSrad_{conc}

Partial parallel coupling in ICON-A

→ Results:

- 25 quantities checked in 9 yr AMIP-like runs
- Deviation is in the range of a few percent
- Radiation accuracy can be increased by a reduced time-step length, without computational overhead
- Computational gain ok, though not terrific:
 - integration of the whole model gets faster by about the amount needed for the radiation computed sequentially
- Very much depending on resource distribution between $\rho\text{Srad}_{\text{conc}}$ and „rest of the model“
- Practical complications and intense testing necessary <- Very costly, **not affordable in our case!**
- And: ρSrad was dropped for other reasons



ICON: Conclusion

- ➔ ICON-A:
 - ➔ More work to be done
 - ➔ Re-structuring and -engineering in full swing (DestinE, Warmworld, others)
 - ➔ For functional concurrency, extra project funding applied for
- ➔ *Dependent on implementation of the concurrency in the component(s)*
- ➔ Seems to have worked for ICON-O/HAMOCC (OBGC)

The Potential of Functional Concurrency

→ Dependent on implementation of the concurrency in the component(s)

The Potential of Functional Concurrency

- ➔ Dependent on implementation of the concurrency in the component(s)
- ➔ No generic approach visible, so case- and model-dependent
- ➔ AND: It is depending on intense scientific evaluation & cooperation
- ➔ But the only potential candidate for more parallelization
- ➔ Huge cost involved

Why it is so hard

- ➔ Radiation is special case since $\Delta t_{\text{rad}} \gg \Delta t_{\text{rest}}$ slow physics
- ➔ Other (sub-time-step) components need to be **modified scientifically** before they can be algorithmically and numerically altered in order to achieve concurrency ▀ **High effort, high risk**
- ➔ However, having the flexibility in the software to be able to test whether it is beneficial to run sub-components concurrently or not and being able to choose a different solution depending on the required configuration and underlying architecture is something that would be beneficial for scientists.
- ➔ In the report different scenarios concerning properties of the sub-components and the resource availability are discussed which would benefit from such flexibility
 - ➔ Flexibility in ordering of SC, or choosing different resolution, rates or precision might also help - discussed in the report
- ➔ Conclusion of this report: A well defined interface (of which a subroutine boundary with all data being passed by argument is considered sufficient) is necessary in order to be able to maintain a more flexibility for concurrency; two of four models have this already
- ➔ Depending on machine architecture, a mix of MPI&OpenMP&OpenACC seems most appropriate - also hard!
- ➔ Buffer space needs to be allocated to store information from former time steps, also depending on the scientific set-up
- ➔ Complexity seems to grow overwhelmingly

Potential Solutions

- ➔ Ideas/recommendations from the report for further investigation:
- ➔ Coupling systems and frameworks show potential
 - ➔ ESMF
 - ➔ DSL
 - ➔ Combination of library code, code-generation and -modification (CCPP@NCAR)
 - ➔ Lib between component and sub-component as a SC-wrapper

Outlook

- ➔ Technology changes
 - ➔ NVIDIA'S Multi-Instance GPU for better load balancing
- ➔ Further problems:
 - ➔ Proliferating programming codes
 - ➔ Heterogeneity
- ➔ Big projects plus pressure for more parallelism might help ...

- ➔ We need to use the concept of functional concurrency to structure the problems in a way to parallelize execution of tailor-made codes of certain physical processes

Thanks!

➔ Questions



ESiWACE2 has received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement No 823988