

# PSyclone in Met Office: Evolution and revolution

Iva Kavcic, Met Office, UK &

Rupert Ford, **Andrew Porter**, Sergi Siso (STFC, UK); Joerg Henrichs (BOM, AU); LFRic Team, Marine Systems Team (Met Office, UK)....

7<sup>th</sup> ENES HPC Workshop, 9-11 May 2022



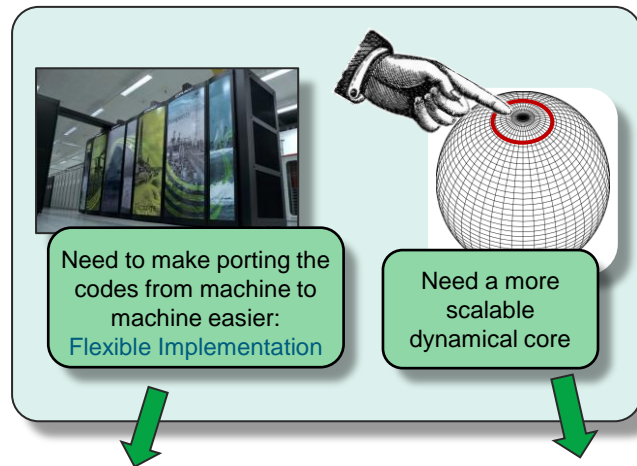
# Overview

**PSyclone: Motivation and intro**

**Revolution: LFRic API**

**Evolution: NEMO API**

# Motivation (GungHo project)



**New Infrastructure  
– LFRic**



**New Dynamical Core  
– GungHo**

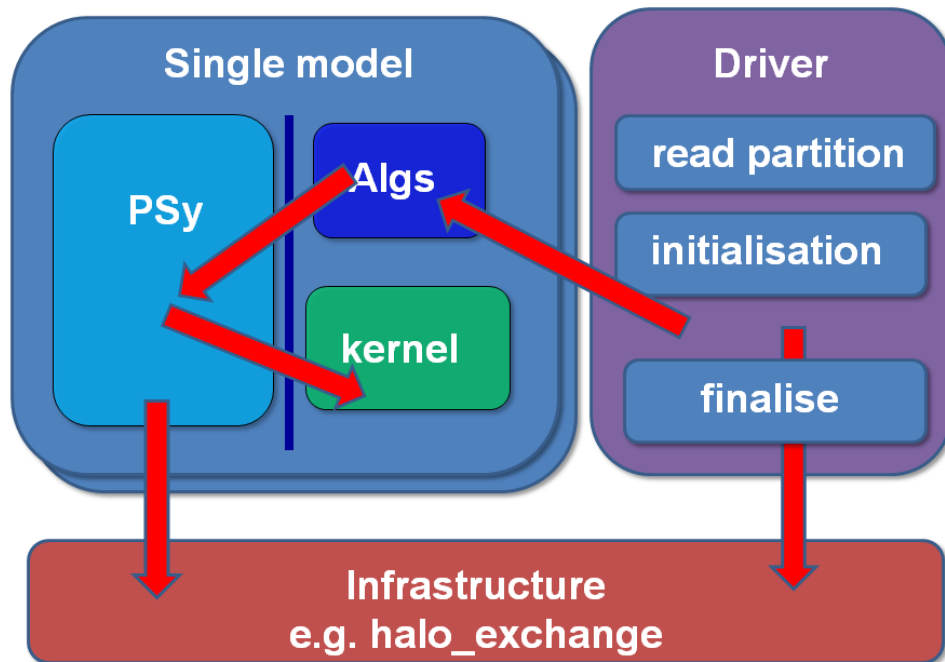


- Increased resolution → exascale computation
- Future architectures - MPI? OpenMP? Accelerators? GPUs? ARM? ...?
- Scientific code can be  $\sim 10^6$  lines (of Fortran).
- Complex parallel code + Complex parallel architectures + Complex compilers = Complex optimisation space => **Single-source optimised code unlikely to be possible**
- 3P's : **Performance, Portability** and **Productivity**
  - Maintainable high performance software
  - **Single-source science code**
  - Performance portability

# Separation of Concerns: Science and code optimisation

## PSyKAI

- **Parallel-Systems:** Computational Science, applies optimisations – *generated code*
- **Kernels:** Natural Science, operations on (columns of) data points
- **Algorithms:** Natural Science, operations on whole **data structures** (e.g. fields)





PSyclone 2.2.0

BSD 3-clause

<https://github.com/stfc/PSyclone>

<https://psyclone.readthedocs.io>

```
> pip install psyclone
```

- A **domain-specific compiler** for **embedded DSL(s)**
  - Configurable: FD/FV NEMO, GOcean, FE LFRic
  - Currently Fortran -> Fortran/OpenCL
  - Supports distributed- and shared-memory parallelism
  - Supports **code generation** and **code transformation**
- A **tool** for use by **HPC experts**
  - Hard to beat a human (debatable)
  - Work round limitations/bugs
  - **Optimisations** encoded as a 'recipe' rather than baked into the scientific source code
  - Different recipes for different computer architectures
  - Enables **scriptable, whole-code optimisation**

# Fparser

## Pure Python Fortran parser:

- Supports Fortran 2003 + some 2008
- Open source BSD3 licence
- Developed on GitHub
- Can fully parse UM, LFRic and NEMO source
- Work-in-progress to parse IFS source
- Used by PSyclone, Stylist, Loki

<https://github.com/stfc/fparser>

<https://fparser.readthedocs.io/>

```
> pip install fparser
```

```
PROGRAM copy_stencil
  IMPLICIT NONE
  INTEGER, PARAMETER :: n = 10, np1 = 11
  INTEGER :: i, j, k
  REAL, DIMENSION(np1, n, n) :: out, in
  DO k = 1, n
    DO j = 1, n
      DO i = 1, n
        out(i, j, k) = in(i + 1, j, k)
      
```



```
child type = <class 'fparser.two.Fortran2003.Execution_Part'>
  child type = <class 'fparser.two.Fortran2003.Block_Nonlabel_Do_Construct'>
    child type = <class 'fparser.two.Fortran2003.Nonlabel_Do_Stmt'>
      child type = <class 'str'> 'DO'
      child type = <class 'fparser.two.Fortran2003.Loop_Control'>
        child type = <class 'NoneType'>
        child type = <class 'tuple'>
        child type = <class 'NoneType'>
      child type = <class 'fparser.two.Fortran2003.Block_Nonlabel_Do_Construct'>
        child type = <class 'fparser.two.Fortran2003.Nonlabel_Do_Stmt'>
          child type = <class 'str'> 'DO'
          child type = <class 'fparser.two.Fortran2003.Loop_Control'>
            child type = <class 'NoneType'>
            child type = <class 'tuple'>
            child type = <class 'NoneType'>
          child type = <class 'fparser.two.Fortran2003.Block_Nonlabel_Do_Construct'>
            child type = <class 'fparser.two.Fortran2003.Nonlabel_Do_Stmt'>
              child type = <class 'str'> 'DO'
```

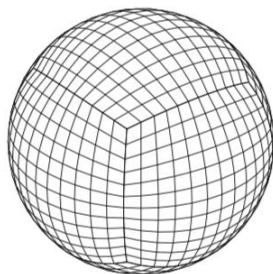
# Revolution: LFRic API\*

- Process code written in a DSL embedded in Fortran
- PSyKAl code structure
- Generated PSy layer

\* *Also GOcean API (FD, 2D structured grid)*



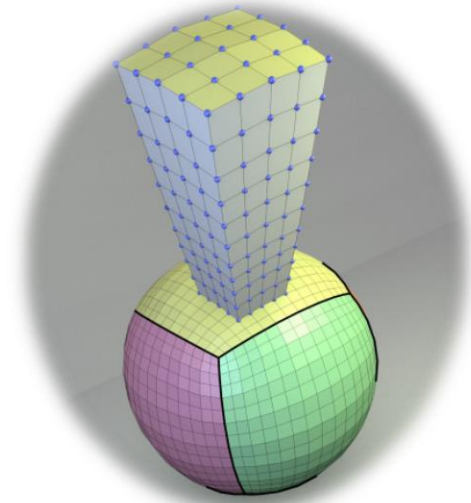
# LFRic system & GungHo dynamical core



- **Mixed Finite Element Method**
- **Horizontally unstructured**, vertically structured quasi-uniform mesh
- **Generated optimisations**

*2D cubed-sphere mesh extruded into 3D levels*

→ **Mesh layout needs to be stored**



- Horizontal adjacency lost
- **Vertically** adjacent cells **contiguous** in memory → operate on **columns** of data

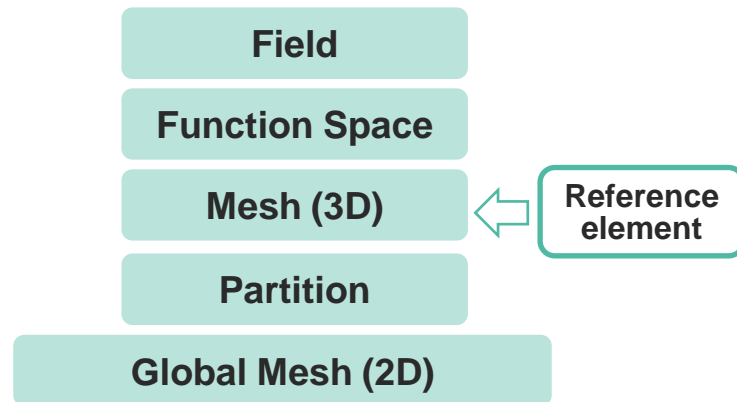


# LFRic infrastructure (Object-orientated Fortran 2003)

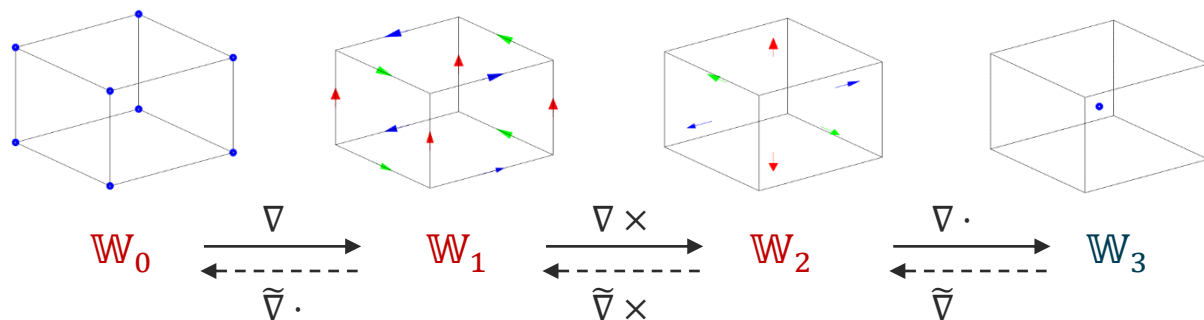
## ➤ LFRic: data classes

- Storing prognostic and diagnostic quantities: **field**;
  - Mathematical operations: **operator** (FEM matrices/FS mappings), **scalar** (global reductions).
- **Data classes** for supporting objects, e.g. mesh, reference element, function space
- **Challenge: Compiler support for OO F2003** is mixed → “compiler league table” (communication with vendors)

## *LFRic infrastructure: Hierarchy of objects*



# Mixed Finite Element Method (**continuous** + **discontinuous**)

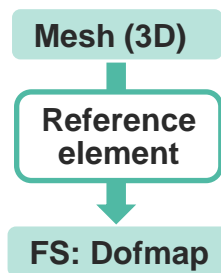


Lowest order ( $k=0$ ) reference element

## Shared DoFs

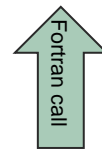
(degrees of freedom):  
 $W_0$  (all),  $W_1$  (tangential components) &  $W_2$  (normal components)

No shared DoFs:  $W_3$



LFRic infrastructure **supports** two main modes of updating fields:

- Looping over **cell columns**,
- Looping over **DoFs**.



PSyclone **generates calls** to LFRic infrastructure support (PSy layer).

# DSL embedded in Fortran: Kernel metadata (how to **access** and **update data**; *kernel code written by scientists, Fortran 90*)

```
module rrho_kernel_mod
```

```
...
```

```
type, public, extends(kernel_type) :: rrho_kernel_type
```

```
private
```

```
type(arg_type) :: meta_args(4) = (/
```

```
arg_type(GH_FIELD, GH_REAL, GH_READWRITE, W3), &
```

```
arg_type(GH_FIELD, GH_REAL, GH_READ, &
```

```
ANY_DISCONTINUOUS_SPACE_1), &
```

```
arg_type(GH_FIELD, GH_REAL, GH_INC, W2), &
```

```
arg_type(GH_FIELD, GH_REAL, GH_READ, ANY_SPACE_1), &
```

```
/)
```

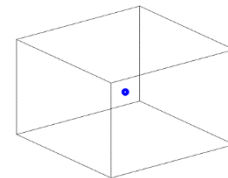
```
integer :: operates_on = CELL_COLUMN
```

```
contains
```

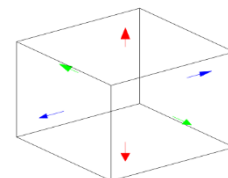
```
procedure, nopass :: rrho_code
```

```
end type
```

```
...
```



**Discontinuous**  
function spaces:  
no shared DoFs



**Continuous**  
function spaces:  
shared DoFs

# DSL embedded in Fortran: Algorithm code (operations on whole fields; *written by scientists, Fortran 2003*)

```
module rhs_rho_alg_mod
...
subroutine on_the_fly_rhs_alg(rhs, state, ref_state, ... )
  use rrho_kernel_mod,          only: rrho_kernel_type
  use matrix_vector_kernel_mod, only: matrix_vector_kernel_type
  implicit none
  type(field_type), target, intent(in) :: state(bundle_size)
  type(field_type), target, intent(inout) :: rhs(bundle_size)
  ...
  call invoke( name = "compute_rhs_rho",
               &
               rtheta_kernel_type( rhs_tmp(igh_t), rho_ref, u, u_ref ), &
               matrix_vector_kernel_type( rhs(igh_t), theta, mm_rho ), &
               inc_X_plus_bY( rhs(igh_t), tau_t_dt, rhs_tmp(igh_t) ) )
  ...
end subroutine on_the_fly_rhs_alg
end module rhs_rho_alg_mod
```

**Global fields:** data layout hidden

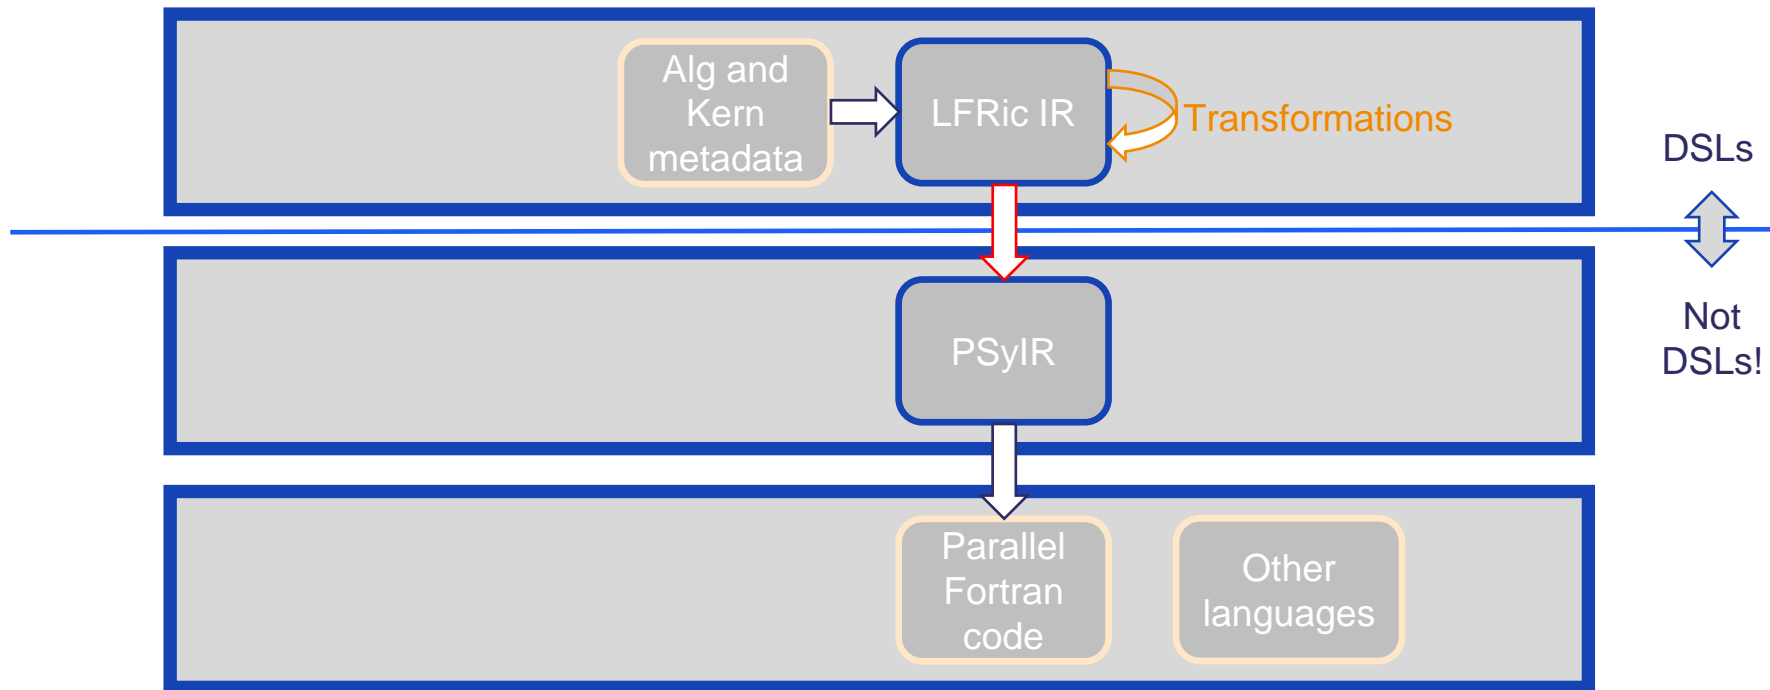
## Kernel (LFRic)

- PSy-layer loop over columns of cells

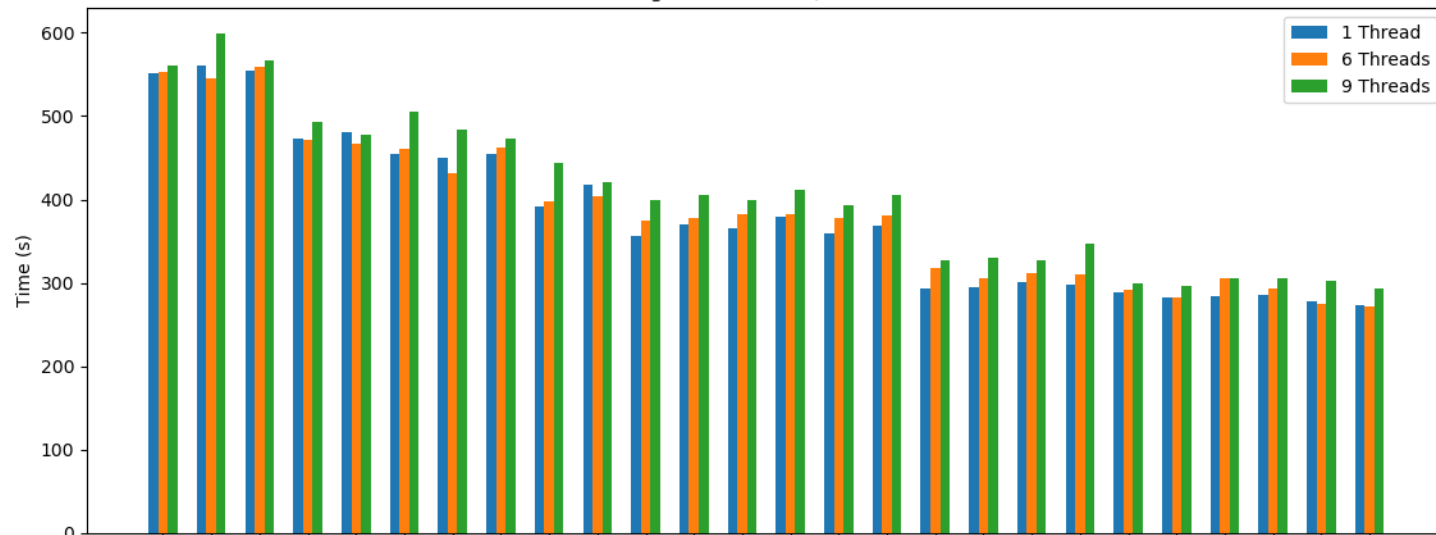
## Built-in (PSyclone)

- PSy-layer loop over all field **DoFs** (arithmetic operations)

# LFRic DSL PSy Layer

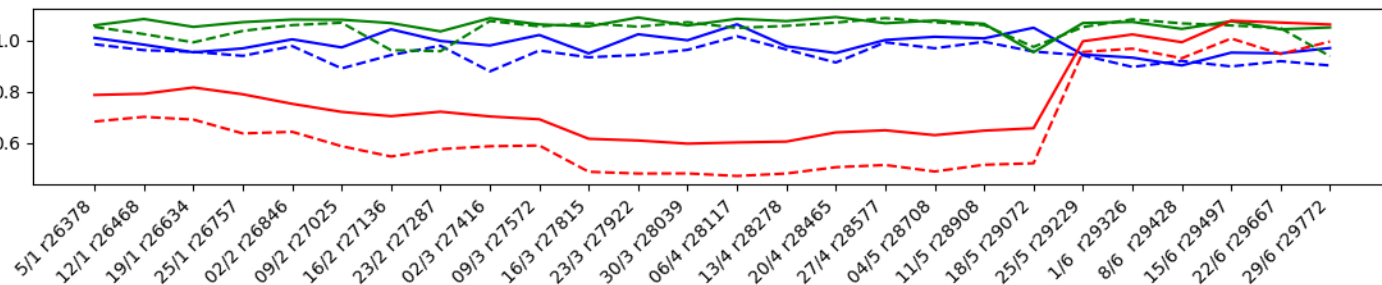


Gungho model cost, 12 nodes



## GungHo dynamical core, Held-Suarez test case:

- C192 MG ( $6 \times 192^2$  columns  $\approx 50$  km hor)
- L30 DCMIP vert
- $dt = 1200$  s (SI)
- Local volume:  $32 \times 16$  (1 OMP),  $64 \times 48$  (6 OMP),  $96 \times 48$  (9 OMP)
- Cray XC40, 2x18-core Broadwell



Courtesy of **Tom Melvin**,  
Dynamics Research, MO

# Evolution: NEMO API\*

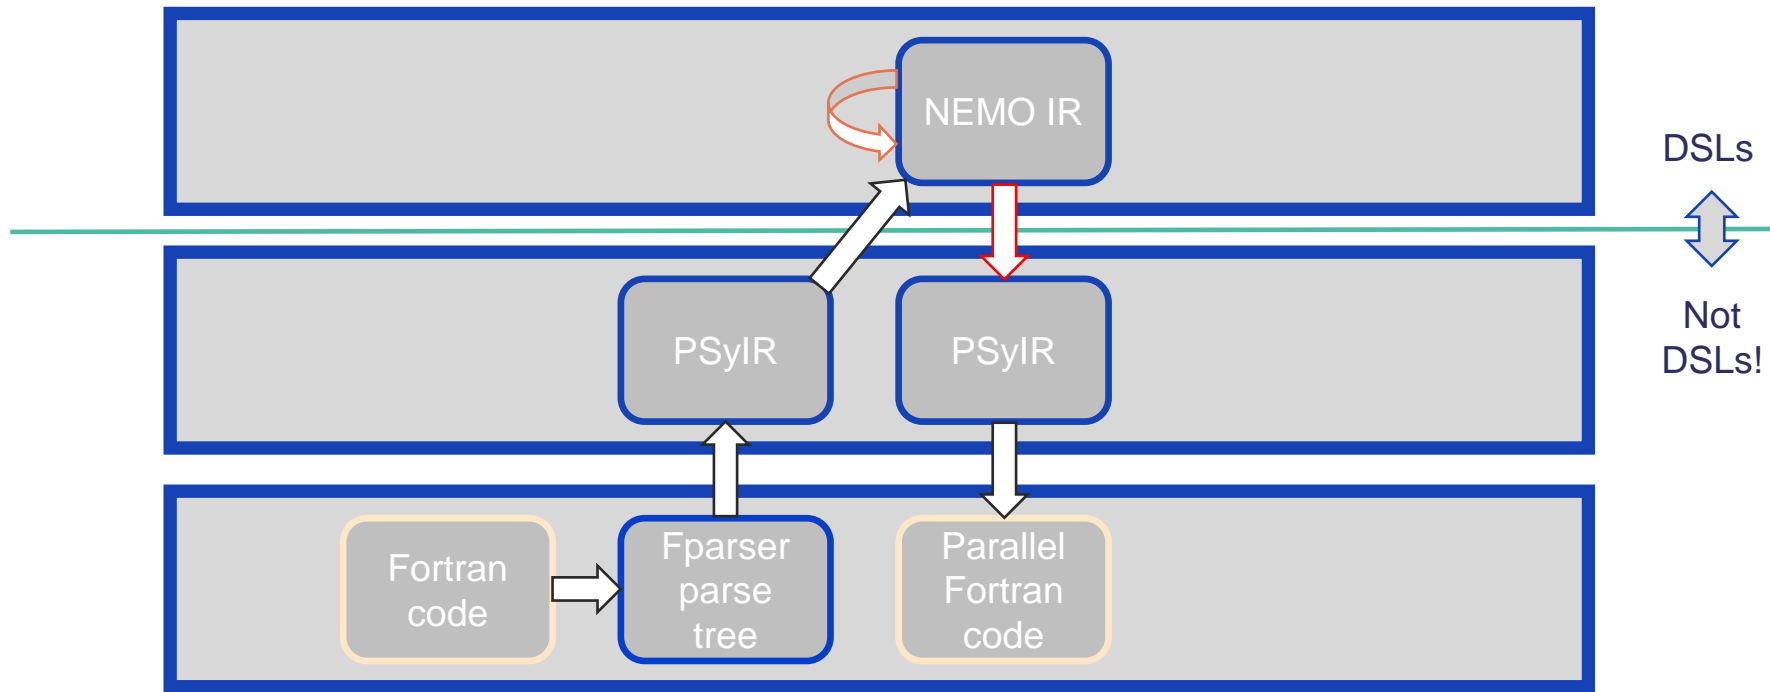
- Process **existing code** that follows strict coding conventions
- Recognise certain code structures and construct higher-level Internal Representation
- Transformations applied to this IR

\* *In development for **NEMO** (plus associated models, e.g.  $SI^3$ , MEDUSA). Also applied to **ROMS**.*



# NEMO DSL

Construct high-level representation of existing source code:



# NEMO transformation example

(PSyclone/examples/nemo/eg2)

Original NEMO  
source code  
(`tral_ldf_iso`  
routine):

*Source code  
treated as a  
“manually written”  
PSy layer with all  
kernels in-lined  
(NEMO coding  
conventions)*

```
!                                     ! =====
DO jn = 1, kjpt                      ! tracer loop
!                                     ! =====
!
!!-----
!!   I - masked horizontal derivative
!!-----
bug.... why (x,,,)?  (1,jpj,:) and (jpi,1,:) should be sufficient....
  zdit (1,,,)= 0._wp      ;      zdit (jpi,,,)= 0._wp
  zdjt (1,,,)= 0._wp      ;      zdjt (jpi,,,)= 0._wp
!!end

! Horizontal tracer gradient
DO jk = 1, jpkml
  DO jj = 1, jpjml
    DO ji = 1, jpiml      ! vector opt.
      zdit(ji,jj,jk) = ( ptb(ji+1,jj ,jk,jn) - ptb(ji,jj,jk,jn) ) * umask(ji,jj,jk)
      zdjt(ji,jj,jk) = ( ptb(ji ,jj+1,jk,jn) - ptb(ji,jj,jk,jn) ) * vmask(ji,jj,jk)
    END DO
  END DO
END DO
IF( ln_zps ) THEN                ! botton and surface ocean correction of the horizontal gradient
  DO jj = 1, jpjml              ! bottom correction (partial bottom cell)
```

## PSyIR

constructed  
by PSyclone:

```
Literal[value: 0. , Scalar<REAL, wp. <Scalar<INTEGER, UNDEFINED>, unresolved>>]
4: Loop[type='levels', field_space='None', it_space='None']
  Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
  Reference[name:'jpk1']
  Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
  Schedule[]
    0: Loop[type='lat', field_space='None', it_space='None']
      Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
      Reference[name:'jpl1']
      Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
      Schedule[]
        0: Loop[type='lon', field_space='None', it_space='None']
          Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
          Reference[name:'fs_jpl1']
          Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
          Schedule[]
            0: InlinedKern[]
              Schedule[]
                0: Assignment[]
                  ArrayReference[name:'zdit']
                  Reference[name:'ji']
                  Reference[name:'jj']
                  Reference[name:'jk']
                  BinaryOperation[operator:'MUL']
```

Hands-on on Binder:

<https://github.com/stfc/PSyclone>

# Transformed PSyIR:

```
Literal[value: 0. , Scalar<REAL, wp: <Scalar<INTEGER, UNDEFINED>, unresolved>]
4: Directive[OMP parallel do]
  Schedule[]
    0: Loop[type='levels', field_space='None', it_space='None']
      Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
      Reference[name:'jpkm1']
      Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
      Schedule[]
        0: Loop[type='lat', field_space='None', it_space='None']
          Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
          Reference[name:'jpjm1']
          Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
          Schedule[]
            0: Loop[type='lon', field_space='None', it_space='None']
              Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
              Reference[name:'fs_jpim1']
              Literal[value:'1', Scalar<INTEGER, UNDEFINED>]
              Schedule[]
                0: InlinedKern[]
                  Schedule[]
                    0: Assignment[]
```

## Generated Fortran with OpenMP directives added

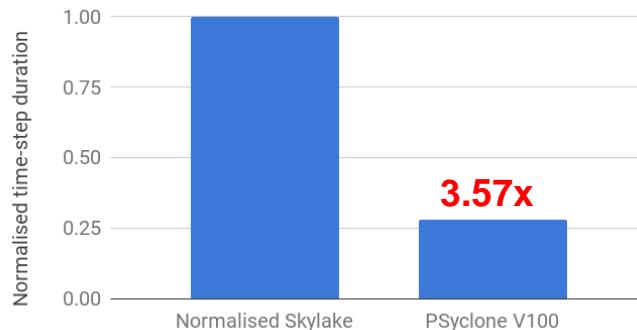
```
DO jn = 1, kjpt
  zdit(1, :, :) = 0._wp
  zdit(jpi, :, :) = 0._wp
  zdjt(1, :, :) = 0._wp
  zdit(jpt, :, :) = 0._wp
  !$OMP parallel do default(shared), private(ji,jj,jk), schedule(static)
  DO jk = 1, jpkm1
    DO jj = 1, jpjm1
      DO ji = 1, fs_jpim1
        zdit(ji, jj, jk) = (ptb(ji + 1, jj, jk, jn) - ptb(ji, jj, jk, jn)) &
          * umask(ji, jj, jk)
        zdjt(ji, jj, jk) = (ptb(ji, jj + 1, jk, jn) - ptb(ji, jj, jk, jn)) &
          * vmask(ji, jj, jk)
      END DO
    END DO
  END DO
  !$OMP end parallel do
```

*All loops  
parallelised  
over vertical  
levels using  
OpenMP*

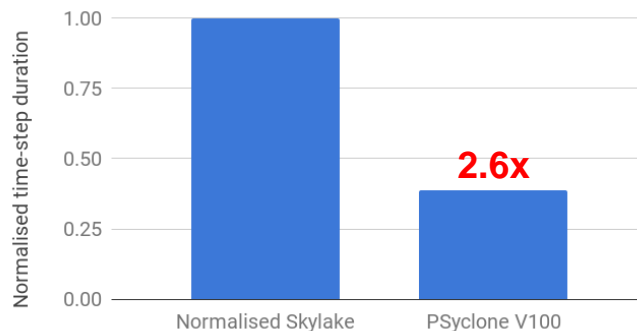
# NEMO performance (May 2022)

(Courtesy of **Chris Dearden**, STFC Hartree Centre)

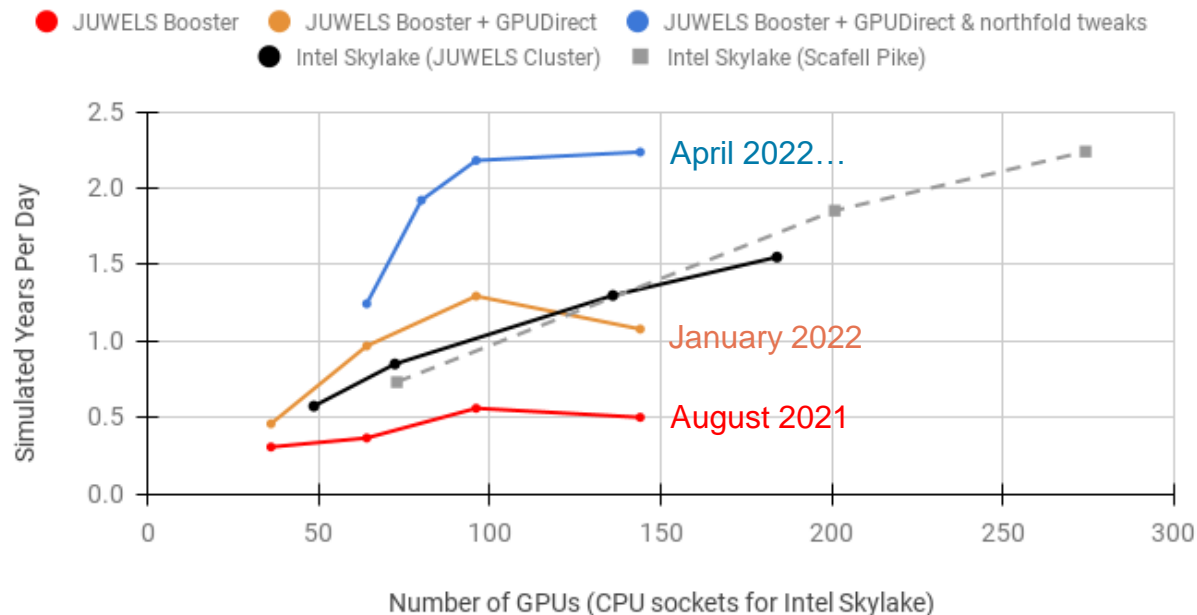
NEMO Ocean, ORCA1



NEMO Ocean + SI3, ORCA1



## Evolution of NEMO ORCA12 GPU+MPI performance



## Summary (revolution meets evolution)

- Separation of concerns → flexible optimisations
- Flexible optimisations + API knowledge → performance improvements
- “Knowledge sharing” between PSyclone APIs in MO
  - NEMO API → Run LFRic on GPU (OpenACC in LFRic API)
  - PSyclone in LFRic build system → Incorporate PSyclone in building NEMO
  - Management of PSyclone in MO



# Questions?

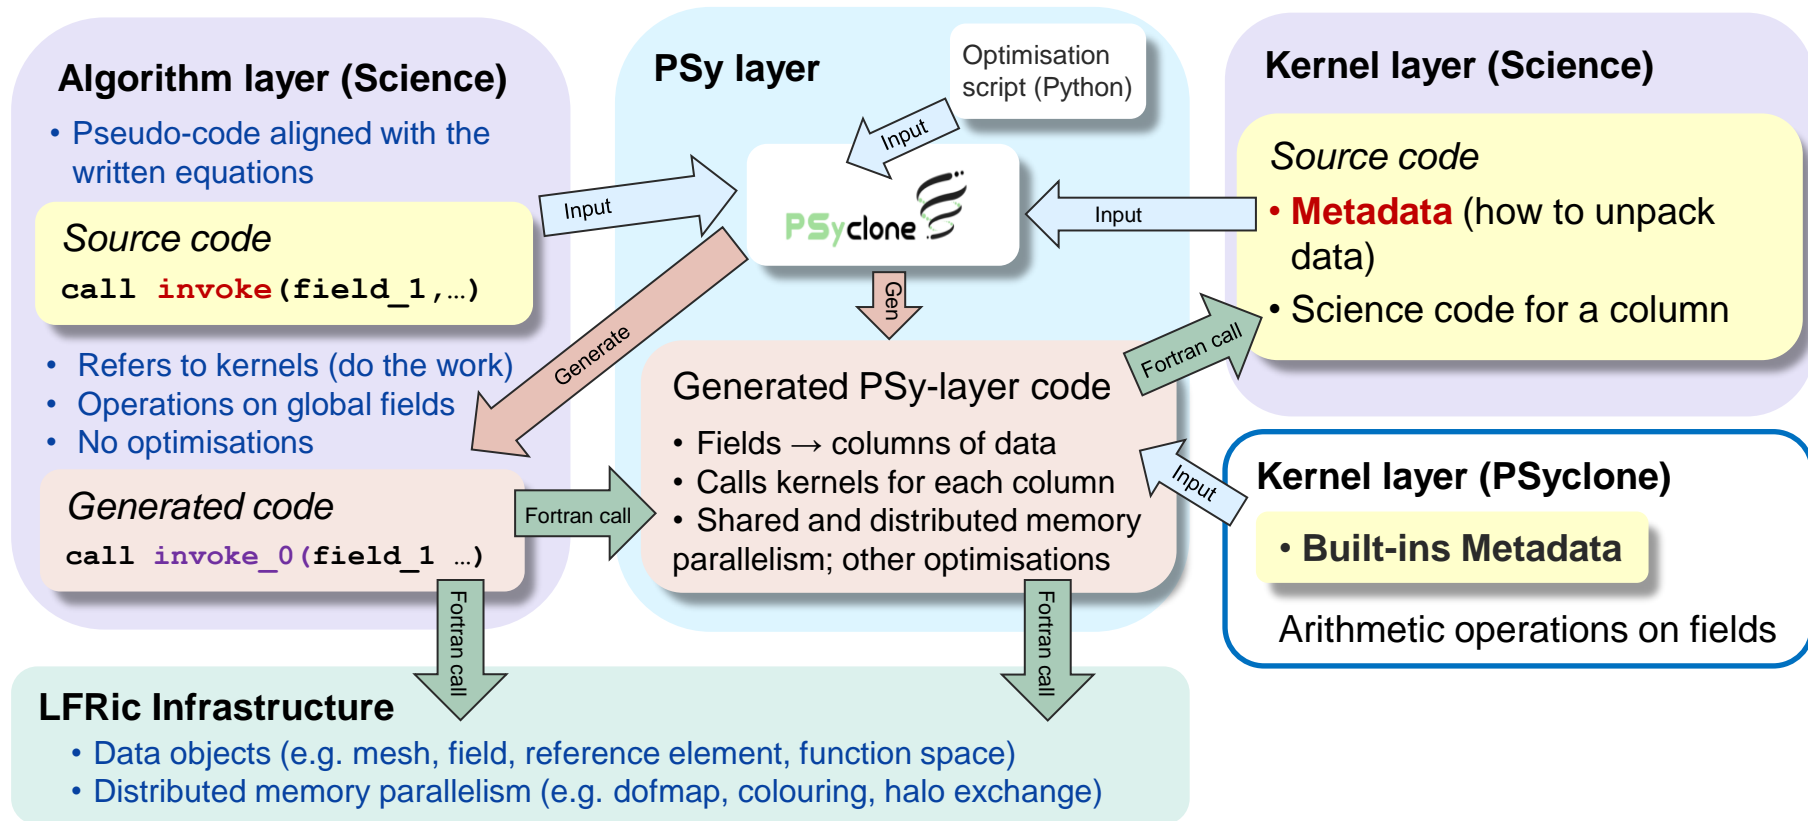


*ESiWACE2 has received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement No 823988*

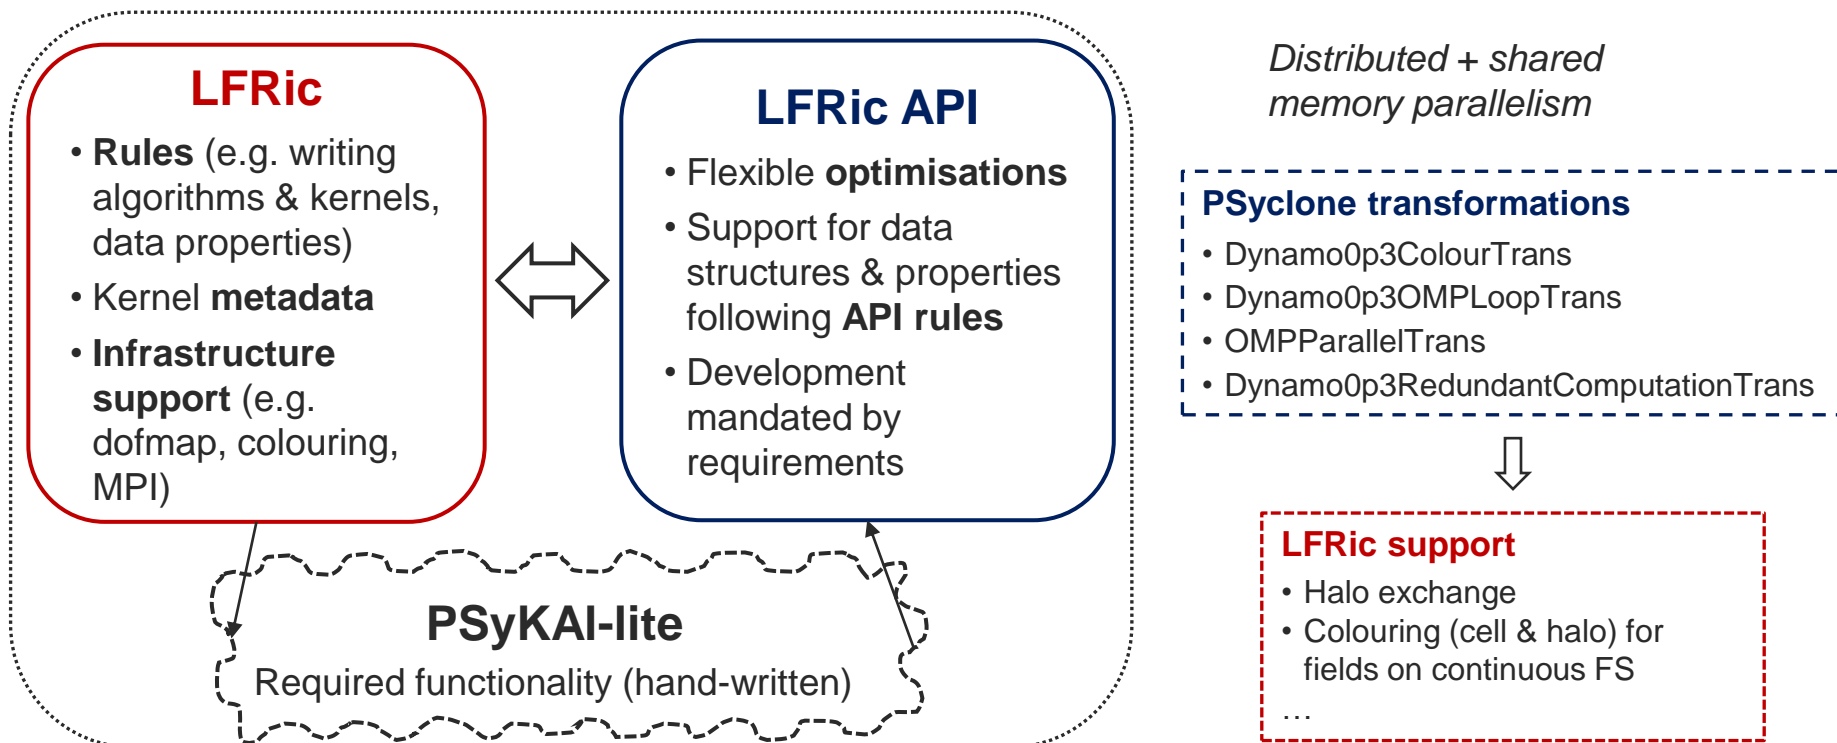
# Links and references

- PSyclone and fparser
  - <https://github.com/stfc/PSyclone>
  - <https://psyclone.readthedocs.io>
  - <https://github.com/stfc/fparser>
  - <https://fparser.readthedocs.io>
  - Hands-on on Binder: <https://github.com/stfc/PSyclone>
- LFRic: <https://code.metoffice.gov.uk/trac/lfric/wiki>
- PSyclone in LFRic: <https://code.metoffice.gov.uk/trac/lfric/wiki/PSycloneTool>
- NEMO Coding Conventions. 2013. URL: [https://forge.ipsl.jussieu.fr/nemo/attachment/wiki/Documentation/NEMO\\_coding.conv\\_v3.pdf](https://forge.ipsl.jussieu.fr/nemo/attachment/wiki/Documentation/NEMO_coding.conv_v3.pdf).
- Stylist: <https://github.com/MetOffice/stylist>
- Adams et al. (2019), [\*LFRic: Meeting the challenges of scalability and performance portability in Weather and Climate models\*](#), Journal of Parallel and Distributed Computing, 132, 383-396

# PSyKAI Infrastructure in LFRic: Parallel-Systems, Kernels, Algorithms



# Rules of engagement: LFRic $\longleftrightarrow$ PSyclone LFRic API



# NEMO transformation script

```
def trans(psy):  
    ''' Transform a specific Schedule by making all loops  
    over levels OpenMP parallel.  
  
    :param psy: the object holding all information on the PSy layer \\  
                to be modified.  
    :type psy: :py:class:`psyclone.psyGen.PSy`  
  
    :returns: the transformed PSy object  
    :rtype: :py:class:`psyclone.psyGen.PSy`  
  
    ...  
    from psyclone.psyGen import TransInfo  
    from psyclone.nemo import NemoKern  
    # Get the Schedule of the target routine  
    sched = psy.invokes.get('tra_ldf_iso').schedule  
    # Get the transformation we will apply  
    ompt = TransInfo().get_trans_name('OMPParallelLoopTrans')  
    # Apply it to each loop over levels containing a kernel  
    for loop in sched.loops():  
        kernels = loop.walk(NemoKern)  
        if kernels and loop.loop_type == "levels":  
            sched, _ = ompt.apply(loop)  
    # Return the modified psy object  
    return psy
```

Parallelises all  
loops over  
vertical levels  
using OpenMP