

# Preparing IFS for HPC accelerators via source-to-source translation

M. Lange, B. Reuter, O. Marsden


[michael.lange@ecmwf.int](mailto:michael.lange@ecmwf.int)

European Centre for Medium-Range Weather Forecasts



# Towards km-scale global models

- Seasonal global simulation at 1.4 km horizontal resolution (INCITE) <sup>1</sup>
- Destination Earth: EC initiative to develop highly accurate digital twin of Earth
- EuroHPC: Access to large-scale computing platforms via Destination Earth
- Many (Pre-)Exascale-class systems feature large GPU partitions




<sup>1</sup> Nils P. Wedi et al. "A Baseline for Global Weather and Climate Simulations at 1 km Resolution". In: *Journal of Advances in Modeling Earth Systems* 12.11 (2020), e2020MS002192. DOI: <https://doi.org/10.1029/2020MS002192>.

# Hybrid 2024 - Preparing IFS for HPC accelerators

**Hybrid 2024:** Internal project to adapt IFS to accelerator-based HPC architectures

## Accelerator-enabled multi-architecture IFS

- Sustainable technical development of accelerator capabilities alongside scientific development
- Incremental adaptation of model components to different accelerators and programming models
- Dedicated build-modes to target different programming models / compiler toolchains
- Continuous testing on available hardware




# Hybrid 2024 - Preparing IFS for HPC accelerators

**The challenge:** Adapt IFS to new programming paradigms without harming CPU performance

## Multiple programming models in a single code

- Use library APIs to hide technical complexities
- Use flexible data structures to deal with complex memory hierarchies
- Increasingly flexible control flow with different parallelisation paradigms
- Source-to-source translation for device-specific optimisation of compute kernels




# ECTTRANS - An open-source Spectral Transform library

- Extracted Spectral Transform routines and packaged as stand-alone project ([🔗ecmwf-ifs/ectrans](https://github.com/ecmwf-ifs/ectrans))
- No dependency on ifs libraries; minimal requirements stand-alone FIAT library ([🔗ecmwf-ifs/fiat](https://github.com/ecmwf-ifs/fiat))

## GPU-enabled version available

- GPU-enabled (NVIDIA) branch is now in the repo
- Fortran + OpenACC + vendor-specific libraries (FFTs rely on CUFFT, GEMM calls on cuBLAS)
- Good scaling behaviour; uses CUDA-aware MPI (GPU-to-GPU MPI communications via NVLink)
- **WARNING: This is work in progress!**  
**Not all features are supported on GPU yet!**

Time per inv+dir timestep at TC01279




# Flexible data structures for complex memory hierarchies

**FIELD API:** Initial adaptation to allow GPU-offload via OpenACC / OpenMP

- **Object-oriented** data structures to encapsulate memory placement of field arrays
- **Separation of concerns:** Explicitly manage data offload to accelerator devices
- Enables restructuring of control flow to adapt to alternative execution modes

## Atlas – A library for NWP and climate modelling


- Modern C++ library with Fortran interfaces <sup>2</sup>
- Data structures for numerical algorithms:
  - Remapping and interpolation
  - Gradient, divergence, laplacian
  - Spherical Harmonics transforms
  - **Increasing accelerator-awareness!**



<sup>2</sup>Willem Deconinck et al. "Atlas : A library for numerical weather prediction and climate modelling". In: *Computer Physics Communications* 220 (2017), pp. 188–204. ISSN: 0010-4655. DOI: <https://doi.org/10.1016/j.cpc.2017.07.006>.

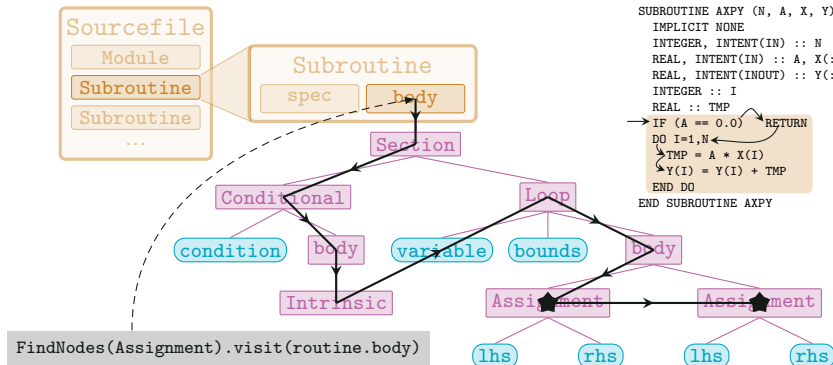
# Source-to-source translation of physical parameterisations

- **Loki:** Programmable source-to-source translation tool written in Python
- Bulk transformation of source tree at compile time (“complex preprocessor”)
- Enable parallel development of science code and performance engineering
- Explore alternative programming model and device-specific code structures



## Under the hood: Traversing expression trees

- Loki's internal representation (IR) is based on two levels of trees
  - **Custom** data structures for control flow (purple)
  - Enhanced **Symbolic** nodes for expression trees (blue)

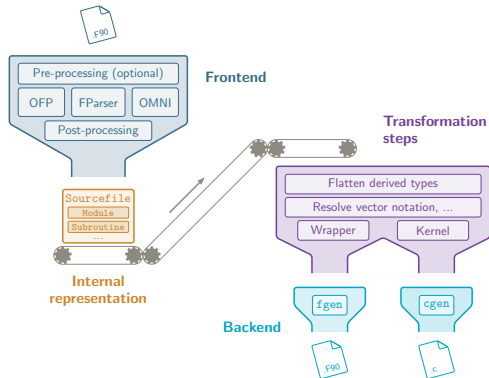




# Loki – freely programmable source-to-source translation



- Two-level internal representation: control flow and expression trees
- Visitors to search and transform code
- Growing library of utilities to aid encoding transformations

```
def remove_loops(self, routine, dimension):  
    """  
    Replace vector loop with it's own body  
    """  
  
    map = {}  
    for loop in FindNodes(Loop).visit(routine.body):  
        if loop.variable == dimension:  
            map[loop] = loop.body  
  
    routine.body = Transformer(map).visit(routine.body)
```



# Single-column layout in physical parameterisations

- No data dependencies between columns:  
Lots of parallelism! <sup>3</sup>
- Scientific kernel can be developed and tested for a single column
- Columns are stored in a block layout traversed with a high OpenMP loop
- Resulting memory layout packs sub-blocks of arrays into cache-friendly chunks (NPROMA)



<sup>3</sup>Valentin Clement et al. "The CLAW DSL: Abstractions for Performance Portable Weather and Climate Models". In: *Proceedings of the Platform for Advanced Scientific Computing Conference. PASC '18*. 2018. ISBN: 9781450358910. DOI: 10.1145/3218176.3218226.


# IFS: Memory data layout and parallelisation

```
!$omp parallel loop  
do ibl=1,nblocks  
  call kernel(var1(:, :, ibl), var2(:, :, ibl), ...)  
end do
```




```
SUBROUTINE KERNEL(nlon,nlev,var1,var2,...)  
  real :: var1(nlon, nlev)  
  real :: var2(nlon)
```

```
do j=1, klon  
  var1(j, 1) = var2(j)  
end do  
  
do k=2, nlev  
  do j=1, klon  
    var1(j,k) = var1(j,k-1) + <update>  
  end do  
end do  
END SUBROUTINE
```

CPU-socket



CPU-socket



# IFS: Memory data layout and parallelisation

```
!$acc parallel loop gang  
do ibl=1, nblocks  
  call kernel(var1(:,ibl), var2(:,ibl), ...)  
end do
```

```
SUBROUTINE KERNEL(nlon,nlev,var1,var2,...)  
  real :: var1(nlon, nlev)  
  real :: var2(nlon)  
!$acc routine vector
```



```
!$acc loop vector
```

```
do j=1, klon  
  var1(j, 1) = var2(j)
```

```
!$acc loop seq
```


```
do k=2, nlev  
  var1(j,k) = var1(j,k-1) + <update>
```

```
end do
```

```
end do
```


```
END SUBROUTINE
```

Host memory



GPU-device

Device memory



# Automatically mapping memory-blocked CPU code to GPUs

## CLOUDSC (<https://github.com/ecmwf-ifs/dwarf-p-cloudsc>)

- Standalone version of cloud microphysics
- Representative parallelisation, memory layout
- Challenging to optimise (high register pressure)


## Evaluation of GPU code transformations

- Pure compute throughput comparison (no offload)
- Comparison of different Fortran GPU loop strategies
- HIP-SYCL: Manual translation from C variant


## Per-chip/socket performance comparison

- Chip-to-chip compute throughput comparison
- PCIe offload cost ignored for single kernel
- FPGA results are memory-bandwidth limited! <sup>4</sup>

CLOUDSC: Comparison of GPU code layouts



CLOUDSC: HPC architecture Comparison



<sup>4</sup>James Stanley Targett et al. "Systematically migrating an operational microphysics parameterisation to FPGA technology". In: *29th IEEE FCCM* 2021, Orlando, FL, USA, May 9-12, 2021. IEEE, 2021, pp. 69–77. DOI: 10.1109/FCCM51124.2021.00016.

# Preparing IFS for HPC accelerators via source-to-source translation

## Hybrid 2024: Internal project to prepare IFS for HPC accelerators

- Develop GPU and accelerator capabilities alongside scientific development
- Incremental adaptation of IFS components to different accelerator types
- Use [library APIs](#), [GPU-enabled data structures](#) and [source-to-source translation](#) to separate technical concerns from scientific user code

## Loki: Freely programmable source-to-source translation

- In-house programmable [source-to-source translation tool](#), under active development to encode (instead of commit) code changes
- Bulk transformation of source tree at compile time (“complex preprocessor”)
- Flexible and powerful, but does require explicit encoding of recipes and thus expertise

# Thank you! Any questions?

✉ [michael.lange@ecmwf.int](mailto:michael.lange@ecmwf.int)  
🐦 [MLange805](https://twitter.com/MLange805)