



Schweizerische Eidgenossenschaft
Confédération suisse
Confederazione Svizzera
Confederaziun svizra

Swiss Confederation

Federal Department of Home Affairs FDHA
Federal Office of Meteorology and Climatology **MeteoSwiss**

Semantic data access to gridded weather data based on zarr

7th ENES HPC Workshop

May 11 2022

Gabriela Aznar Siguan, Néstor Tarín, Marcus Schulte, Manuel Moser, Philipp Falke, Mathieu Bernard, Christoph Spirig, Mathieu Schaer, Mark Liniger


MOTIVATION




Postprocessing and downscaling as a typical data science application

refining numerical weather predictions and climate simulations using large archives of past data and statistics/machine learning

COSMO @ 2km




DEM @ 50m



observations



wind speed @ 50m




Machine learning

manipulating ~ 100 TB + 35 TB/yr heterogeneous data

Daniele Nerini et al. EMS 2021



Unified access (API) to feed data pipelines





GrideFix database in a nutshell

- GrideFix is a **cloud-native** database: it can be deployed close to the data and accessed from everywhere.





GrideFix database in a nutshell

- GrideFix is a **cloud-native** database: it can be deployed close to the data and accessed from everywhere.
- It handles **multi-dimensional arrays** using compressed binary chunks following the **Zarr** specification.





GrideFix database in a nutshell

- GrideFix is a **cloud-native** database: it can be deployed close to the data and accessed from everywhere.
- It handles **multi-dimensional arrays** using compressed binary chunks following the **Zarr** specification.
- GrideFix offers a **unified** and **semantic data access**. Users can query data by tags and standard variable names, independently of the data source and original file structure.



Amazon S3





GrideFix database in a nutshell

- GrideFix is a **cloud-native** database: it can be deployed close to the data and accessed from everywhere.
- It handles **multi-dimensional arrays** using compressed binary chunks following the **Zarr** specification.
- GrideFix offers a **unified** and **semantic data access**. Users can query data by tags and standard variable names, independently of the data source and original file structure.
- Timestamped data can be imported and deleted efficiently, it is optimized to implement a **rolling archive**.



Amazon S3





GrideFix database in a nutshell

- GrideFix is a **cloud-native** database: it can be deployed close to the data and accessed from everywhere.
- It handles **multi-dimensional arrays** using compressed binary chunks following the **Zarr** specification.
- GrideFix offers a **unified** and **semantic data access**. Users can query data by tags and standard variable names, independently of the data source and original file structure.
- Timestamped data can be imported and deleted efficiently, it is optimized to implement a **rolling archive**.
- A **Python API** provides the requested data as **dask** arrays, allowing to slice and manipulate the data along any arbitrary dimension with **parallel** computations.






GrideFix database in a nutshell

- GrideFix is a **cloud-native** database: it can be deployed close to the data and accessed from everywhere.
- It handles **multi-dimensional arrays** using compressed binary chunks following the **Zarr** specification.
- GrideFix offers a **unified** and **semantic data access**. Users can query data by tags and standard variable names, independently of the data source and original file structure.
- Timestamped data can be imported and deleted efficiently, it is optimized to implement a **rolling archive**.
- A **Python API** provides the requested data as **dask** arrays, allowing to slice and manipulate the data along any arbitrary dimension with **parallel** computations.
- **Web services** exposing the API provide cloud-native **import**, **export** (via OPeNDAP) and **catalogue** services to facilitate language independent event-driven architectures.






High level architecture



HOW DOES GRIDFIX WORK?




Data model






Labeled data to ease search





Data format defined by GrideFix arrays with Zarr





Data format defined by GrideFix arrays with Zarr

```
COSMO-1
├── 30a23917-11e7-4950-82b2-e8690184ac55
│   ├── .zattrs
│   ├── coordinates
│   │   ├── latitude
│   │   ├── longitude
│   │   ├── time
│   │   └── vertical
│   ├── 2020-01-01T00:00:00
│   │   ├── air_temperature
│   │   ├── eastward_wind
│   │   ├── mass_fraction_of_cloud_liquid_water_in_air
│   │   └── northward_wind
│   ├── 2020-01-01T06:00:00
│   │   ├── air_temperature
│   │   └── eastward_wind
│   └── 9ed433bb-ebfc-4f00-8f31-f0441d8c703f
│       ├── .zattrs
│       ├── coordinates
│       │   ├── latitude
│       │   ├── longitude
│       │   ├── time
│       │   └── vertical
│       ├── air_temperature
│       ├── eastward_wind
│       ├── mass_fraction_of_cloud_liquid_water_in_air
│       └── northward_wind
└── 482505bf-f2c3-43bf-867c-e34f6156bb21
...

```

← Source group
← Array group identified by uuid
← Tags go here
← Array coordinates group

← Time axis group
← Variable array

← Array uuid
← Tags go here

← Static data

Advantages:

- rolling archive efficient
- time-dependent attributes

Data not changing often.

Advantages:


- more chunking freedom



GrideFix arrays are combined into datasets

```
COSMO-1
├── 30a23917-11e7-4950-82b2-e8690184ac55
│   ├── .zattrs
│   ├── coordinates
│   │   ├── latitude
│   │   ├── longitude
│   │   ├── time
│   │   └── vertical
│   ├── 2020-01-01T00:00:00
│   │   ├── air_temperature
│   │   ├── eastward_wind
│   │   ├── mass_fraction_of_cloud_liquid_water_in_air
│   │   └── northward_wind
│   ├── 2020-01-01T06:00:00
│   │   ├── air_temperature
│   │   └── eastward_wind
│   └── 9ed433bb-ebfc-4f00-8f31-f0441d8c703f
│       ├── .zattrs
│       ├── coordinates
│       │   ├── latitude
│       │   ├── longitude
│       │   ├── time
│       │   └── vertical
│       ├── air_temperature
│       ├── eastward_wind
│       ├── mass_fraction_of_cloud_liquid_water_in_air
│       └── northward_wind
└── 482505bf-f2c3-43bf-867c-e34f6156bb21
...


```



Arrays with **same source and tags** are concatenated along the major time axis into a **dataset**.



Metadata store to improve performance



Metadata store:


- Fast catalog queries
- Fast identification of queried data to retrieve
- Data consistency using transactional operations
- No need to consolidate metadata

Chunk store:

- Store available from everywhere



Zarr metadata store data model




PostgreSQL

Zarr nodes



APIs





Example: start using the catalogue web service

language independent
access

cloud-native → scalable

get fast information
about available data

```
In [1]: import requests
DISCOVER_URL = 'http://discover.meteoswiss.ch/gridefix/api'
```

```
In [3]: # 1. get sources which contain variables of interest for surface values
variables = 'air_temperature,eastward_wind,northward_wind,wind_speed_of_gust'
tags = 'surface'

sources = requests.get(f'{DISCOVER_URL}/sources?tags={tags}')
sources.json()
```

```
Out[3]: ['COSMO-2E', 'COSMO-1E', 'ECMWF_IFS']
```

```
In [4]: # 2. choose source and see details
source = 'COSMO-2E'

source_details = requests.get(f'{DISCOVER_URL}/sources/{source}')
source_details.json()
```

```
Out[4]: [{'bounding_box': [5.304744720458984,
45.49021530151367,
10.898223876953125,
48.099998474121094],
'crs': 'EPSG:4326',
'description': 'MeteoSwiss ensemble forecasting system',
'level_type': 'depth',
'name': 'COSMO-2E',
'source_crs': 'EPSG:4326',
'tags': ['numerical weather prediction',
'operational',
'forecast',
'ensemble',
'depth'],
'timestamps': ['2022-03-09T12:00:00',
'2022-03-09T18:00:00',
'2022-03-10T00:00:00',
'2022-03-10T06:00:00',
'2022-03-10T12:00:00',
'2022-03-10T18:00:00']}]
```

```
In [5]: # 3. get available timestamps for chosen source and tags
tags = 'operational,numerical weather prediction,ensemble,surface,forecast'

timestamps = requests.get(f'{DISCOVER_URL}/sources/{source}/timestamps?variable={variables}&tags={tags}')
timestamps = timestamps.json()

print(f'\n Timestamps available for variables {variables} in dataset ({source},{tags}):')
print(f'{timestamps[0]}, {timestamps[1]}, {timestamps[2]} ... {timestamps[-1]}')
```

```
Timestamps available for variables air_temperature,eastward_wind,northward_wind,wind_speed_of_gust in dataset (COSMO-2E,operational,numerical weather prediction,ensemble,surface,forecast):
2022-03-09T12:00:00, 2022-03-09T18:00:00, 2022-03-10T00:00:00 ... 2022-04-28T00:00:00
```



Example: read desired dataset using the Python API



Chunked, compressed multi-dimensional arrays in flexible data stores



Flexible, general-purpose parallel computing framework.



High-level API for analysis of multidimensional labelled arrays.

```
In [1]: database = GridfixDatabase(metadata_store, chunk_store)

In [2]: source = 'COSMO-1E'
tags = ['operational', 'numerical weather prediction', 'ensemble', 'surface', 'forecast']
timestamps = ('2022-01-01', '2022-04-27') # range of timestamps when the data was generated
variables = ['eastward_wind', 'northward_wind', 'wind_speed_of_gust'] # variables of interest
```

```
In [3]: dataset = database.get_xarray_dataset(source, tags, timestamps, variables)
```

```
In [4]: dataset
```

Out[4]: xarray.Dataset

► Dimensions: (forecast_reference_time: 395, realization: 11, t: 46, y: 252, x: 376)

▼ Coordinates:

forecast_refer...	(forecast_reference_time)	datetime64[ns]	2022-03-08T15:...		
realization	(realization)	float64	0.0 1.0 2.0 3.0		
time	(forecast_reference_time, t)	datetime64[ns]	...		
longitude	(y, x)	float64	...		
latitude	(y, x)	float64	...		

▼ Data variables:

northward_wind	(forecast_reference_time, t, realization, y, x)	float32	...		
wind_speed_of...	(forecast_reference_time, t, realization, y, x)	float32	...		
eastward_wind	(forecast_reference_time, t, realization, y, x)	float32	...		

dask arrays: - convenient for manipulating big data → development/training use case.
- currently with single machine: local resources used.



Example: read desired dataset using export web service

language independent
access with OPeNDAP

cloud-native → scalable

xarray dataset:

- lazy loading
- `dask.compute()` in service pod → not for manipulating big data → use case inference

slice across any dimension

```
In [7]: import xarray as xr
RETRIEVE_URL = 'http://retrieve-gridefix-main-prod.apps.cp.meteoswiss.ch/gridefix/api'

In [8]: # 4. get variables for last available timestamp containing all of them
dataset = xr.open_dataset(f'{RETRIEVE_URL}/sources/{source}/tags/{tags}/timestamps/{timestamps[-1]}/variables/{variables}')
dataset

Out[8]: xarray.Dataset
  Dimensions:  (forecast_reference_time: 1, realization: 21, t: 121, y: 127, x: 188)
  Coordinates:
    forecast_refer...  (forecast_reference_time)  datetime64[ns]  2022-04-28
    realization        (realization)              float64         0.0 1.0 2.0 3.0 ...
    time               (forecast_reference_time, t)  datetime64[ns]  ...
    longitude          (y, x)                    float64         ...
    latitude           (y, x)                    float64         ...
  Data variables:
    air_temperature    (forecast_reference_time, t, realization, y, x)  float32  ...
    northward_wind     (forecast_reference_time, t, realization, y, x)  float32  ...
    wind_speed_of...   (forecast_reference_time, t, realization, y, x)  float32  ...

In [9]: # 5. load only what you need
data_load = dataset.air_temperature.isel(forecast_reference_time=0, t=0).mean(dim='realization').load()

In [10]: data_load.plot()
Out[10]: <matplotlib.collections.QuadMesh at 0x7f0e8db7e350>
```



Summary and Future work

- GridFix is a **data store** and **service** which offers unified access to multi-dimensional data.
- The cloud-native services allow to define **event-driven architectures**.
- Computations with large amount of data (e.g. training) possible through the Python API with **parallel computations** in dask.

Work in progress and future work:

- Async services: calls to PostgreSQL database not involving zarr.
- Use dask distributed on several machines.
- Quantify thoughtfully performance to quantify tradeoff: less performance at expense of handling metadata and rolling archive.
- Configurations for big archives
 - chunk store options: use several buckets in the cloud object store or local object store.
 - chunks: adapt chunking by use case of dataset in import functions



Schweizerische Eidgenossenschaft
Confédération suisse
Confederazione Svizzera
Confederaziun svizra

Swiss Confederation

Federal Department of Home Affairs FDHA
Federal Office of Meteorology and Climatology MeteoSwiss

MeteoSwiss

Operation Center 1

CH-8058 Zurich-Airport

T +41 58 460 91 11 www.meteoswiss.ch

MeteoSvizzera

Via ai Monti 146

CH-6605 Locarno-Monti

T +41 58 460 92 22

www.meteosvizzera.ch

MétéoSuisse

7bis, av. de la Paix

CH-1211 Genève 2

T +41 58 460 98 88

www.meteosuisse.ch

MétéoSuisse

Chemin de l'Aérologie

CH-1530 Payerne

T +41 58 460 94 44

www.meteosuisse.ch

MeteoSwiss



Metadata store glue code

```
class PostgreSQLStore(MutableMapping):

    def __getitem__(self, key: str) -> str:
        with self.con().cursor() as c:
            c.execute('SELECT value::text FROM zarr_nodes WHERE path = %s', (key,))
            r = c.fetchone()

            if r:
                return r[0]
            else:
                raise KeyError(key)

    def __setitem__(self, key: str, value: bytes) -> None:
        # e.g.
        # key = /COSMO-1E/98e81709-57dd-4bed-b318-ed261188056c/2022-03-30T12:00:00/northward_wind/.zarray
        # value = {"dtype": "<f4", ...}
        # -> will insert the zarr_node and the zarr_variable northward_wind with the timestamp 2022-03-30T12:00:00

        self.upsert({key: value})
        ...

class GridfixDatabase:
    metadata_store: PostgreSQLStore

    def __init__(self, metadata_store, chunk_store=None):

        self.metadata_store = metadata_store
        self.chunk_store = chunk_store

        self.root = zarr.group(store=self.metadata_store,
                               chunk_store=self.chunk_store)
        ...
```



GrideFix: Cloud-native database

