



New Capabilities in MAPL

T. Clune

Atanas Trayanov, Arlindo da Silva,
Ben Auer, William Jamieson, Matt Thompson
Hamid Oloso, and Weiyuan Jiang



Outline

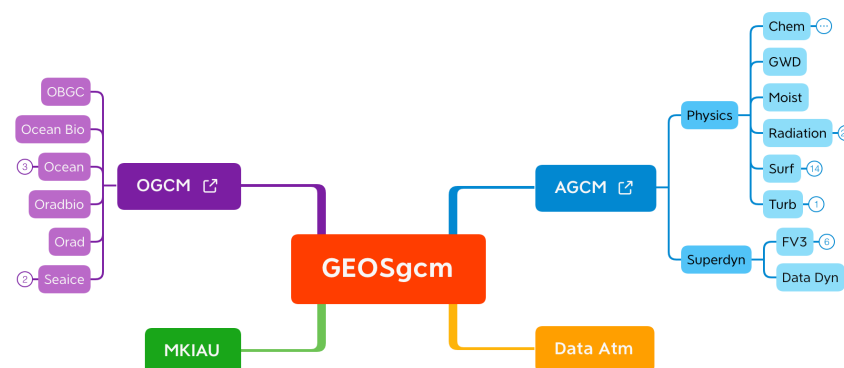
- ❖ What is MAPL
- ❖ What's new
 - ❖ mepo – holistic approach to multi-repo git
 - ❖ ESMF-enabled shortcut for hybrid MPI + OpenMP
 - ❖ MAPL-NUOPC interoperability layer
 - ❖ Provider-subscriber services
- ❖ What's next



What is MAPL


MAPL: a *usability layer* on top of ESMF to simplify the creation and use of ESMF gridded components in a hierarchical architecture

- Easy specification of import, export, and internal states
- Easy addition of child components
- Default implementation for checkpoint/restart.
- Framework standard gridded components:
 - History – run time configurable output
 - ExtData - configurable input – imports of last resort
 - Cap – drives hierarchy; manages History+ExtData
- Wrap user-specified ESMF entry points to
 - Manage instantiation of fields, esp. field connections
 - Profile (time and memory)
 - Insert couplers (when needed)
- Enforce conventions



Similar motivations have led to NUOPC but unfortunately the details diverged.

2020-09-15





mepo

- ❖ **Driver:** Simplify routine development with nested Git repositories
- ❖ **Available technologies:**
 - ❖ Git submodules
 - ❖ Git subtree
 - ❖ Manage Externals
- ❖ **Solution:** Extensible Python package **mepo** (<https://github.com/GEOS-ESM/mepo>)
 - ❖ Typical syntax: `$ mepo <git-command> <git-options> <repo-list>`
- ❖ Currently supported commands

- | | | |
|-------------------|-----------------------------|------------------------|
| • clone | • init | • <i>whereis</i> |
| • status | • stage | • <i>stage</i> |
| • diff | • tag | • <i>unstage</i> |
| • fetch | • stash | • <i>save</i> |
| • checkout | • <i>list</i> | • <i>restore-state</i> |
| • branch | • <i>checkout-if-exists</i> | • <i>fetch-all</i> |
| • commit | • <i>develop</i> | • <i>pull-all</i> |
| • push | • <i>compare</i> | |

M. Thompson, T. Clune,
P. Chakraborty, & A. da Silva



```
~/Models/GEOSgcm on main at 10:45:16 AM
> mepo status
Checking status...
env | (t) v2.1.6 (DH)
  | CMakeLists.txt: modified, not staged
cmake | (t) v3.1.3 (DH)
  | esma.cmake: modified, staged
ecbuild | (t) geos/v1.0.5 (DH)
NCEP_Shared | (t) v1.0.0 (DH)
GMAO_Shared | (t) v1.1.9 (DH)
MAPL | (t) v2.1.6 (DH)
  | include/unused_dummy.H: renamed as headers/unused_dummy.H
FMS | (t) geos/2019.01.02+noaff.1 (DH)
GEOSgcm_GridComp | (t) v1.10.2 (DH)
  | GEOSwgcm_GridComp/GEOS_WgcmGridComp.F90: added, staged but deleted, not staged
FVdycoreCubed_GridComp | (t) v1.2.5 (DH)
  | LatLonToCubeRegridder.F90: renamed, staged as LatLon2CubeRegridder.F90 with unstaged changes
fvdycore | (t) geos/v1.1.3 (DH)
GEOSchem_GridComp | (t) v1.3.5 (DH)
  | GOCART_GridComp/BC_GridComp/BC_GridCompMod.F90: modified, staged with unstaged changes
mom | (t) geos/5.1.0+1.1.1 (DH)
mom6 | (t) geos/v1.1.0 (DH)
GEOSgcm_App | (t) v1.3.7 (DH)
  | unified_setup: untracked file
UMD_Etc | (t) v1.0.3 (DH)
CPLFCST_Etc | (t) v1.0.1 (DH)
~/Models/GEOSgcm on main at 10:45:19 AM
>
```



mepo (cont'd)

- ❖ Top level YAML config file
 - ❖ Mount point in src tree
 - ❖ URL of repo to clone
 - ❖ Tag/branch
 - ❖ Optional path to sparse checkout
 - ❖ Name of “develop” branch (for GitFlow)
- ❖ Caveats
 - ❖ Single config – not nested (design tradeoff)
 - ❖ Just Git; no SVN (or CVS)

```
GMAO_Shared:  
  local: ./src/Shared/@GMAO_Shared  
  remote: ../GMAO_Shared.git  
  tag: v1.1.8  
  sparse: ./config/GMAO_Shared.sparse  
  develop: main
```

```
MAPL:  
  local: ./src/Shared/@MAPL  
  remote: ../MAPL.git  
  tag: v2.1.6  
  develop: develop
```


```
FMS:  
  local: ./src/Shared/@FMS  
  remote: ../FMS.git  
  tag: geos/2019.01.02+noaff.1  
  develop: geos/release/2019.01
```



ESMF-enabled MPI+OpenMP

Driver:

1. Hybrid MPI+OpenMPI improves perf of FV core by up to 50% at extreme scales.
2. Legacy parameterizations not (properly) instrumented with OpenMP.




Solution: Use ESMF to couple pure-MPI physics with hybrid FV.

- ESMF idles subset of PETs; FV launches OMP threads during
- Data transferred using ESMF shared-image Field data storage
 - Low overhead – copy required due to halos anyway

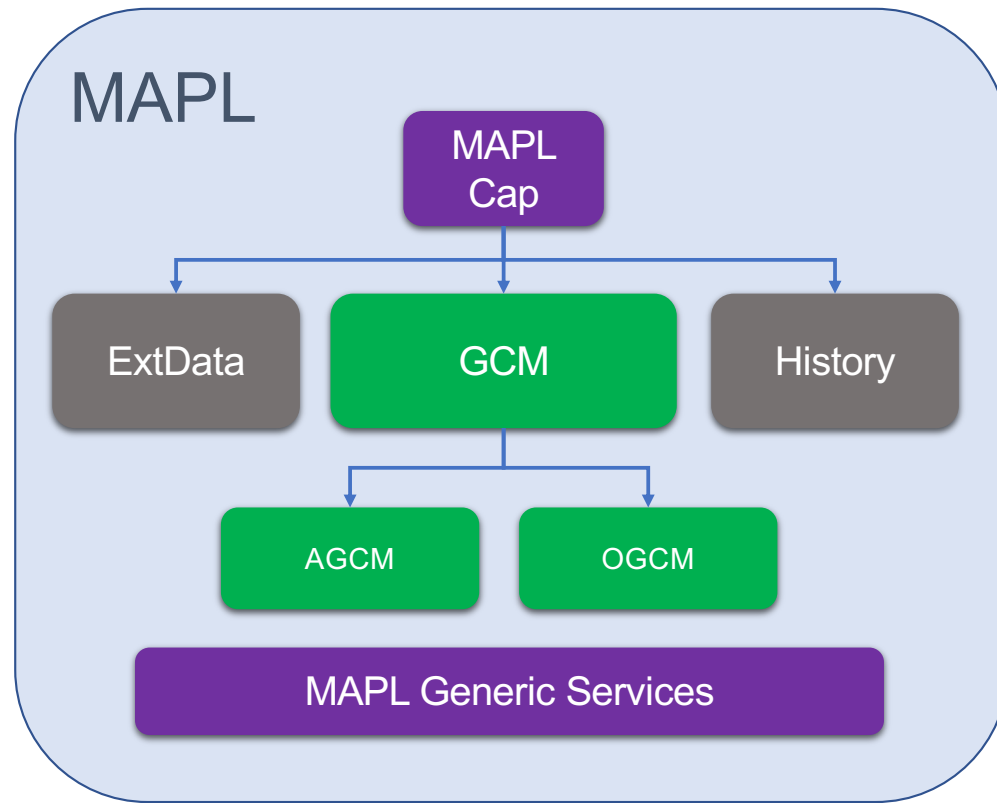

Results encouraging on synthetic example

Investigating why pinning seems to go wonky with full use case.

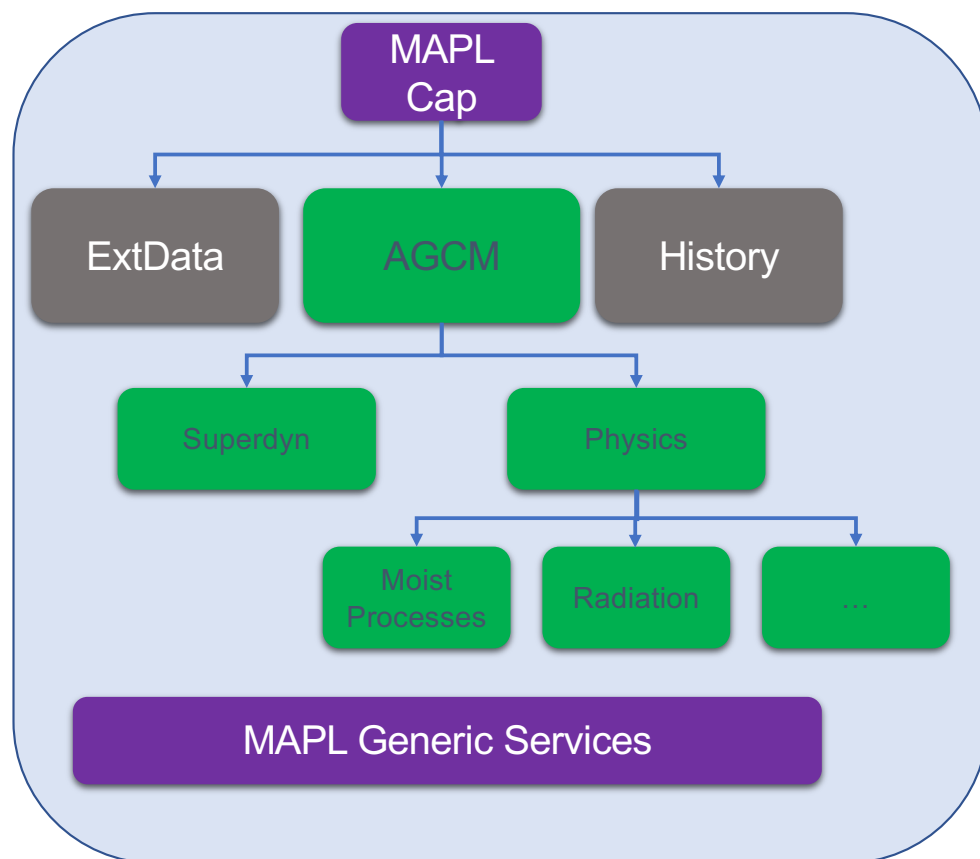
A. Oloso, T.Clune, & G. Theurich



NUOPC – MAPL Interoperability



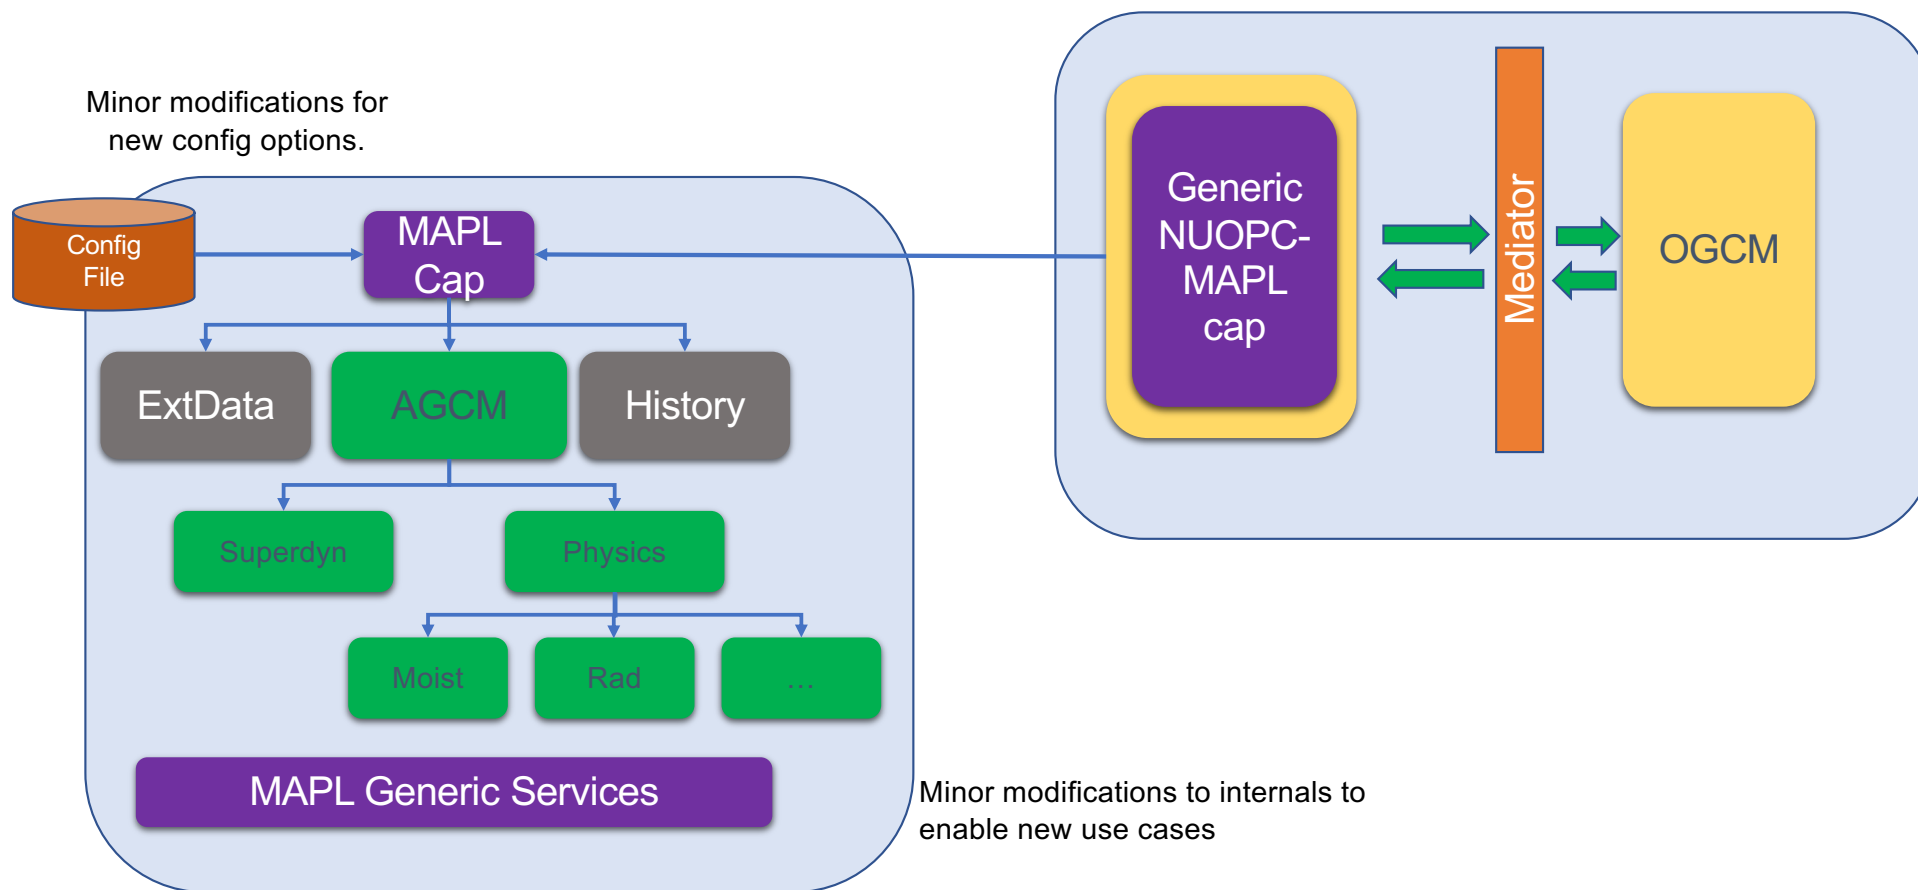
NUOPC –MAPL Interoperability (cont'd)



Drivers:

- MAPL component used in NUOPC architecture
 - E.g. GOCART aerosols exported to NOAA UFS
- Exploiting unique NUOPC capabilities
 - E.g. concurrent execution of MAPL components

NUOPC –MAPL Interoperability (cont'd)





Provider – Subscriber Services

Driver: Consider a simple implementation of advection of chemistry tracers

1. Chemistry tracers passed as imports to dynamics
2. Advection tendencies passed as imports to chemistry

Can require > 150 3D arrays to store tendencies!

Concept: Allow components to “outsource” services as a mechanism to save memory

- Subscriber components (e.g. chemistry,) give explicit permission for state to be modified
- Provider components (e.g., DynCore) directly update bundle of tracers

Legacy mechanism:

- Little direct support from MAPL
- Some logic appears in otherwise uninvolved components (brokers)
- Somewhat inflexible/hardwired
- Fragile

A. Trayanov, T. Clune, & A. da Silva



Provider – Subscriber Services (cont'd)

New approach: Generalize existing MAPL “data service”

- MAPL provides high-level interface – simple to use/understand
- Elevates services to a component-level interface
- Flexible – e.g., send some tracers to DynCore and others to AdvecCore
- Applicable to many processes beyond advection
- Can detect unsatisfied service requests

Pseudo code:

1. “provider” component advertises service
`call advertise_service(self, 'advection', bundle_name)`
2. “subscriber” component requests service with label and list of field names
`call subscribe_to_service(self, 'advection', my_fields)`
3. common ancestor component specifies connections: {service, provider, subscriber}
`call add_connection('advection', provider=child_a, subscriber=child_b)`
4. Under-the-hood: MAPL combines subscriber bundles into single import bundle for provider

A. Trayanov, T. Clune, & A. da Silva



What's Next?


- ❖ Major refactoring of MAPL core layers underway
 - ❖ Clean OO design and comprehensive unit tests
 - ❖ Generalize a few increasingly inconvenient low-level assumptions
 - ❖ Switch from static to dynamic libraries
 - ❖ Enhanced run-time configurability
 - ❖ Eliminate dependencies between components (loose coupling)
- ❖ Improved interoperability with NUOPC
 - ❖ Either next-gen MAPL components are NUOPC caps, *or*
 - ❖ Next-gen MAPL components have factory method to generate a NUOPC cap
- ❖ Share ExtData and History gridded components with community
 - ❖ Establish comprehensive tests and refactor code base
 - ❖ Enable server-side ESMF regridding
 - ❖ Produce thorough tutorials/examples/documentation



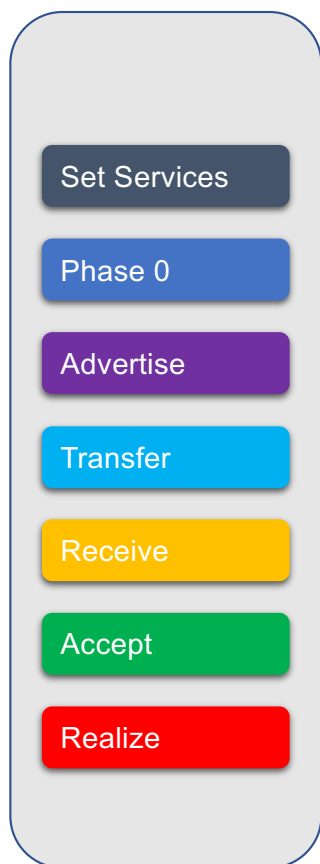
Questions?



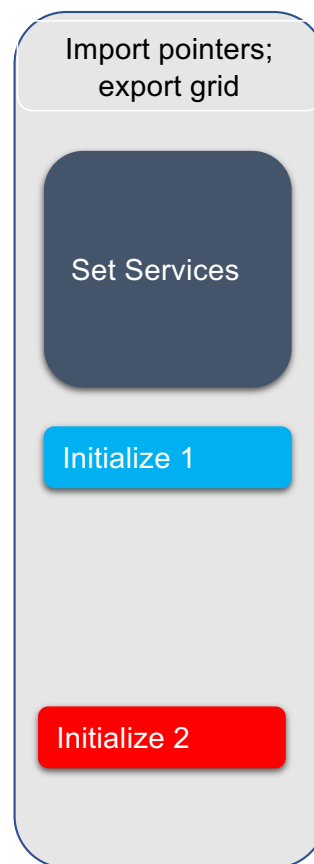

Extra materials



NUOPC



NUOPC-MAPL Cap





MAPL NUOPC Compatibility Layer

Drivers:

1. Concurrent execution of MAPL components – e.g. CTM driven by GCM
2. Run MAPL components within a NUOPC architecture – e.g. GOCART aerosols inside NOAA UFS

Major issues:

- Conflicts between initialization phasing
- Grid sharing
- Data pointer sharing (performance)

Near term solution – Develop a generic NUOPC-MAPL Cap

- New top-level config options:
 - Specify which fields are to imported/exported (using “standard” names).
 - Specifier for whether pointer is shared and/or which side allocates
- Extend MAPL VarSpec to include new config options
- Extend MAPL internal logic to
 - handle variant allocation cases
 - receive externally specified grid

W. Jamieson, T. Clune, & A. da Silva