

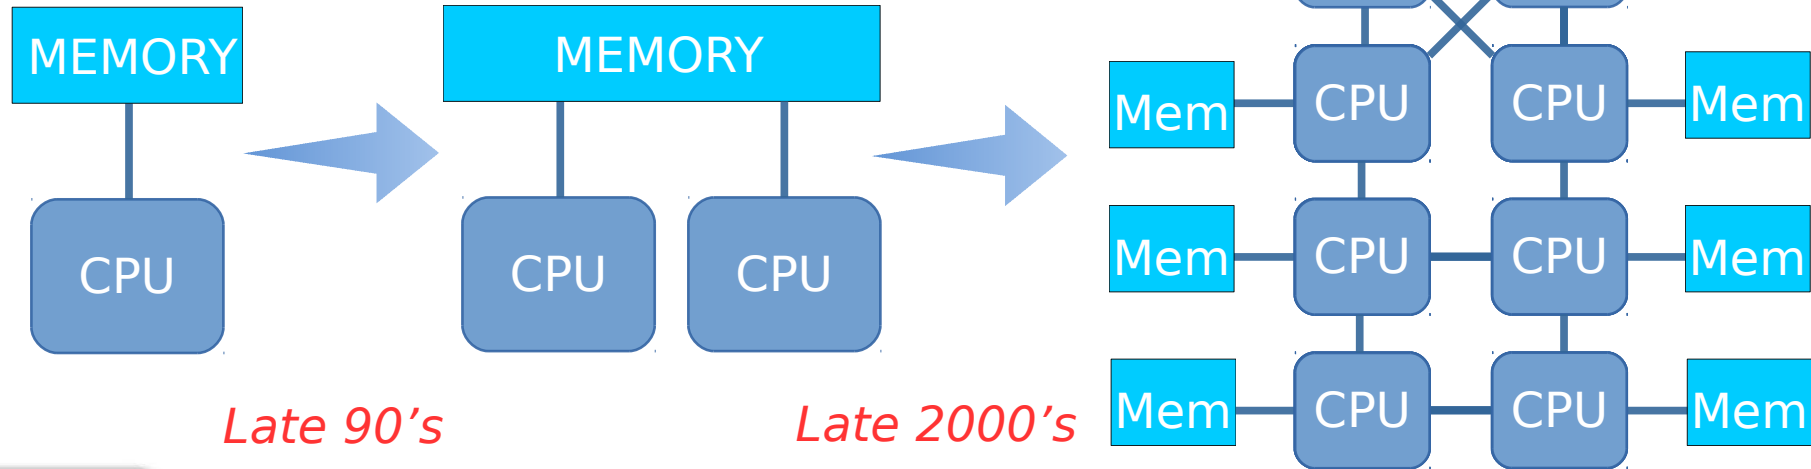


# Hardware Topologies Management in Message-Passing Based Parallel Applications


**Guillaume Mercier**

# Parallel Computers are increasingly complex


- Current trend in computing nodes
  - Increasing core count
  - Deeper memory hierarchies
- Regular users cannot understand/exploit all of this
- Some help is needed to deal with this complexity
  - exploit the hierarchies



# The Performance Heterogeneity Issue



- Depending on process **location** on **cores**
- **Communication performance** can differ
  - **Memory access costs** can change (NUMA effects)



Haswell E5-2680v3 + Infiniband QDR interconnect

# Motivations

- Application developers need abstract features to:
  - Deal with hardware characteristics (Caches, Interconnect, Cores, NUMA nodes, etc.)
    - **Example: Selection of a set of hardware resources in an application**
  - Deal with existing low level tools
    - hwloc
    - numactl
    - Etc.
- Expected performance improvements
  - Improved locality
  - Improved communication performance

# The Message Passing Interface and its programming model

- MPI is widely used for parallel applications since the mid-90's
- Relies on **message exchanges** between processing entities (MPI processes)
- MPI programming model is **flat**
  - Any MPI process can communicate with any other MPI process
  - No routing scheme
  - No locality/hierarchy in the model originally
- MPI is **hardware-agnostic**
  - No assumptions made about the underlying hardware
  - Does not mean that it cannot be accessed at the application level

# Hardware and Locality Management in the current (version 3.1) MPI Standard

- Some additions since MPI\_Get\_processor\_name
- Currently available features:
  - Virtual topologies
  - Neighborhood Collective Communications
  - Shared Memory Support in MPI

# MPI Virtual Topologies

- **“Software locality”**: characterize the application behavior (e.g., its communication pattern)
  - HW independent feature (interface-wise)
  - Virtual to physical mapping “outside of the scope of MPI”
- HW can be taken into account with **the reordering of processes**: possibility to match the HW topology to the virtual topology
  - **Implementation-dependent feature**
    - **Not standard**
    - No guarantee of behavior from one implementation to the other, or even from one version to the other
- Implementation issues
  - Virtual topologies poorly implemented
  - Chicken-and-egg problem

# Neighborhood Collective Communications

- Leverage virtual topologies
  - A virtual topology is attached to a communicator
  - It represents the application's **communication pattern**
- This communicator is an argument of the NCC function
  - Users can **define their own communication pattern** for the collective communication
    - **Better control** of how communications are scheduled
  - Refines the flat programming model of MPI
    - **Improves scalability and locality**



# MPI and Shared Memory Support

- Until MPI 3.0, shmem support is **hidden** by the implementation
  - **Gain control** over the actual shared memory transfers
- **Hybrid programming** (e.g., MPI + X) is now commonplace
  - MPI + OpenMP (i.e. multithreading) is often used
    - Because of multicore nodes
    - Sensible approach (sometimes counter-intuitive)
  - **MPI + MPI is also possible!**
    - Message Passing for internode communications
    - MPI Shared Memory for intranode communications

# MPI Shared Memory Functionalities

- Joint use of:
  - (1) MPI\_Comm\_split\_type with MPI\_COMM\_TYPE\_SHARED split value
    - **Process isolation** on computing nodes
  - (2) MPI\_Win\_allocate\_shared
- Result: **an explicit two-level hierarchy** in the application
  - No MPI overhead for intranode transfers
    - **The MPI stack is bypassed!**
    - **Use load/store operations instead of send/recv**
  - Performance discrepancies within a node not leveraged
    - Use of an other programming model?

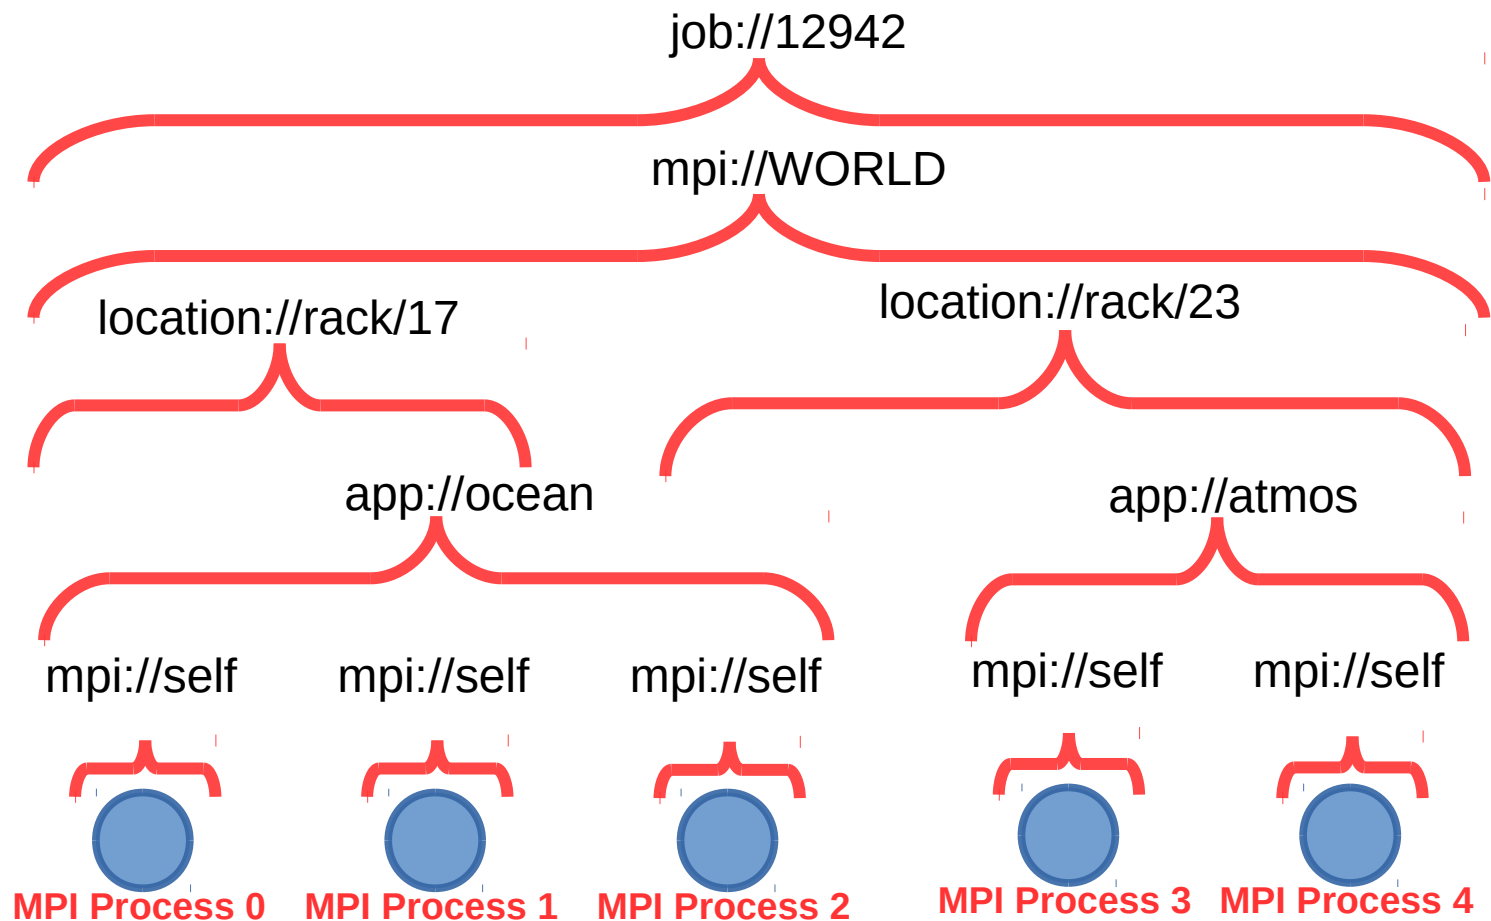
# New features in the MPI Standard 4.0

- ~~Cartesian topologies optimized for hierarchical hardware~~
  - ~~– Very few virtual topologies routines are optimized~~
  - ~~– Not passed in votes yet!~~
- MPI Sessions
- Hardware communicators

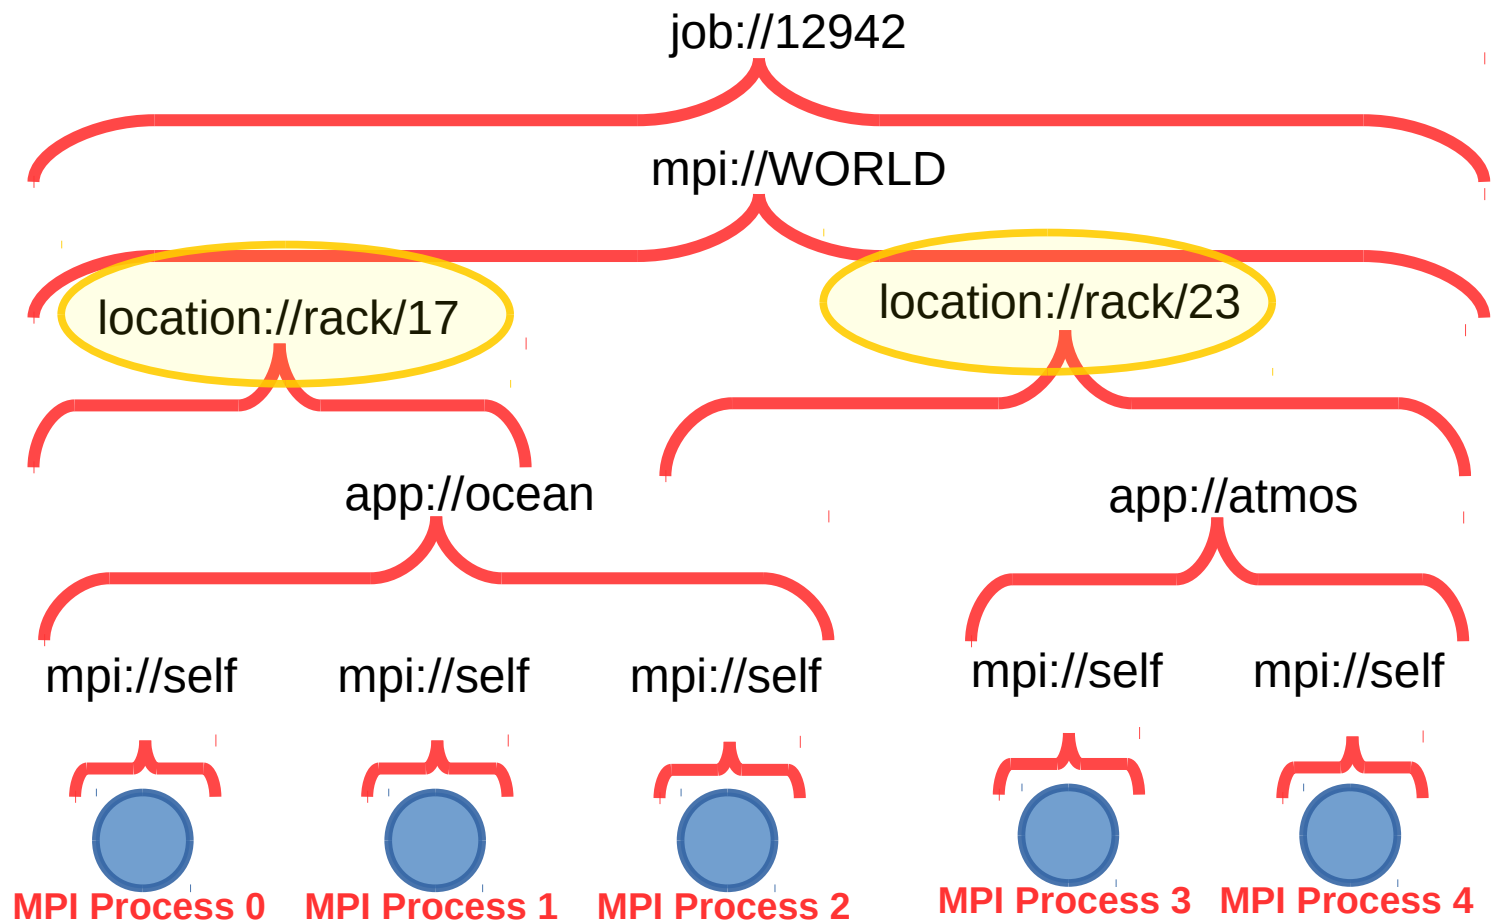
# MPI Sessions

- Original goals
  - Solve composability issues
  - Solve multi init/finalize situations
- Principles
  - Process isolation in “sessions”
    - ➔ **No communications between sessions**
  - Based on **process sets**: psets → groups → communicators
    - Lighter objects than groups (and communicators)
  - Legacy “World Model” to ease the transition to “Sessions Model”
- Possibility to give “names” to sessions: URIs
  - Could extend MPI\_Get\_processor\_name for hardware resources
  - **Naming scheme is problematic**

# MPI Sessions Example: URIs



# MPI Sessions Example: URIs



# Hardware Communicators Purpose

- **Create communicators that convey locality/the sharing of resources between MPI processes**

# Hardware Communicators Example

Machine (377GB total)

Package L#0

Die L#0

L3 (36MB)

Group0

NUMANode L#0 P#0 (46GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#0 Core L#1 Core L#11

PU L#0 P#0 PU L#1 P#1 PU L#11 P#20

Group0

NUMANode L#1 P#1 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#12 Core L#13 Core L#23

PU L#12 P#4 PU L#13 P#5 PU L#23 P#23

Die L#1

L3 (36MB)

Group0

NUMANode L#2 P#2 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#24 Core L#25 Core L#35

PU L#24 P#24 PU L#25 P#25 PU L#35 P#44

Group0

NUMANode L#3 P#3 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#36 Core L#37 Core L#47

PU L#36 P#28 PU L#37 P#29 PU L#47 P#47

Package L#1

Die L#2

L3 (36MB)

Group0

NUMANode L#4 P#4 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#48 Core L#49 Core L#59

PU L#48 P#48 PU L#49 P#49 PU L#59 P#68

Group0

NUMANode L#5 P#5 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#60 Core L#61 Core L#71

PU L#60 P#52 PU L#61 P#53 PU L#71 P#71

Die L#3

L3 (36MB)

Group0

NUMANode L#6 P#6 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#72 Core L#73 Core L#83

PU L#72 P#72 PU L#73 P#73 PU L#83 P#92

Group0

NUMANode L#7 P#7 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#84 Core L#85 Core L#95

PU L#84 P#76 PU L#85 P#77 PU L#95 P#95



# Hardware Communicators Example

Machine (377GB total)

Package L#0

Die L#0

L3 (36MB)

Group0

NUMANode L#0 P#0 (46GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#0 Core L#1 Core L#11

Group0

NUMANode L#1 P#1 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#12 Core L#13 Core L#23

Die L#1

L3 (36MB)

Group0

NUMANode L#2 P#2 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#24 Core L#25 Core L#35

Group0

NUMANode L#3 P#3 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#36 Core L#37 Core L#47

Package L#1

Die L#2

L3 (36MB)

Group0

NUMANode L#4 P#4 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#48 Core L#49 Core L#59

Group0

NUMANode L#5 P#5 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#60 Core L#61 Core L#71

Die L#3

L3 (36MB)

Group0

NUMANode L#6 P#6 (47GB)

L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#72 Core L#73 Core L#83

Group0

NUMANode L#7 P#7 (47GB)

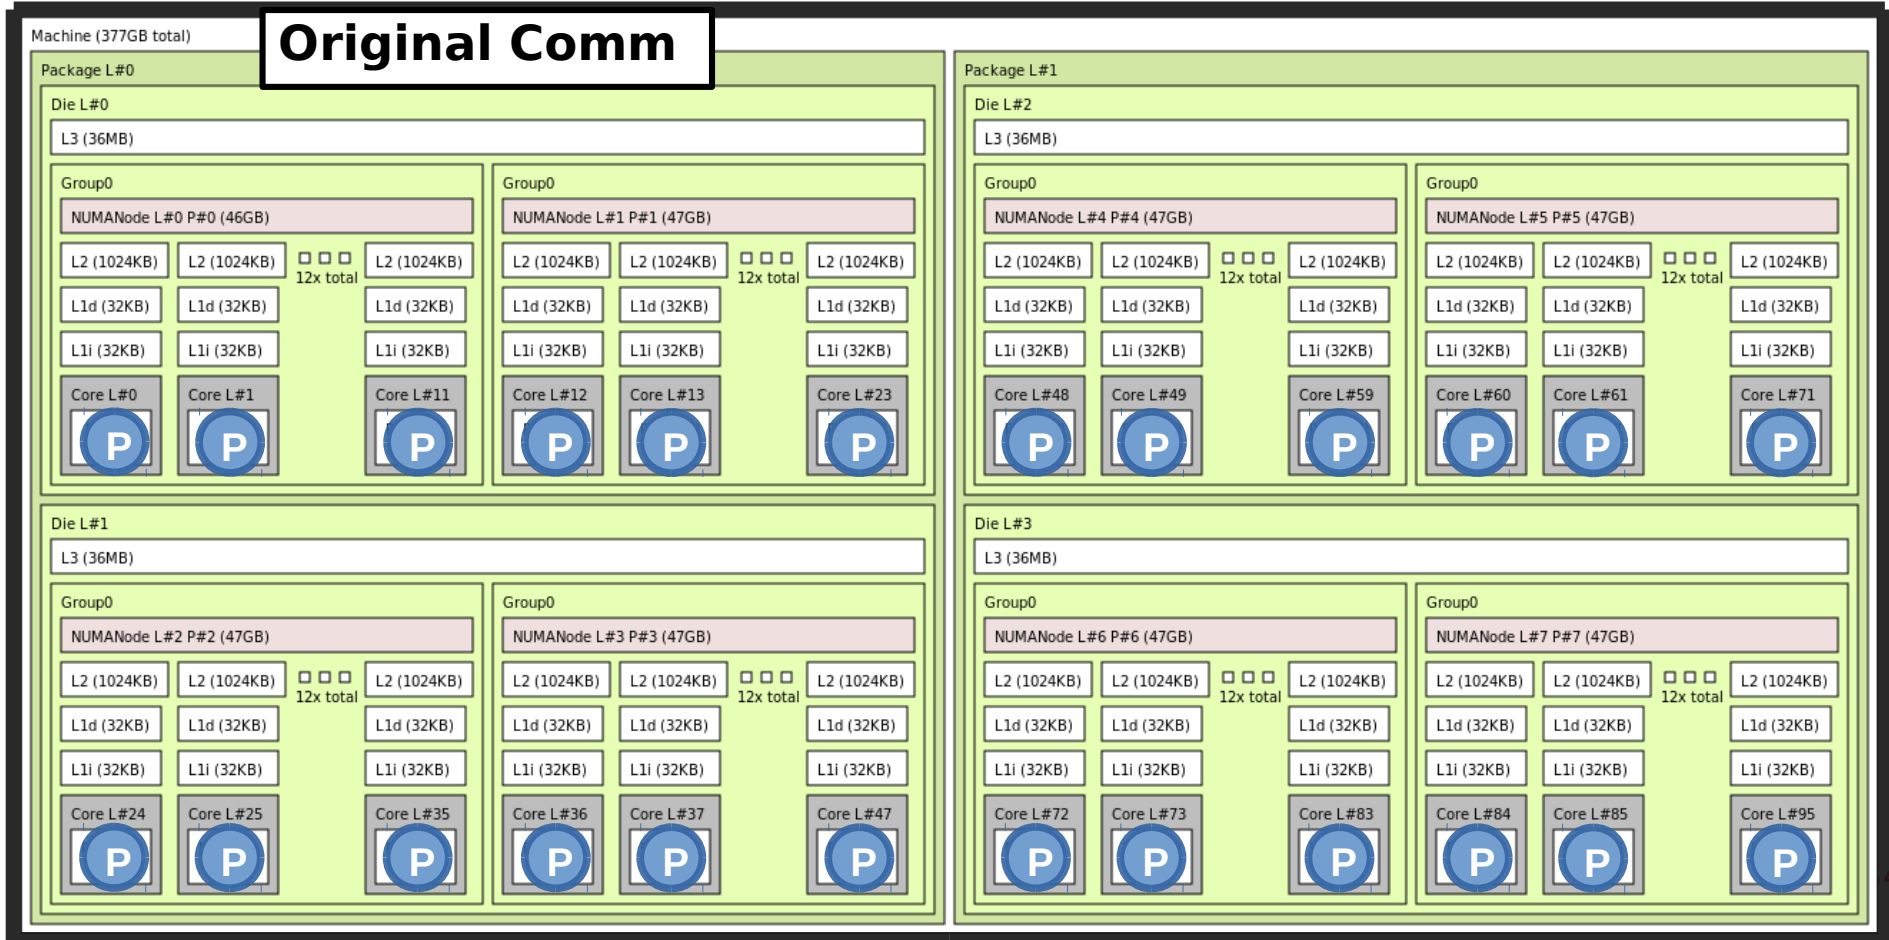
L2 (1024KB) L2 (1024KB) 12x total L2 (1024KB)

L1d (32KB) L1d (32KB) L1d (32KB)

L1i (32KB) L1i (32KB) L1i (32KB)

Core L#84 Core L#85 Core L#95

# Hardware Communicators Example



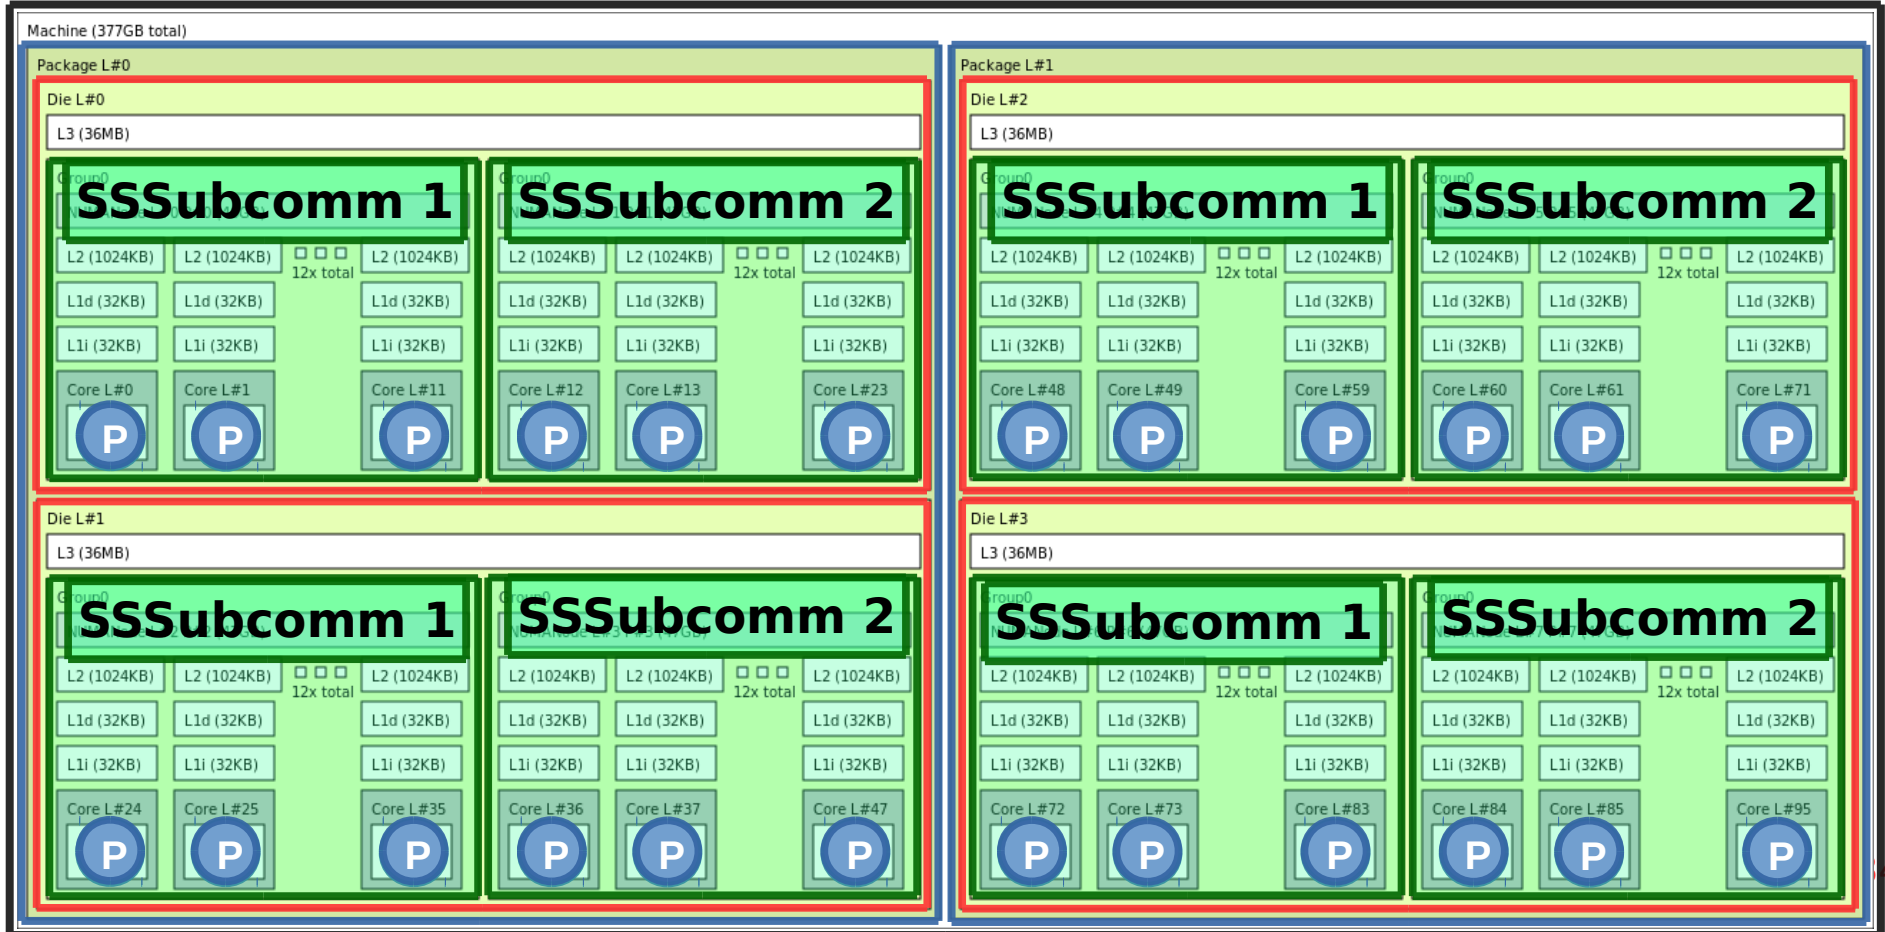
# Hardware Communicators Example



# Hardware Communicators Example



# Hardware Communicators Example



# Hardware Communicators Example



# Hardware Communicators Proposal

- **Extension of an already existing** communicator creation operation:

```
MPI_Comm_split_type(MPI_Comm comm    /* IN */,  
                    int split_type /* IN */,  
                    int key         /* IN */,  
                    MPI_Info info    /* INOUT */,  
                    MPI_Comm newcomm /* OUT */);
```

# Hardware Communicators Proposal

- **Extension of an already existing** communicator creation operation:

```
MPI_Comm_split_type(MPI_Comm comm    /* IN */,  
                    int split_type    /* IN */,  
                    int key           /* IN */,  
                    MPI_Info info     /* INOUT */,  
                    MPI_Comm newcomm /* OUT */);
```

- With:

- Two new `split_type` values:
  - **MPI\_COMM\_TYPE\_HW\_UNGUIDED**: split at the next level in the platform
    - Valid ( i.e., non `COMM_NULL` ) new comms are strict subsets of the old comm
  - **MPI\_COMM\_TYPE\_HW\_GUIDED**: split for a specified resource type
- One new info key: `mpi_hw_resource_type`
  - To constrain the splitting operation (guided mode)
  - To query the type of resource represented by `newcomm` (unguided mode)



# Hardware Communicators Proposal

- **Extension of an already existing** communicator creation operation:

```
MPI_Comm_split_type(MPI_Comm comm /* IN */,  
                    int split_type /* IN */,  
                    int key /* IN */,  
                    MPI_Info info /* INOUT */,  
                    MPI_Comm newcomm /* OUT */);
```

- With:

- Two new `split_type` values:
  - **MPI\_COMM\_TYPE\_HW\_UNGUIDED**: split at the next level in the platform
    - Valid ( i.e., non `COMM_NULL` ) new comms are strict subsets of the old comm
  - **MPI\_COMM\_TYPE\_HW\_GUIDED**: split for a specified resource type
- One new info key: **`mpi_hw_resource_type`**
  - To constrain the splitting operation (guided mode)
  - To query the type of resource represented by newcomm (unguided mode)

# Unguided Mode Example

```
#define MAX_NUM_LEVELS 32
```

```
MPI_Comm hwcomm[MAX_NUM_LEVELS];  
int rank, level_num = 0;
```

```
hwcomm[level_num] = MPI_COMM_WORLD;
```

```
while((hwcomm[level_num] != MPI_COMM_NULL) &&  
      (level_num < MAX_NUM_LEVELS-1))
```

```
{
```

```
    MPI_Comm_rank(hwcomm[level_num], &rank);
```

```
    MPI_Comm_split_type(hwcomm[level_num],  
                        MPI_COMM_TYPE_HW_UNGUIDED,  
                        rank,  
                        MPI_INFO_NULL,  
                        &hwcomm[level_num+1]);
```

```
    level_num++;
```

```
}
```

*Splitting  
operation*

*Recursive  
Splitting of  
MPI\_COMM\_WORLD*

# Guided Mode Example

Splitting MPI\_COMM\_WORLD into NUMANode subcommunicators:

```
MPI_Info info;  
int rank;  
MPI_Comm hwcomm;  
  
MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
  
MPI_Info_create(&info);  
MPI_Info_set(info, "mpi_hw_resource_type", "NUMANode");  
  
MPI_Comm_split_type(MPI_COMM_WORLD,  
                    MPI_COMM_TYPE_HW_GUIDED,  
                    rank,  
                    info,  
                    &hwcomm);  
  
/* Use hwcomm now */
```

# Guided Mode Example

Splitting MPI\_COMM\_WORLD into NUMANode subcommunicators:

```
MPI_Info info;  
int rank;  
MPI_Comm hwcomm;  
  
MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
  
Info object set-up { MPI_Info_create(&info);  
                    MPI_Info_set(info, "mpi_hw_resource_type", "NUMANode");  
  
                    MPI_Comm_split_type(MPI_COMM_WORLD,  
                                         MPI_COMM_TYPE_HW_GUIDED,  
                                         rank,  
                                         info,  
                                         &hwcomm);  
  
                    /* Use hwcomm now */
```

# Guided Mode Example

Splitting MPI\_COMM\_WORLD into NUMANode subcommunicators:

```
MPI_Info info;  
int rank;  
MPI_Comm hwcomm;
```

*Implementation dependent name*



```
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

*Info object set-up* {  
MPI\_Info\_create(&info);  
MPI\_Info\_set(info, "mpi\_hw\_resource\_type", "NUMANode");

```
MPI_Comm_split_type(MPI_COMM_WORLD,  
MPI_COMM_TYPE_HW_GUIDED,  
rank,  
info,  
&hwcomm);
```

```
/* Use hwcomm now */
```

# Guided Mode Example

Splitting MPI\_COMM\_WORLD into NUMANode subcommunicators:

```
MPI_Info info;  
int rank;  
MPI_Comm hwcomm;
```

*Implementation dependent name*



```
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

*Info object  
set-up*


```
{ MPI_Info_create(&info);  
  MPI_Info_set(info, "mpi_hw_resource_type", "NUMANode");
```

*Splitting  
operation*

```
{ MPI_Comm_split_type(MPI_COMM_WORLD,  
                     MPI_COMM_TYPE_HW_GUIDED,  
                     rank,  
                     info,  
                     &hwcomm);
```

```
/* Use hwcomm now */
```

# Results: Hierarchical Matrix Multiplication



# MPI 4.X: Hardware Information Query Proposal

- **Resource types are implementation dependent**

- Query function to retrieve the types recognized by the MPI implementation:

```
MPI_Get_hw_resource_types(MPI_Info info /* OUT */);
```

- User can also check if the **resource type is supported** by the MPI implementation with:

```
MPI_Get_hw_resource_status(char *res_type /* IN */,  
                           int result /* OUT */);
```



# On the Standardization Front...

- Continue current efforts in MPI for 4.X
- **Interact** more with other HPC communities (e.g., OpenMP)
- **Process Mapping/Binding is a mess!**
  - Discussions started in the Hardware Topologies Working Group of the MPI Forum
  - CEA initiative to set up a **dedicated working group**, outside of the MPI Forum
    - Target: the whole HPC ecosystem, not just MPI libraries
    - Goal: identify **necessary components** and their **interactions** in the HPC ecosystem



*That's all Folks!*