

# Adaptive Parallelism for Coupled, Hybrid Programs

## CW2020



Samuel K. Gutiérrez

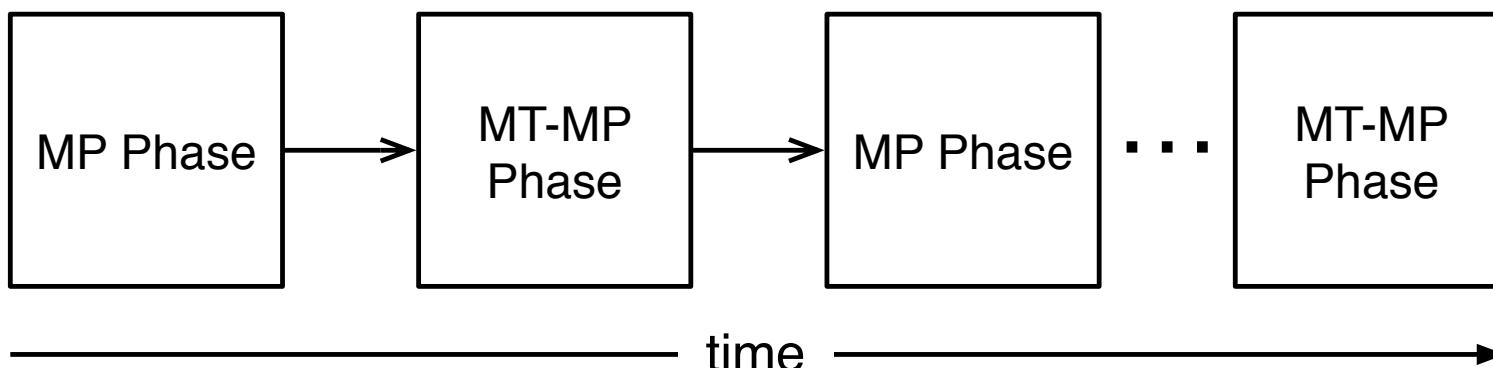
September 22, 2020



Operated by Los Alamos National Security, LLC for the U.S. Department of Energy's NNSA

# Typical Hybrid Programming Models in HPC

- Multi-threaded message-passing (**MT-MP**) applications becoming more popular (e.g., **MPI+OpenMP**)
- **MP + MT-MP** commonplace in *coupled applications*
  - MP (e.g., **MPI Everywhere**) and MT-MP libraries linked together with interleaved execution

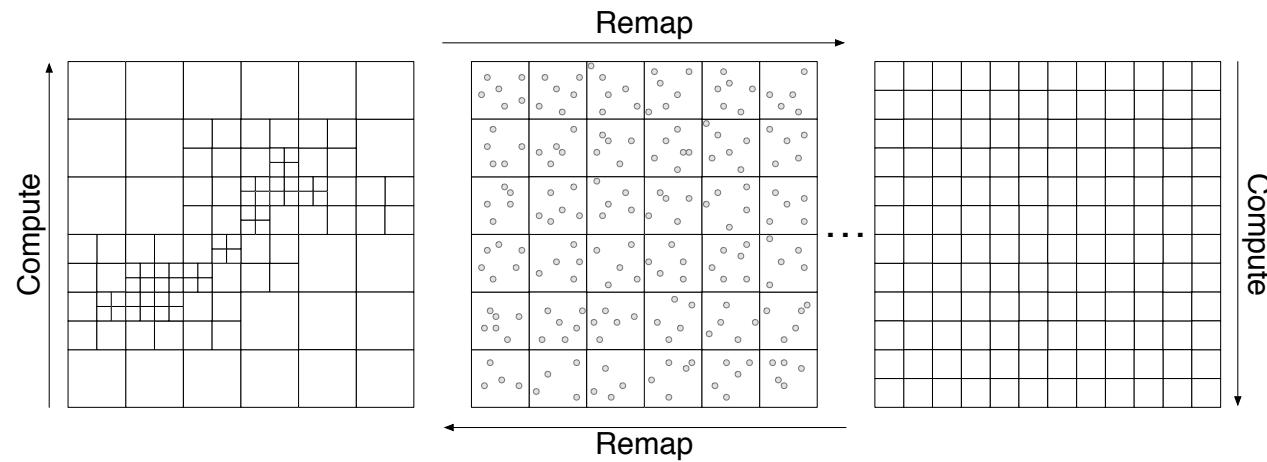


# Where does this Type of Coupling Occur?

## Parallel and distributed multi-physics applications

- Crucial in science and engineering
- Often interdisciplinary effort
  - Built by integration (or *coupling*) of independently developed and tuned software libraries linked into a single application

Phases of  
single, coupled →  
application



# Challenges of Coupling Disparate Libraries

Each library should execute based on its design and tuning

- Each has own optimal *runtime configuration*
  - E.g., number and placement of *tasks* (processes and threads)
  - Poor library configuration → poor library performance → poor application performance
- **Static parameters**—often found manually, heuristically, and offline
- *Configuration conflicts* arise when optimal library configurations conflict with each other

# State of the Practice

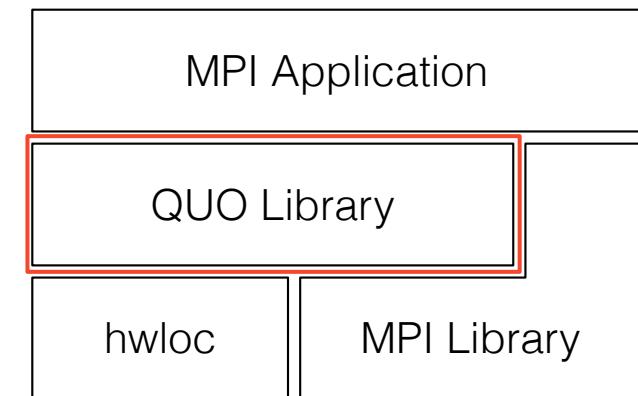
- Task-to-hardware bindings can improve performance
- Parallel launchers (e.g., srun, orte run) support only static binding and placement
  - Launch-time configuration persists for **entire execution**
- Two basic, static configuration options:
  - *Over-subscription* or *under-subscription*

# Quo Approach: Efficient Dynamic MPI+X

- Dynamic, run-time configuration conflict resolution for performant coupled thread-heterogeneous MP apps.
- Specifically, dynamic, coupled applications based on the Message Passing Interface (MPI)
  - MPI+X, where X is a Pthread-based runtime library
    - E.g., MPI-everywhere plus MPI+OpenMP or MPI+std::thread

# Quo Features

- Programmer driven
- Hardware/software state queries (hardware topology, task config.)
- Efficient task reconfiguration (task placement, task affinity)
- Composable and general
- Open-source library ([github.com/lanl/libquo](https://github.com/lanl/libquo))
- C and Fortran interfaces



# Quo Main Concepts

- Contexts (instance handles)      `QUO_create(&ctx, MPI_COMM_WORLD)`

`QUO_free(ctx)`

# Quo Main Concepts

- Contexts (instance handles)
- Hardware and software environment queries

// Hardware topology queries

QUO\_nobjs\_by\_type()

QUO\_nobjs\_in\_type\_by\_type()

// Intra-task state queries

QUO\_cpuset\_in\_type()

QUO\_bound()

// Inter-task state queries

QUO\_qids\_in\_type()

QUO\_create(&ctx, MPI\_COMM\_WORLD)

// Dynamically determine target resource

tres = QUO\_OBJ\_NUMANODE

// Query hardware/software at run-time

QUO\_auto\_distrib(ctx, tres,

max\_pe, &in\_dset)

QUO\_free(ctx)

# Quo Main Concepts

- Contexts (instance handles)
- Hardware and software environment queries
- Dynamic process affinities
  - Stack-based semantics

```
QUO_create(&ctx, MPI_COMM_WORLD)
// Dynamically determine target resource
tres = QUO_OBJ_NUMANODE
// Query hardware/software at run-time
QUO_auto_distrib(ctx, tres,
                  max_pe, &in_dset)
```

```
QUO_bind_push(ctx, tres)
A_library_call(in_args, &result)
QUO_bind_pop(ctx)
```

```
QUO_free(ctx)
```

# Quo Main Concepts

- Contexts (instance handles)
- Hardware and software environment queries
- Dynamic process affinities
  - Stack-based semantics
- Efficient node-local process quiescence

```
QUO_create(&ctx, MPI_COMM_WORLD)
// Dynamically determine target resource
tres = QUO_OBJ_NUMANODE
// Query hardware/software at run-time
QUO_auto_distrib(ctx, tres,
    max_pe, &in_dset)
```

```
QUO_bind_push(ctx, tres)
A_library_call(in_args, &result)
QUO_bind_pop(ctx)
```

```
QUO_barrier(ctx)
```

```
QUO_free(ctx)
```

# Quo Main Concepts

- Contexts (instance handles)
- Hardware and software environment queries
- Dynamic process affinities
  - Stack-based semantics
- Efficient node-local process quiescence
- Data dependencies

```
QUO_create(&ctx, MPI_COMM_WORLD)
// Dynamically determine target resource
tres = QUO_OBJ_NUMANODE
// Query hardware/software at run-time
QUO_auto_distrib(ctx, tres,
                  max_pe, &in_dset)
// Satisfy outstanding data dependencies
if (in_dset)
    QUO_bind_push(ctx, tres)
    A_library_call(in_args, &result)
    QUO_bind_pop(ctx)

QUO_barrier(ctx)
// Satisfy outstanding data dependencies
QUO_free(ctx)
```

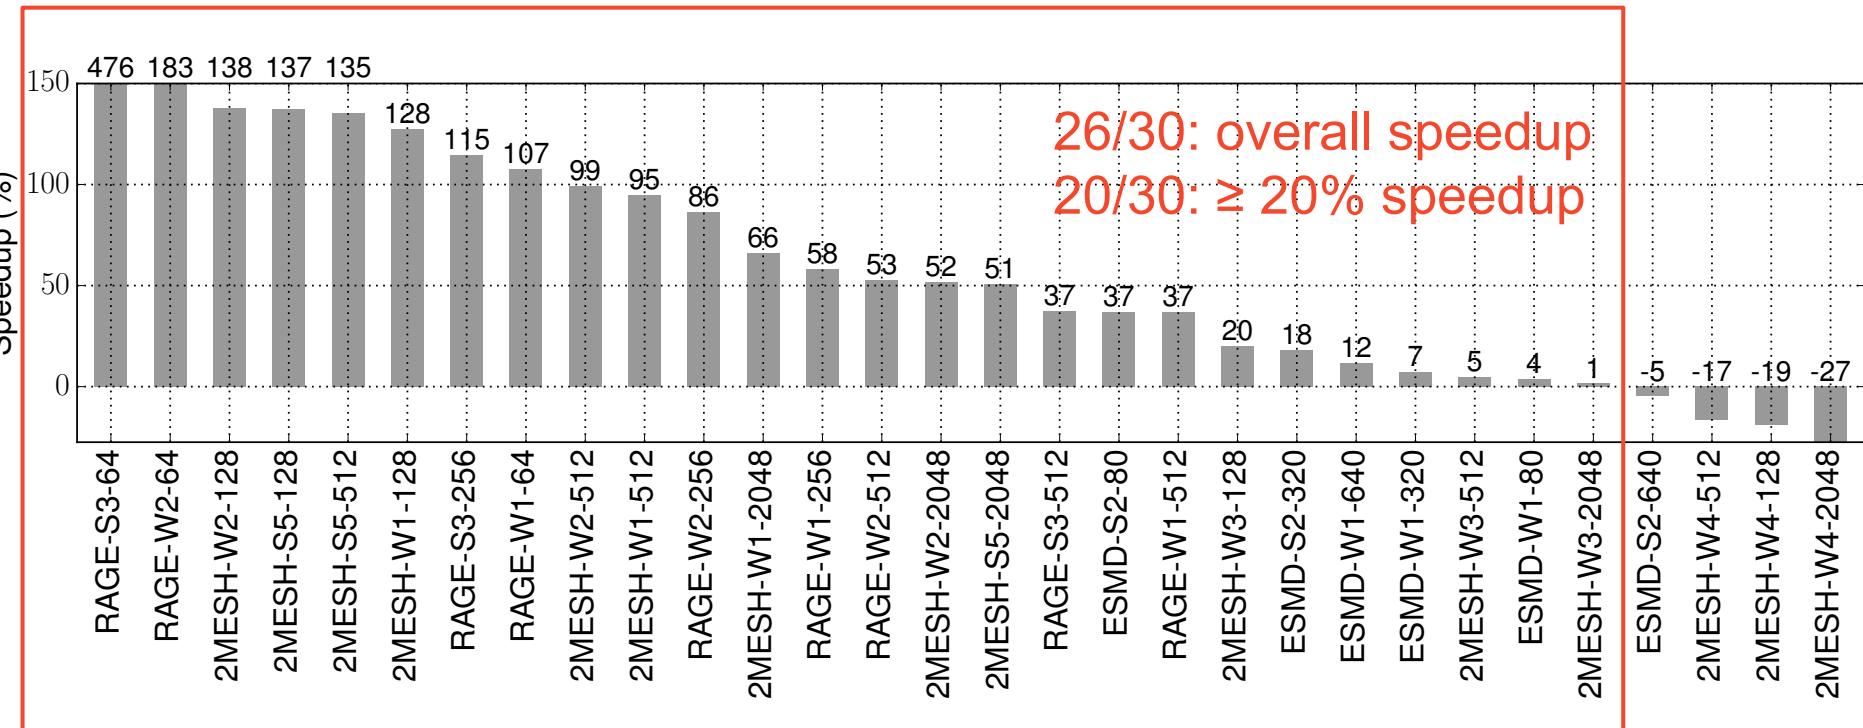
# Quo Main Concepts

- Contexts (instance handles)
- Hardware and software environment queries
- Dynamic process affinities
  - Stack-based semantics
- Efficient node-local process quiescence
- Data dependencies
- Policy management

```
QUO_create(&ctx, MPI_COMM_WORLD)
// Dynamically determine target resource
tres = QUO_OBJ_NUMANODE
// Query hardware/software at run-time
QUO_auto_distrib(ctx, tres,
                  max_pe, &in_dset)
// Satisfy outstanding data dependencies
if (in_dset)
    QUO_bind_push(ctx, tres)
    A_library_call(in_args, &result)
    QUO_bind_pop(ctx)

QUO_barrier(ctx)
// Satisfy outstanding data dependencies
QUO_free(ctx)
```

# Quo's Efficacy



\*Details Available in: Gutiérrez et al. *Accommodating Thread-Level Heterogeneity in Coupled Parallel Applications.* (IPDPS 2017). [ieeexplore.ieee.org/abstract/document/7967136](http://ieeexplore.ieee.org/abstract/document/7967136)

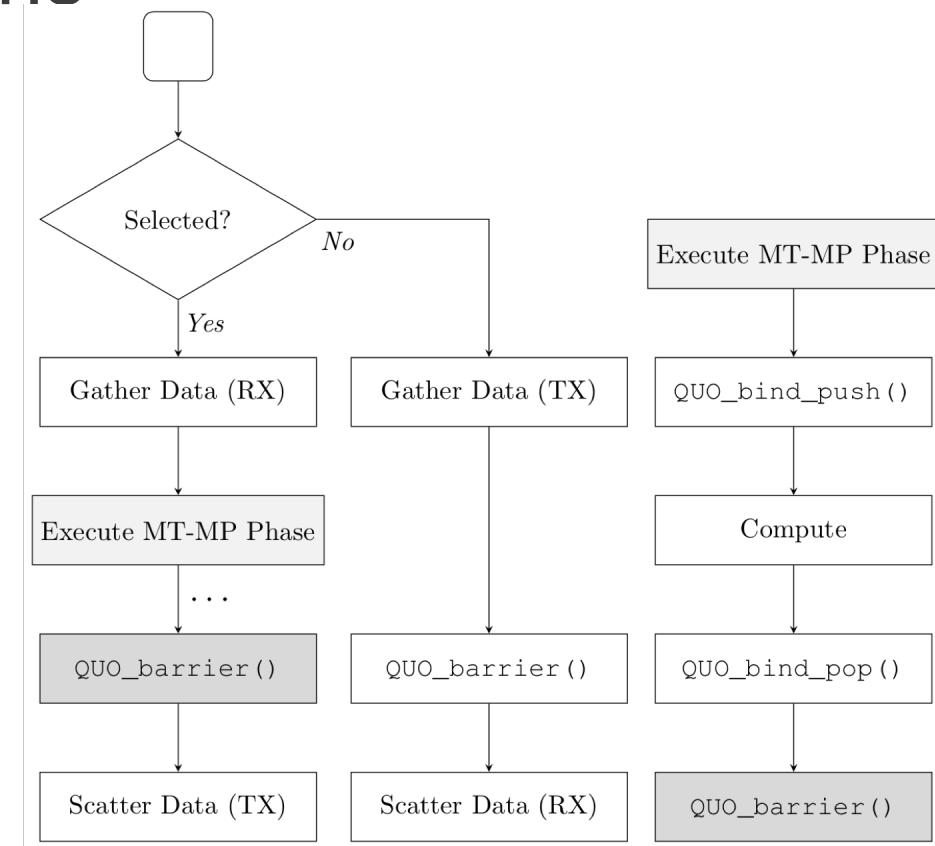
# Practical Considerations

## Increased code complexity

- Task quiescing + resumption  
→ data remapping

## Encapsulating code regions

- QUO\_bind\_push() +  
QUO\_bind\_pop()



\*Details Available in: Samuel K. Gutiérrez. *Adaptive Parallelism for Coupled, Multithreaded Message-Passing Programs*. (2018). [digitalrepository.unm.edu/cs\\_etds/95](https://digitalrepository.unm.edu/cs_etds/95)

# In Summary

Coupled applications increasingly comprise MP libraries with differing preferred degrees of threading

Static execution environments are often suboptimal for resolving dynamic, phased configuration conflicts

Quo's approach can help improve overall application performance with dynamic execution environments

# Thank You

## Questions?

This research was supported by the Advanced Simulation and Computing Program of the U.S. Department of Energy's NNSA and by the Exascale Computing Project (17-SC-20-SC), a collaborative effort of the U.S. Department of Energy Office of Science and the National Nuclear Security Administration.