

# XIOS-3

## Toward a new infrastructure of HPC services and model coupling



**Y. Meurdesoif (IPSL - CEA/DRF/LSCE)**


**J. Dérouillat (IPSL - CEA/DRF/LSCE)**

**A. Caubel (IPSL - CEA/DRF/LSCE)**




## XIOS-1 (2014)

- I/O description outsourced of models in external XML files
  - Simple fortran interface : `xios_send_field("id", field)`
  - Compact and flexible XML description using hierarchical concept
- Asynchronous transfer to dedicated parallel I/O Servers
  - Overlap transfer and writing time by computation
- Parallel write using parallel file system capability
- Targeted for coupled models
  - Interfaced with OASIS



## XIOS-2 (2017)

- Add asynchronous reading capability from servers
  - `xios_rcv_field("id", field)`
- Add "in-situ" parallel workflow computing, developing filters for :
  - Time integration (instant, averaging, min, max...)
  - Arithmetic combination of fields
  - Spatial transformation
    - horizontal and vertical interpolation, sub-domain extraction, reductions, etc.
    - Interpolation : weight computation "on the fly"
- Complex workflow can be achieved by chaining filters before data flux are sent to servers or returned to model (reading)




## XIOS-2.5 (2018)

- Add second levels of servers in order to increase file writing concurrency between servers
  - Activating netcdf writing compression in parallel runs
  - Time series management

### ⇒ Reference version for CMIP6 experiments

- DR2XML (CNRM) : translate automatically CMIP6 data Request into xml xios files
- ~1000 of different variables generated for one CMIP6 deck
- All post-treatments done “in the fly”, automatically CMORized (IPSL and Météo-France/CNRM ESM)



## XIOS-3 (end-2022) : total rewrite of the internal XIOS core engine

- 3 years work of intense developments, touching more than the half part of code lines
  - 514 file modified, 244 SVN commits, 60 000 code lines modified/added/deleted (over 110 000 of total code lines)


```
irene171 work~/XIOS3/src>svn diff -r 1749 | diffstat -m -s
514 files changed, 36824 insertions(+), 19441 deletions(-), 5209 modifications(!)
```

- Goals
  - Cleaning code and rationalizing internal concept due to years of eclectic development
  - Improvement of workflow performance and memory footprint reduction
  - Improvement of robustness and reliability
  - New infrastructure introducing XIOS HPC services concept and model coupling

- Implementation of a non regression suite testcase for continuous integration


Table of XIOS unit tests results

Revision	Jean-Zay	Irene						
2384	X64_JEANZAY_PGL_prod <input checked="" type="checkbox"/> X64_JEANZAY_prod <input type="checkbox"/>	X64_Irene_prod <input checked="" type="checkbox"/>						
		test_scalar_algo <input type="checkbox"/>				✗		
		test_grid_algo <input type="checkbox"/>					✓	
			CONFIG_AxisMask_false_DomMask_false_Dom_Imdz_Duration_1d_NDITr_false_NbCnt_4_NbSrv_2_OneSided_false_PctSrv2_50_ScaMask_sparse_Srv2_false <input type="checkbox"/>				✓	
		test_dynamic_algo <input type="checkbox"/>					✓	
		test_nemo_algo <input type="checkbox"/>					✓	
			CONFIG_AxisMask_false_DomMask_true_Dom_nemo_Duration_1d_NDITr_false_NbCnt_4_NbSrv_2_OneSided_false_PctSrv2_50_ScaMask_none_Srv2_false <input type="checkbox"/>				✓	
							atm_output_reduce_axis.nc <input type="checkbox"/>	✓
			CONFIG_AxisMask_false_DomMask_false_Dom_nemo_Duration_1d_NDITr_false_NbCnt_4_NbSrv_2_OneSided_false_PctSrv2_50_ScaMask_none_Srv2_false <input type="checkbox"/>					✓
		test_domain_algo <input type="checkbox"/>					✓	
		test_function <input type="checkbox"/>					✓	
		test_axis_algo <input type="checkbox"/>					✓	
		test_memory <input type="checkbox"/>					✗	
			X64_Irene_GNU_prod <input checked="" type="checkbox"/>					✓



EXAMPLE OF THE XIOS WORKFLOW VISUALISATION.

- Representation of XIOS workflow execution in the form of graphs, viewable through a web browser
- Tools to track internal memory usage and memory leak, time line visualization through web browser



- Help for debugging : output of the XIOS software stack in case of a crash, with relevant information
- Additional internal output timers at the end of the simulation for better performance profiling

- Development of new client/server transfer protocols based on passive one-sided MPI3 communication
- New concepts of 'views' and 'connectors' for distributed management of workflow grids
- Reduction of the memory footprint by applying tensor product properties onto elements (domains, axis, scalars) composing a grid
- Full rewrite of transformation engine
- Full rewrite of the chaining filters engine

## ⇒ Increase the transfer protocol fluidity, performance improvement

- Under evaluation : testcase : NEMO 4 configuration 1440 x 1680 x 75, 20000 timesteps, 2688 process, 80 XIOS servers, 2 levels of server, write every 50 ts
  - Whole time NEMO no IO (without initialization) Reference => **3051 s**
  - Whole time NEMO (without initialization) XIOS 2 => **3462 s : XIOS overhead 411 s => 13% overhead**
  - Whole time NEMO (without initialization) XIOS 3 => **3186 s : XIOS overhead 135 s => 4.4% overhead**

⇒ Reduction of the XIOS overhead by a factor 3

**3.5 Tb generated over 3000s => 1.2 Gb/s**

⇒ **8% speed-up on this configuration**

- Preliminary results, can be configuration dependant



## ⇒ Memory footprint reduction

- Same NEMO configuration
- XIOS2 Vs XIOS3 Client+Model : **reduction of 20% of whole virtual memory**
- XIOS2 Vs XIOS3 Client part : **reduction by a factor 3 of virtual memory consumption**
- XIOS2 Vs XIOS3 Server side : **reduction of virtual memory consumption up to a factor 3**

	XIOS2	XIOS3
Client + Model	150 Mo	120 Mo
Server N1	3.75 Go	2 Go
Server N2	30 Go	10 Go

## What is an XIOS service ?


- A parallel and asynchronous task running over a fraction of the dedicated pool of server processes
  - XIOS schedules dynamically the launching of the required services in free resources
  - Interconnection between models and services are done through the XIOS middleware which provide mechanism for grid and data flux exchange
  - A model is saw by the XIOS middleware like a specific service which generate data periodically
- XIOS-2 was providing IO services, but that was hardly coded



- XIOS-3 rewrote XIOS-2 functionalities in term of interconnected services
  - Rationalized way to exchange data flux through MPI partition
    - model<->service, service<->service, model<->model
    - Enabling model coupling
  - Description of services launching and models coupling remains described in a flexible way through external XLM files
  - Flexible management : services can run in separate resources or totally overlapping an other service resources

# XIOS-3 services : how is it working ?


mpirun -np 11 atm : -np 23 ocean : -np 20 land : -np 76 xios\_server.exe




mpirun -np 11 atm : -np 23 ocean : -np 20 land : -np 76 xios\_server.exe

MPI\_COMM\_WORLD


CALL `xios_initialize("model_id")` => Communicator splitting



mpirun -np 11 atm : -np 23 ocean : -np 20 land : -np 76 xios\_server.exe



CALL xios\_initialize("model\_id") => Communicator splitting



- split servers into pools through XML definition

- Allocated ressources defined by "pool" attributes

- global\_fraction (double) : fraction of whole server resources
- remain\_fraction (double) : fraction of remaining free servers
- nprocs (int) : number of servers
- remain (bool) : remaining free servers

```

<context id="xios">
  <variable id="using_server2"> false </variable>

  <pool_definition>      => 76 free
  <pool name="pool_atm" global_fraction="0.25">
    => 32 allocated, remaining 44


  </pool>
  <pool name="pool_ocean" nprocs="25">
    => 25 allocated, remaining 19

  </pool>


</pool_definition>

</context>

```



- “default\_pool” is created on remaining free resources
  - => 19 servers




- Launching services into allocated pools
  - Same attributes that for allocating pools

```

<context id="xios">
  <variable id="using_server2"> false </variable>

  <pool_definition>
    <pool name="pool_atm" global_fraction="0.25">
      <service name="gatherer" global_fraction="0.31" type="gatherer"/>
    </pool>
    <pool name="pool_ocean" nprocs="25">
      </pool>
    </pool_definition>
  </context>

```




- Launching services into allocated pools
  - Same attributes that for allocating pools

```

<context id="xios">
  <variable id="using_server2"> false </variable>

  <pool_definition>
    <pool name="pool_atm" global_fraction="0.25">
      <service name="gatherer" global_fraction="0.31" type="gatherer"/>
      <service name="writer1" remain_fraction="0.5" type="writer"/>
    </pool>
    <pool name="pool_ocean" nprocs="25">
      </pool>
    </pool_definition>
  </context>
    
```




- Launching services into allocated pools
  - Same attributing that for allocating pools

```

<context id="xios">
  <variable id="using_server2"> false </variable>

  <pool_definition>
    <pool name="pool_atm" global_fraction="0.25">
      <service name="gatherer" global_fraction="0.31" type="gatherer"/>
      <service name="writer1" remain_fraction="0.5" type="writer"/>
      <service name="writer2" nprocs="6" type="writer"/>
    </pool>
    <pool name="pool_ocean" nprocs="25">
      </pool>
    </pool_definition>
  </context>
    
```




- Lauching services into allocated pools
  - Same attributes that for allocating pools

```

<context id="xios">
  <variable id="using_server2"> false </variable>

  <pool_definition>
    <pool name="pool_atm" global_fraction="0.25">
      <service name="gatherer" global_fraction="0.31" type="gatherer"/>
      <service name="writer1" remain_fraction="0.5" type="writer"/>
      <service name="writer2" nprocs="6" type="writer"/>
      <service name="reader" remain="true" type="reader"/>
    </pool>
    <pool name="pool_ocean" nprocs="25">

    </pool>
  </pool_definition>
</context>
    
```




- Launching services into allocated pools
  - Same attributes that for allocating pools

```

<context id="xios">
  <variable id="using_server2"> false </variable>

  <pool_definition>
    <pool name="pool_atm" global_fraction="0.25">
      <service name="gatherer" global_fraction="0.31" type="gatherer"/>
      <service name="writer1" remain_fraction="0.5" type="writer"/>
      <service name="writer2" nprocs="6" type="writer"/>
      <service name="reader" remain="true" type="reader"/>
    </pool>
    <pool name="pool_ocean" nprocs="25">
      <service name="writer" nprocs="13" type="writer">
        <service name="reader" type="reader"/>
      </service>
    </pool>
  </pool_definition>
</context>
    
```

# S-3 services : how is it working ?



- Default services are launched on unallocated servers
- In order to define default behaviour
- Retro-compatibilty with XIOS2

- Targeted service is identified by pool name and service name => **id = pool\_name::service\_name**
- Can be assigned at context level => default behaviour
- Or Can be assigned at file level

```
<context id="atm" default_pool_writer="pool_atm" default_pool_writer="pool_atm" default_pool_reader="pool_atm" >
  <file_definition output_freq="1d">
```

```
    <file name="out1" mode="write" using_server2="true" gatherer="gatherer" writer="writer1">
```

```
      <field field_ref="field_out1"/>
    </file>
```

↳ Sent to service **pool\_atm::gatherer** chained to service **pool\_atm::writer1**

```
    <file name="out2" mode="write" writer="writer2">
```

```
      <field field_ref="field_out2"/>
    </file>
```

↳ Sent to service **pool\_atm::writer2**

```
    <file name="in" mode="read" reader="reader">
```

```
      <field field_ref="field_in"/>
    </file>
```

↳ Received from service **pool\_atm::reader**

```
</file_definition>
</context>
```

- With no service specification, we find the XIOS2 behaviour

```
<context id="land" >
```

```
  <file_definition output_freq="1d">
```

```
    <file name="land_out" mode="write" / >
```

```
    <file name="land_in" mode="read" />
```

⇒ Sent to service **default\_pool::default\_writer**

⇒ Received service **default\_pool::default\_reader**

```
  </file_definition>
```

```
</context>
```

## New middleware infrastructure to manage I/O services in a flexible way

- Only gatherer, writer and reader services are currently implemented
- Main interest is for performance tuning
  - using dedicated services for models, aggregating more parallelism

Potentiality will be fully exploited with future development of new services, which can be interconnected with I/O services

## Future plans are developing :

- Offload service : a piece of costly XIOS workflow diagnostic can be offloaded on dedicated resources
  - Short term
- Ensemble service : dedicated to efficient management of large ensemble runs
- AI services : deep learning training and inference could be done “In situ” and asynchronously
  - Making the bridge between the Fortran world of models and the Python world of deep learning technology
- User defined services
  - Users can write their own service for specific diagnostic
  - Could be written in Python to fully benefit of the software stack of python library

## New service infrastructure enable exchange of grid and data flux between different XIOS contexts

- Context can be attached to a service
- Context can be attached to a model
  - A model is saw like a service that produce specific data periodically

=> Exchange is now possible between 2 contexts running onto 2 different models

## XIOS coupling time scheme

- Fields and associated grids are described as usually in XML context file
- Field to be sent from source to a destination context are imbedded into “coupler\_out” elements

```
<coupler_out context="ocean::ocean_context">
  <field id="field3D_oce" field_ref="field3D"/>
</coupler_out>
```

Targeted context = dest\_model\_id::dest\_context\_id

- Field to be received are embedded into “coupler\_in” elements



```
<coupler_in context="atm::atm_context">
  <field id="field3D_atm" grid_ref="grid_atm"/>
</coupler_in>
```

Source context = src\_model\_id::src\_context\_id


- At context initialisation (close\_context\_definition) grid are sent and redistributed from source to destination context
- Remapping can be achieved by chaining existing transformation filters (horizontal and vertical remapping)
- In time loop, coupling fields can be sent and received from/into models using the standard Fortran interface
  - CALL xios\_send\_field("field\_out\_id", field\_out)
  - CALL xios\_recv\_field("field\_in\_id", field\_in)

Source model id = "atm"

Target model id = "ocean"



CALL xios\_send\_field("field3d",field)



**More complex configurations can easily be achieved by combining more of the XIOS workflow functionalities**


- 2 way coupling
- Coupling at different time step
- Exchanging averaged or cumulated fields...

**Some works is remaining to have a stable and efficient coupler**

- Couple from previous times step
- Restartability
- More interpolation methods
  - Currently first and second order conservative
- Dead-lock hunting

**Vision of future : a multitude of model components and services fully interconnected through an single middleware**

- Simple minimalist Fortran interface
- Flexible management
- Asynchronous data exchange through the MPI partition to exhibit more parallelism and concurrency
- Light weight coupling written in Python or Fortran
  - ⇒ User defined service



- **Stabilization and consolidation of the services and coupling functionalities**
  - Must be implemented and tested on a full ESM model (IPSL-ESM)
- **Development of new kinds of service**
  - Offload, ensemble, AI...
- **Development of a Python interface**
  - User defined services
- **Revisit the XIOS timeline management**
  - Time interpolations
  - Adaptative time step
  - Make XIOS restartable
- **GPU porting, accelerators**
  - Will be the main priority for the next years
  - Be easier by new recoding
  - CPU consumption in time loop is now localized in small fractions of code : connectors and filters
  - Incremental approach, filters after filters...
  - Not decided which technology to use : language based directive (OpenAcc, OpenMP), kokkos or others..